

DEREK RAYSIDE & ECE351 STAFF

ECE351 LAB MANUAL

UNIVERSITY OF WATERLOO

Copyright © 2024 Derek Rayside & ECE351 Staff
Compiled September 3, 2024

ACKNOWLEDGEMENTS:

- Prof Paul Ward suggested that we look into something with VHDL to have synergy with ECE327.
- Prof Mark Aagaard, as the ECE327 instructor, consulted throughout the development of this material.
- Prof Patrick Lam generously shared his material from the last offering of ECE251.
- Zhengfang (Alex) Duanmu & Lingyun (Luke) Li [1B Elec] wrote solutions to most labs in TXL.
- Jiantong (David) Gao & Rui (Ray) Kong [3B Comp] wrote solutions to the VHDL labs in ANTLR.
- Aman Muthrej and Atulan Zaman [3A Comp] wrote solutions to the VHDL labs in Parboiled.
- Michael Thiessen [3A Comp] improved the VHDL grammar in Parboiled.
- TA's Jon Eyolfson, Vajih Montaghani, Alireza Mortezaei, Wenzhu Man, and Mohammed Hassan.
- TA Wallace Wu developed the VHDL labs.
- High school students Brian Engio and Tianyu Guo drew a number of diagrams for this manual, wrote Javadoc comments for the code, and provided helpful comments on the manual.

Licensed under Creative Commons Attribution-ShareAlike (CC BY-SA) version 2.5 or greater.
<http://creativecommons.org/licenses/by-sa/2.5/ca/>
<http://creativecommons.org/licenses/by-sa/3.0/>

Contents

0	Overview	15	
0.1	How the Labs Fit Together	15	
0.2	Learning Progressions	17	
0.3	How this project compares to CS241, the text book, etc.	19	Compiler Concepts: call stack, heap Programming Concepts: version control, push, pull, merge, SSH keys, IDE, debugger, objects, pointers
0.4	Student work load	20	
0.5	How this course compares to MIT 6.035	21	
0.6	Where do I learn more?	21	
0.7	Full-Adder Circuit Example	22	
0.7.1	Picturing Git [†]	24	
0.8	What To Do This Week	25	
0.9	Metadata	27	
0.10	Checklist for Every Lab	27	
0.11	How to do these labs	28	
0.12	I think I found a bug in the skeleton code	29	
0.13	I want to change the skeleton code for my own usage	29	
0.14	Testing [†]	30	
0.15	Phases of Compilation [†]	31	
0.16	Engineers and Other Educated Persons	32	
1	Recursive Descent Parsing of \mathcal{W}	35	
1.1	Grammars & EBNF [†]	35	Compiler Concepts: regular languages, regular expressions, EBNF, recognizers, parsing, recursive descent, lexing, pretty-printing, abstract syntax tree (AST)
1.1.1	Summary of Grammar Notation Conventions	36	Programming Concepts: classes, objects, variables, aliasing, immutability, test- driven development
1.1.2	Derivations [†]	37	

1.2	Write a regular expression recognizer for \mathcal{W}	38
1.2.1	Test-Driven Development [†]	39
1.2.2	Complete the regular expression recognizer for \mathcal{W}	39
1.3	Write a recursive descent recognizer for \mathcal{W}^+	39
1.4	Write a pretty-printer for \mathcal{W}	40
1.5	Write a recursive descent parser for \mathcal{W}	40
1.6	Object Diagram of Example \mathcal{W} AST	41
1.7	Steps to success	42
1.8	Evaluation	44
1.9	Reading	44
1.9.1	Tiger Book	44
1.9.2	Programming Language Pragmatics	44
1.9.3	Web Resources	44
2	Transforming $\mathcal{W} \rightarrow \text{SVG}$ for Visualization	45
2.1	Write a translator from $\mathcal{W} \rightarrow \text{svg}$	46
2.2	Write a translator from $\text{svg} \rightarrow \mathcal{W}$	46
2.3	Introducing the Object Contract [†]	47
2.3.1	Engineering Apparatus and Controls [†]	47
2.3.2	Learning to write professional code [†]	49
2.3.3	Reading	49
2.4	Evaluation	50
2.5	Steps to success	50
3	Recursive Descent Parsing of \mathcal{F}	53
3.1	A Tale of Three Hierarchies [†]	53
3.1.1	Parse Tree	54
3.1.2	AST	54
3.1.3	A peak at the Expr class hierarchy	56

Compiler Concepts: trees, transformations, XML, SVG

Programming Concepts: object contract, object equality, mathematical equivalence classes, DOM *vs.* SAX parser styles, call-backs, iterator design pattern

Compiler Concepts: context-free grammars, LL(1) grammars, predict sets, parse trees, precedence, associativity, commutativity, program equivalence
Programming Concepts: inheritance, polymorphism, dynamic dispatch, type tests, casting, memory safety, composite design pattern, template design pattern, singleton design pattern, recursive functions, recursive structures, higher-order functions

3.2	Polymorphism & Dynamic Dispatch [†]	57
3.3	Write a recursive-descent recognizer for \mathcal{F}	58
3.4	Write a pretty-printer for the AST	58
3.5	Equals, Isomorphic, and Equivalent [†]	58
3.6	Write Equals and Isomorphic for the \mathcal{F} AST classes	59
3.7	Write a recursive-descent parser for \mathcal{F}	59
3.8	Steps to success	60
3.9	Common Missteps	60
3.10	Evaluation	61
3.11	Background & Reading	61
4	Circuit Optimization: \mathcal{F} Simplifier	63
4.1	Iteration to a Fixed Point and Termination [†]	64
4.2	Confluence [†]	65
4.3	Mathematical Properties of Binary Operators [†]	66
4.3.1	Commutativity	66
4.3.2	Associativity	66
4.4	Transforming BinaryExprs to NaryExprs	67
4.5	Identity Element & Absorbing Element [†]	68
4.6	Simplify Once	68
4.7	Object sharing in the simplifier	70
4.8	Class Representation Invariants [†]	71
4.9	Logical equivalence for boolean formulas [†]	72
4.10	BinaryExpr.examine() is a Higher-Order Function [†]	73
4.11	Choice of Data Structures for NaryExpr.children [†]	73
4.12	Evaluation	74
5	Parsing \mathcal{W} with Parboiled	75
5.1	Introducing Parboiled	76
5.2	Write a recognizer for \mathcal{W} using Parboiled	77
5.3	Add actions to recognizer to make a parser	77
5.3.1	An alternative Name() rule	78
5.3.2	The stack for our \mathcal{W} parser will have two objects	79

Compiler Concepts: intermediate languages, identity element, absorbing element, equivalence of logical formulas, term rewriting, termination, confluence, convergence

Programming Concepts: interpreter design pattern, template design pattern, representation invariants

Compiler Concepts: parser generators, Parsing Expression Grammars (PEG), push-down automata

Programming Concepts: domain specific languages (DSL): internal vs. external, debugging generated code, stacks

5.4	<i>Parsing Expression Grammars (PEGs)[†]</i>	79	
5.5	<i>Expressions, Statements, and Side Effects[†]</i>	80	
5.6	<i>A few debugging tips</i>	81	
5.7	<i>Evaluation</i>	81	
5.8	<i>Reading</i>	81	
6	<i>Parsing \mathcal{F} with Parboiled</i>	83	
6.1	<i>Write a recognizer for \mathcal{F} using Parboiled</i>	83	
6.2	<i>Add actions to recognizer to make a parser</i>	84	
6.3	<i>Composite Design Pattern[†]</i>	84	
6.4	<i>Template Method Design Pattern[†]</i>	85	
6.5	<i>Singleton Design Pattern[†]</i>	87	
6.6	<i>Evaluation</i>	88	
7	<i>Technology Mapping: $\mathcal{F} \rightarrow \text{Graphviz}$</i>	89	
7.1	<i>Interpreters vs. Compilers[†]</i>	90	
7.2	<i>Introducing the Interpreter Design Pattern[†]</i>	90	
7.3	<i>Introducing the Visitor Design Pattern[†]</i>	92	
7.3.1	<i>Is Visitor always better than Interpreter?[†]</i>	95	
7.4	<i>Hash Structures, Iteration Order, and Object Identity[†]</i>	99	
7.5	<i>Non-determinism: Angelic, Demonic, and Arbitrary[†]</i>	102	
7.6	<i>Translate one formula at a time</i>	102	
7.7	<i>Common Subexpression Elimination</i>	102	
7.8	<i>Designing a Common Subexpression Eliminator for \mathcal{F}</i>	103	
7.8.1	<i>An n^2 design using isomorphic and memory address</i>	104	
7.9	<i>Evaluation</i>	105	
8	<i>Simulation: $\mathcal{F} \rightarrow \text{Java}$</i>	107	
8.1	<i>Name Collision/Capture[†]</i>	107	
8.2	<i>Testing a code generator[†]</i>	110	
8.3	<i>Evaluation</i>	110	

Compiler Concepts: parser generators, Parsing Expression Grammars (PEG), push-down automata
Programming Concepts: domain specific languages (DSL): internal vs. external, debugging generated code, stacks

Compiler Concepts: common subexpression elimination
Programming Concepts: hash structures, iteration order, object identity, non-determinism, Visitor design pattern, tree traversals: preorder, postorder, inorder

Compiler Concepts: program generation, name capture
Programming Concepts:

9	<i>Simulation: $\mathcal{F} \rightarrow \text{x64 Assembly}$</i>	111	
9.1	<i>Reading the \mathcal{W} Program</i>	111	
9.2	<i>Register Allocation[†]</i>	111	
9.2.1	<i>Register Allocation in This Lab</i>	112	
9.3	<i>Assembling & Linking</i>	112	
9.4	<i>Evaluation</i>	112	
9.4.1	<i>Bonus Marks</i>	112	
10	<i>VHDL Recognizer & Parser</i>	115	
10.1	<i>Keywords and Whitespaces</i>	118	
10.2	<i>VHDL Recognizer</i>	118	
10.3	<i>VHDL Parser</i>	118	
10.4	<i>Engineering a Grammar[†]</i>	119	
10.5	<i>Evaluation</i>	120	
11	<i>VHDL \rightarrow VHDL: Desugaring & Elaboration</i>	121	
11.1	<i>VHDL \rightarrow VHDL: Desugaring</i>	121	
11.2	<i>Elaboration</i>	122	
11.2.1	<i>Inlining Components without Signal List in Architecture</i>	122	
11.2.2	<i>Inlining Components with Signal List in Architecture</i>	124	
11.2.3	<i>Inlining Components with Processes in Architecture</i>	126	
11.2.4	<i>Notes</i>	126	
11.3	<i>Evaluation</i>	126	
12	<i>VHDL Process Splitting & Combinational Synthesis</i>	127	
12.1	<i>Process Splitter</i>	127	
12.2	<i>Splitter Notes</i>	129	
12.3	<i>Synthesizer</i>	130	
12.3.1	<i>If/Else Statements</i>	130	
12.3.2	<i>Example Synthesizing Signal Assignment Statements</i>	131	
12.3.3	<i>Example Synthesizing If/Else Statements</i>	131	

Compiler Concepts:
Programming Concepts:

Compiler Concepts: desugaring, function
inlining
Programming Concepts:

12.4	<i>Evaluation</i>	132	
13	<i>Simulation: $\mathcal{F} \rightarrow \text{Assembly}$</i>	133	
13.1	<i>Which assembler to use?</i>	133	
13.2	<i>MASM32 Configuration</i>	134	
13.3	<i>MASM32 JNI and x86 Assembly</i>	134	
13.4	<i>Evaluation</i>	137	
14	<i>Simulation: $\mathcal{F} \rightarrow \text{JVM}$</i>	139	
14.1	<i>How javac compiles boolean logical operators[†]</i>	139	<i>Compiler Concepts: code generation</i>
14.2	<i>How we will generate code</i>	141	<i>Programming Concepts: assembly</i>
14.3	<i>Negation</i>	144	
14.4	<i>Evaluation</i>	144	
14.5	<i>Bonus Lab: \mathcal{F} to x86/x64/etc</i>	144	
A	<i>Extra VHDL Lab Exercises</i>	145	
A.1	<i>Define-Before-Use Checker</i>	145	
A.1.1	<i>Some Tips</i>	146	
A.2	<i>Inline \mathcal{F} intermediate variables</i>	146	
B	<i>Advanced Programming Concepts</i>	147	
B.1	<i>Immutability</i>	147	
B.2	<i>Representation Invariants</i>	148	
B.3	<i>Functional Programming</i>	148	
C	<i>Design Patterns</i>	149	
C.1	<i>Iterator</i>	149	See the <i>Design Patterns</i> book or
C.2	<i>Singleton</i>	149	Wikipedia or SourceMaking.com
C.3	<i>Composite</i>	149	
C.4	<i>Template Method</i>	150	
C.5	<i>Interpreter</i>	150	
C.6	<i>Visitor</i>	150	

D	Lab Instructor Notes (GitLab)	151
D.1	Repository URLs	151
D.2	Setting Up	151
D.2.1	Creating the Wildcard Repo ece351/term	151
D.2.2	Notes/PDF Repository	152
D.2.3	Lib Repository	152
D.2.4	Course Offering Metadata Repository	152
D.2.5	Skeleton Repository	152
D.3	Forking Student Repos	154
D.4	Marking Script: Build.xml	154
D.5	Web User Interface for Gitolite Server	154
D.6	Exporting to Skeleton: export.sh	155
D.6.1	Update header.txt at the start of term	155
D.6.2	Export on an as-needed basis	155
D.6.3	Exporting to skeleton dev branch	155
D.6.4	Files to Release for Each Lab	156
E	Lab Instructor Notes (gitolite)	159
E.1	Repository URLs	159
E.2	Setting Up	159
E.2.1	Creating the Wildcard Repo ece351/term	159
E.2.2	Notes/PDF Repository	160
E.2.3	Lib Repository	160
E.2.4	Course Offering Metadata Repository	160
E.2.5	Skeleton Repository	160
E.3	Forking Student Repos	162
E.4	Marking Script: Build.xml	162
E.5	Web User Interface for Gitolite Server	162
E.6	Exporting to Skeleton: export.sh	163
E.6.1	Update header.txt at the start of term	163
E.6.2	Export on an as-needed basis	163
E.6.3	Exporting to skeleton dev branch	163
E.6.4	Files to Release for Each Lab	164

F Bibliography 167

[†]denotes conceptual sections

List of Figures

1	Overview of ECE351 Labs	15
2	Lab dependencies	16
3	Descriptions of individual labs with the compiler and programming concepts introduced in each	18
4	Student hours spent on labs in winter 2013	20
5	Student hours spent on labs in summer 2013. (Partial data to lab 9.)	20
6	Source code for full adder circuit (VHDL)	22
7	Input waveform for full adder (\mathcal{W})	22
8	Boolean formulas for full adder (\mathcal{F} generated from source code in Figure 6)	22
9	Gates for full adder (generated from formulas in Figure 8)	22
10	Input and output waveforms for full adder (generated from formulas in Figure 8 and input waveform in Figure 7)	23
11	Git topology for ECE351	24
12	Phases of compilation and the labs in which we see them	31
1.1	Example waveform file for an or gate	35
1.2	Grammar for \mathcal{W} .	35
1.3	Object diagram for example \mathcal{W} AST	41
1.4	Steps to success for §1	42
1.5	Legend for Figure 1.4	43
2.1	Example waveform file for an or gate	45
2.2	Rendered svg of \mathcal{W} file from Figure 2.1	45
2.3	First 10 lines of svg text of \mathcal{W} file from Figure 2.1	46
2.4	Code comprehension exercises	47
2.5	Infographic and photo of the International Kilogram Prototype	48
2.6	Steps to success for §2	50
2.7	How §1 and §2 fit together	51
3.1	LL(1) Grammar for \mathcal{F}	53
3.2	Object diagram of AST for \mathcal{F} program $X \leq A \text{ or } B;$	54
3.3	Object diagram for example \mathcal{F} AST	55
3.4	Some highlights from the Expr class hierarchy	56

3.5	Steps to success for §3	60
4.1	Code listing for Expr.simplify() showing iteration to a fixed point	64
4.2	Example of two confluent rewrites converging on a common solution	65
4.3	Four trees that represent logically equivalent circuits. The sorted n-ary representation makes equivalence comparisons easier.	67
4.4	NaryExpr.simplifyOnce()	68
4.5	Simplifications for \mathcal{F} programs	69
4.6	Object sharing in the \mathcal{F} simplifier	70
4.7	NaryExpr.repOk()	71
4.8	Equivalent truth tables	72
4.9	BinaryExpr.examine() is a higher-order function	73
5.1	Snippet of a recognizer written in Parboiled's DSL	77
5.2	Snippet of a parser written in Parboiled's DSL	78
5.3	Stack of Parboiled parser from Figure 5.2 while processing input 'Ren Stimpy '.	78
5.4	An alternative definition of the Name() rule in Figure 5.2	78
5.5	Snippet of a parser written in Parboiled's DSL (Figure 5.2). Extended with debugging commands. Corresponding to the recognizer snippet in Figure 5.1.	79
6.1	LL(1) Grammar for \mathcal{F} (reproduced from Figure 3.1)	83
6.2	UML class diagram for Composite Design Pattern	84
6.3	Template Design Pattern illustrated by UML Class Diagram	85
6.4	The only two instances of class ConstantExpr	87
6.5	Preventing clients from instantiating class ConstantExpr	87
7.1	Example Graphviz input file and rendered output	89
7.2	A UML class diagram for a subset of \mathcal{F} expression class hierarchy	91
7.3	UML class diagram for \mathcal{F} Expr hierarchy, Interpreter design pattern	91
7.4	Matrix view of Interpreter Design Pattern	91
7.5	UML class diagram for \mathcal{F} Expr hierarchy, Visitor design pattern	92
7.6	UML class diagram for \mathcal{F} Expr hierarchy, Visitor design pattern	93
7.7	Matrix view of Visitor Design Pattern	93
7.8	Abstract super class for Visitors for \mathcal{F} expression ASTs	93
7.9	Signature and implementation of accept method.	94
7.10	The traverse methods of ExprVisitor	96
7.11	Implementation of PostOrderExprVisitor	97
7.12	Implementation of ExtractAllExprs	98
7.13	Iteration order for different data structures	99
7.14	Object identity for different data structures	100
7.15	Synthesized circuit <i>without</i> common subexpression elimination	104
7.16	Synthesized circuit <i>with</i> common subexpression elimination	104

8.1 Simulator for \mathcal{F} program $x \leq a$ or b ;	108
8.2 Implementation of DetermineInputVars	109
8.3 Dataflow of LAB8	110
9.1 x64 assembly for $x \leq a$;	114
10.1 Grammar for VHDL	116
10.2 Object diagram for example VHDL AST	117
10.3 Implementation of rule used to match the keyword 'signal'	118
11.1 VHDL program used to illustrate elaboration.	123
11.2 Elaborated architecture body, four_port_structure.	124
11.3 Extension of the VHDL program shown in Figure 11.1.	125
11.4 Elaborated architecture body, eight_port_structure.	125
12.1 Example VHDL program used to illustrate process splitting.	128
12.2 The resulting VHDL program of Figure 12.1 after process splitting.	129
12.3 Example used to illustrate synthesizing assignment statements.	131
12.4 Synthesized output of the program in Figure 12.3.	131
12.5 Example used to illustrate synthesizing assignment statements.	132
12.6 Synthesized output of the program in Figure 12.5.	132
13.1 Generate this method in assembly rather than in Java	133
13.2 Example assembly for an F Program	135
13.3 Example definitions file for a dynamic link library (DLL)	136
A.1 VHDL program used to illustrate signal declarations and the use of undefined signals in signal assignment statements.	146

Compiler Concepts: call stack, heap
 Programming Concepts: version control,
 push, pull, merge, SSH keys, IDE,
 debugger, objects, pointers

Overview

0.1 How the Labs Fit Together

The overall structure of our VHDL synthesizer and simulator is depicted in Figure 1.

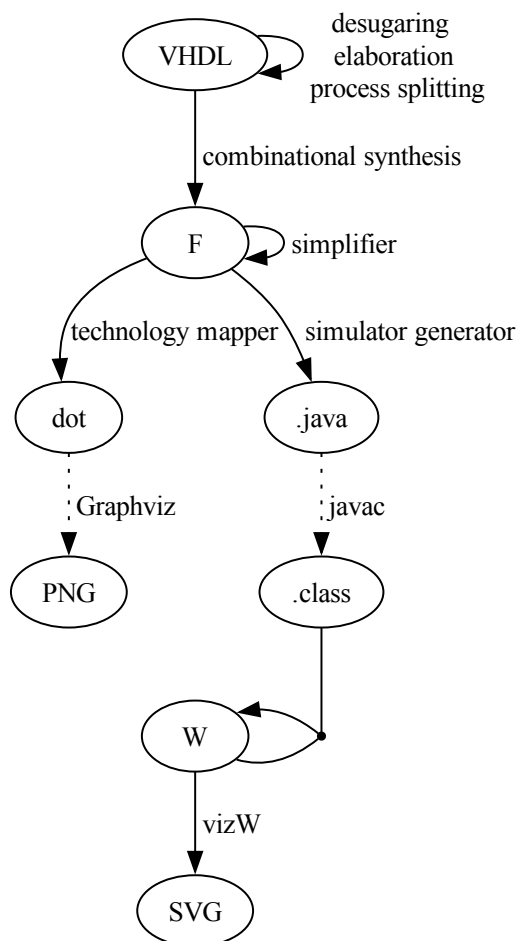


Figure 1: Overview of ECE351 Labs.

Nodes represent file types (e.g., VHDL). All of these file types, with the exception of .class and PNG files, are text files.

Edges represent translators between different file types. Solid edges represent translators that we will implement in ECE351. Dotted edges represent translators provided by third-parties such as Sun/Oracle (javac) or AT&T Research (dot).

The three-part edge between .class and W nodes is intended to indicate that the .class file we generate will read a waveform (W) file as input and write a waveform (W) file as output.

Labels on edges describe the translation(s) performed.

Numbers on edges indicate the order in which we will implement those translators in ECE351. For example, the first translator that we will implement will transform waveform files to svg files (svg is an XML-based graphics file format).

The general direction of our work in ECE351 will be from the bottom of the figure towards the top. We will start with file types that have the simplest grammars and work towards file types with more complicated grammars.

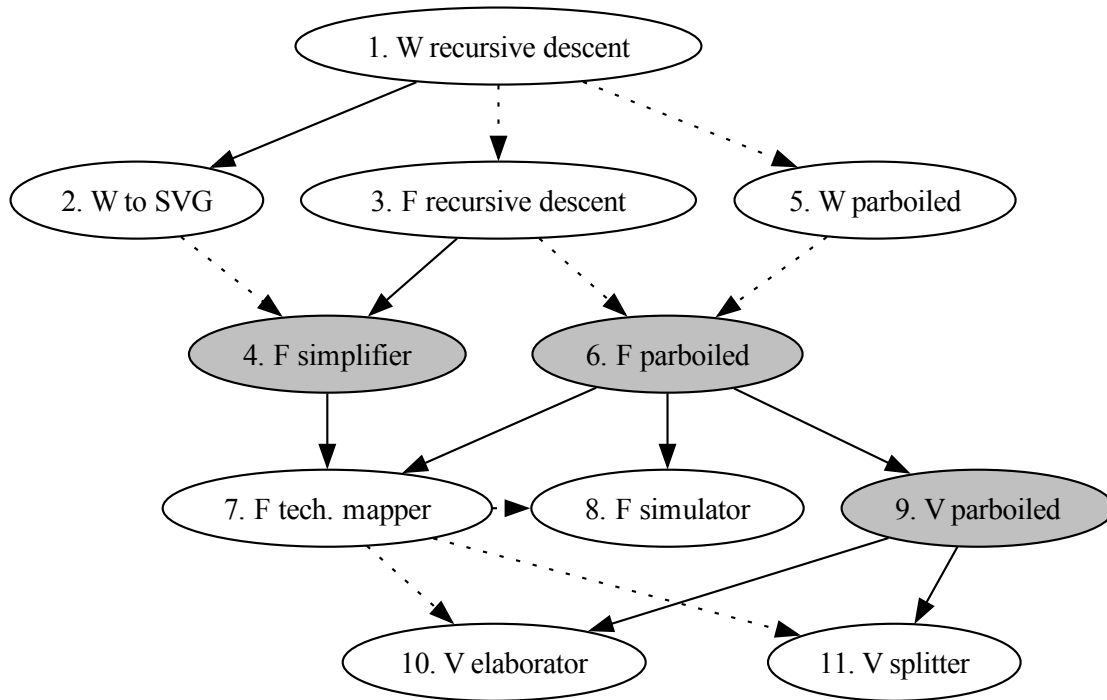


Figure 2: Lab dependencies. A solid line indicates that the code is needed to test the future lab. For example, you need to have a working implementation of LAB1 in order to test LAB2.

A dotted line indicates that ideas from one lab feeds into the next lab. For example, you learn the grammar of \mathcal{W} in LAB1 and then use that idea again in LAB5. Not all of the dotted lines are shown, in order to simplify the graph.

The shaded nodes indicate labs that you need to focus on because they are used by future labs. If you must skip a lab, skip something like 2, 5, 8, 10, or 11 that are not used as much by future labs.

0.2 Learning Progressions

There are a variety of undergraduate compiler projects. Almost all of them involve compiling a subset of some imperative programming language to assembly. In the past Pascal was the popular language to compile, whereas these days Java is.

The ECE351 labs are different. On a superficial level, the difference is that these labs use a subset of VHDL as the input language, and produce circuit gate diagrams and simulations as outputs. The deeper difference is that our project is designed around two parallel *learning progressions* rather than around the logical structure of a compiler. This project comprises both a *programming skills progression* and a *compiler concepts progression*.

The key technical decision that enables these progressions is to restrict our subset of VHDL to combinational circuits: no loops, no timing, *etc.*. From this decision flows the design of a simple intermediate language, \mathcal{F} , for boolean formulas, and the design of a simple auxiliary language, \mathcal{W} , for boolean waveforms. The project progresses from the simplest language (\mathcal{W}) to the most complex (VHDL), performing parsing, transformation, and translation on each of the three languages. Repetition with increasing complexity (hopefully) leads to mastery.

The idea of a learning progression has been used in hockey for several decades. It is recently attracting attention in educational circles. For example, suppose that the goal of the practice is to have the players skate quickly around the circles with the puck. To progress to that goal the team might start with skating around circles, then skating in a straight line with the puck, then skating around the circles slowly with the puck, and finally skating around the circles quickly with the puck. The skills are first practiced in isolation and at slower speeds, and then finally all together at high speed.

#	Description	Compiler Concepts	Programming Concepts
0	Prelab	call stack, heap	version control, push, pull, merge, SSH keys, IDE, debugger, objects, pointers
1	Parsing \mathcal{W} by recursive descent	regular languages, regular expressions, EBNF, recognizers, parsing, recursive descent, lexing, pretty-printing, abstract syntax tree (AST)	classes, objects, variables, aliasing, immutability, test-driven development
2	Translating \mathcal{W} to SVG (visualization)	trees, transformations, XML, SVG	object contract, object equality, mathematical equivalence classes, DOM <i>vs.</i> SAX parser styles, call-backs, iterator design pattern
3	Parsing \mathcal{F} by recursive descent	context-free grammars, LL(1) grammars, predict sets, parse trees, precedence, associativity, commutativity, program equivalence	inheritance, polymorphism, dynamic dispatch, type tests, casting, memory safety, composite design pattern, template design pattern, singleton design pattern, recursive functions, recursive structures, higher-order functions
4	Simplifying \mathcal{F} programs (optimization)	intermediate languages, identity element, absorbing element, equivalence of logical formulas, term rewriting, termination, confluence, convergence	interpreter design pattern, template design pattern, representation invariants
5	Parsing \mathcal{W} with a parser generator	parser generators, Parsing Expression Grammars (PEG), push-down automata	domain specific languages (DSL): internal <i>vs.</i> external, debugging generated code, stacks
6	Parsing \mathcal{F} with a parser generator	parser generators, Parsing Expression Grammars (PEG), push-down automata	domain specific languages (DSL): internal <i>vs.</i> external, debugging generated code, stacks
7	Translating \mathcal{F} to Graphviz (technology mapping)	common subexpression elimination	hash structures, iteration order, object identity, non-determinism, Visitor design pattern, tree traversals: preorder, postorder, inorder
8	Translating \mathcal{F} to Java (circuit simulation)	program generation, name capture	
9	Parsing VHDL with a parser generator		
10	VHDL elaboration		
11	VHDL process splitting and translation to \mathcal{F} (combinational synthesis)	desugaring, function inlining	
B	Translating \mathcal{F} to assembly	instruction selection, register allocation	assembly, linking

* 'DP' stands for 'design pattern'; 'B' stands for bonus lab

Figure 3: Descriptions of individual labs with the compiler and programming concepts introduced in each

0.3 How this project compares to CS241, the text book, etc.

	ECE351	CS241	Tiger	CS444	MIT 6.035
Language(s)	VHDL, \mathcal{F} , \mathcal{W}	Java	Java	Java	Java
Compiler Phases					
Parsing	✓	✓	✓	✓	✓
Symbol tables		✓	✓	✓	✓
Type checking		?	✓	✓	✓
Dataflow analysis			○	?	✓
Optimization	✓		○	?	✓
Translation	✓	✓	✓	✓	✓
Assembly	○	✓	✓	✓	✓
Pedagogy					
Skills Progression	✓				
Concept Progression	✓				
Background	✓				
Tests	✓				
Workload	✓	✓	✓	×2	×2

0.4 Student work load




Figure 4: Student hours spent on labs in winter 2013. The target is five hours per lab. About half the students hit this target on most labs, as indicated by the bars in the middle of the boxes. On most labs, 75% of the students completed the lab within eight hours, as indicated by the top of the boxes.

Some students took much longer than eight hours. If you are one of those students, please consider taking advantage of our course collaboration policy. The point of the collaboration policy is for you to learn more in less time.

The exceptional labs that took more time were 4 and 9. This term we will be increasing the TA resources for lab 4, and also changing the way we teach lab 9, in an effort to get these numbers down.




Figure 5: Student hours spent on labs in summer 2013. (Partial data to lab 9.)

o.5 How this course compares to MIT 6.035

MIT's (only) undergraduate compilers course is 6.035. It differs from ECE351 in a two important ways:

- a. 6.035 is rated at 12 hours per week for 14 weeks, whereas ECE351 is rated at 10 hours per week for 12 weeks: a 48 hour nominal difference. Moreover, 6.035 makes no effort to keep the workload near the course rating, whereas ECE351 actively tracks student hours and over half the students get the work done in the rated time. So the actual difference is much larger than the nominal.

6.035 also comes in an 18 hour per week variant.

- b. 6.035 is an elective course, whereas ECE351 is a required course. Elective courses only get students who have the interest and ability; required courses get everyone. On the flipside, not every graduate of MIT EECS will know compilers: some will be totally ignorant. At UW we guarantee a minimum level of quality in our graduates (this approach is also required for CEAB accreditation).

o.6 Where do I learn more?

If you are interested in learning more about compilers at UW your next step is CS444. If you take ECE351 and CS444 then you will know more than if you took MIT 6.035. CS462 covers the theoretical end of formal languages and parsing, although this knowledge won't greatly improve your practical skills. In the graduate curriculum CS644 and CS744 are offered on a semi-regular basis. There are also a number of graduate courses in both ECE and CS on program analysis that are offered irregularly.

0.7 Full-Adder Circuit Example

```

entity full_adder is port (
    A, B, Cin: in bit;
    S, Cout: out bit
);
end full_adder;

architecture full_adder_arch of full_adder is
begin
    S <= (A xor B) xor Cin;
    Cout <= ((A xor B) and Cin) or (A and B);
end full_adder_arch;

```

A: 0 1 0 1 0 1 0 1;
 B: 0 0 1 1 0 0 1 1;
 Cin: 0 0 0 0 1 1 1 1;

```

S <= (((not (((not A) and B) or ((not B) and A))) and Cin) or
      (((not A) and B) or ((not B) and A)) and (not Cin));
Cout <= (((not A) and B) or ((not B) and A)) and Cin
      or (A and B);

```




Figure 6: Source code for full adder circuit (VHDL)

Figure 7: Input waveform for full adder (\mathcal{W})

Figure 8: Boolean formulas for full adder (\mathcal{F} generated from source code in Figure 6)

Figure 9: Gates for full adder (generated from formulas in Figure 8)




Figure 10: Input and output waveforms for full adder (generated from formulas in Figure 8 and input waveform in Figure 7)

0.7.1 Picturing Git[†]

Git is a powerful *distributed version control system*. Figure 11 will help you understand how we use it in ECE351. If you want to read up on Git, here are some references:

- <http://git-scm.com/book/en/Getting-Started-Git-Basics>
- <http://git-scm.com/book/en/Git-Basics-Getting-a-Git-Repository>

While Git is quite powerful, its usability leaves something to be desired. Usually the fastest way to learn Git is to have someone who knows it sit with you and teach you what you need to do. Come in to the lab for help. Or ask a classmate or friend — anyone who is willing to help you. Getting help with Git is outside the scope of the lab collaboration policy.

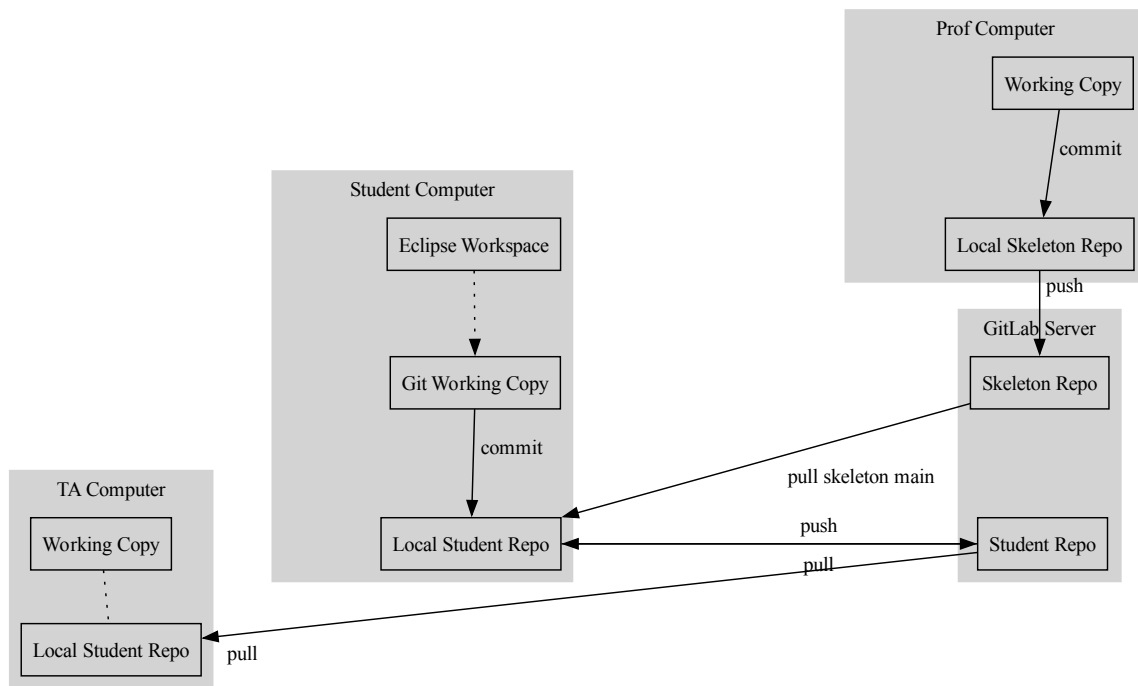


Figure 11: Git topology for ECE351

o.8 What To Do This Week

Install and Configure Software

- ☐ Create an ssh key (with a password!)
`ssh-keygen -t rsa -C "USERNAME@uwaterloo.ca"`
- ☐ GitLab (git.uwaterloo.ca)
 - ☐ Make account on git.uwaterloo.ca
 - ☐ Upload your ssh public key to git.uwaterloo.ca
 - ☐ Set up your ssh agent
 - * Might be PuTTY's ssh agent
 - * Might be adding eval 'ssh-agent' to your dotfiles
- ☐ Eclipse (or other IDE you prefer)
 - ☐ Download Eclipse JDT from <http://www.eclipse.org/>
 - ☐ Might need to install Java JDK (not JRE)
 - ☐ Configure Git at command line (if you are using it that way)
 - * `git config --global user.name "FIRSTNAME LASTNAME"`
 - * `git config --global user.email "USERNAME@uwaterloo.ca"`
 - * `git config --global push.default simple`
 - * `git config --global color.ui "true"`
 - ☐ Connect Eclipse to GitLab (might need your ssh key)
 - ☐ Configure Eclipse's JUnit Launcher to enable assertions:
 - * Go to **Window / Preferences / Java / JUnit**
 - * Check **Add -ea to VM arguments for new launch configurations**
- ☐ Graphviz (you can live without this, but nice to have)

Common Problem 1: JDK vs. JRE test fails. In order to fix this, do the following. Click Window and go to Preferences. In the tree view select Java, then Installed JREs and click Add. Hit Next (Standard VM should be selected already), then Directory, navigate to C:\Program Files\Java\jdk, or something similar, and hit Finish. Finally, check the jdk in the Installed JREs pane and click OK.

Common Problem 2: testJUnitConfiguration fails, but you have followed these steps. Probably you tried to run TestPrelab before configuring Eclipse's JUnit launcher. The configuration only adds '-ea' to new run configurations; it does not change existing run configurations. You now need to change the TestPrelab run configuration individually so that it includes '-ea' in the VM arguments.

Common Problem 3: Lots of compiler errors. Ensure that Eclipse is configured to use compiler/language level 11 or greater — and that the runtime JDK/JRE is set to 11.

Clone and Connect to Course Repositories

See §0.7.1 above.

- ☐ `git clone ist-git@git.uwaterloo.ca:ece351-TERM/pdfs ece351-pdfs`
- ☐ `git clone --recursive ist-git@git.uwaterloo.ca:ece351-TERM/class/USERID ece351-labs`
- ☐ `cd ece351-labs`
- ☐ `git remote add skeleton ist-git@git.uwaterloo.ca:ece351-TERM/skeleton`

Replace TERM with the term number, e.g. '1201'. Replace USERID with your user ID, e.g. 'p24gill'.

Show us that you can do stuff

- ☐ `git pull skeleton main`
- ☐ `git submodule init`
- ☐ `git submodule update`
- ☐ Check that there is a *non-empty* lib directory
- ☐ Run TestPrelabConfig.java
- ☐ Run TestPrelabExceptions.java
- ☐ Run TestImports.java
- ☐ Run build.xml (this is the marking script)

- ☐ Right-click on a single test to run it individually
- ☐ Edit `meta/hours.txt` to specify your time spent getting set up
- ☐ Edit `meta/collaboration.txt`
- ☐ `git commit -m 'edited meta files' meta/*.txt`
- ☐ `git push`

See §0.9 below.

In `collaboration.txt`, indicate that for LAB0 you were mentored by one of the course staff (this is just an exercise to check that you know who the staff are and that you know how to edit the `collaboration.txt` file properly).

0.9 Metadata

Your workspace has a directory named *meta* that contains the following two files in which you can describe a few things about how your work on the lab went.

collaboration.txt To record your collaborators. Each line is a triple of lab number, collaboration role, and userid. Legal values for collaboration roles are: converser, partner, mentor, protege. The role field describes the role of the *other* person, so lab2 mentor jsmith says that J Smith was your mentor for lab2. Similarly, lab3 protege jsmith says that J Smith was your protégé for lab3 (*i.e.*, you were his mentor). Both parties are required to report collaborations. If you collaborated with more than one person on a lab then you should put multiple lines into this file for that lab: one line for each collaborator on each lab.

hours.txt Estimate of the hours you worked on each lab. This file will have a line for each lab like so: lab1 5 (indicating five hours spent on lab 1). For pre-lab / computing environment time, use lab0.

You must edit this file for every lab, even if all of your collaborations were just conversation.

These data will be used solely for the staff to assess the difficulty of the labs. This assessment will be made in aggregate, and not on an individual basis. These data will not be used to assess your grade. However, we will not mark your lab until you report an estimate of your hours.

0.10 Checklist for Every Lab

BEFORE YOU START WORKING ON THE LAB:

- ☐ `cd ~/git/ece351-labs`
- ☐ `git checkout main`
- ☐ `git pull`
- ☐ `git pull skeleton main`
- ☐ `git submodule update`
- ☐ manually resolve any conflicts

DURING/AFTER EVERY WORKING SESSION:

- ☐ Run tests in Eclipse
- ☐ Run tests via `build.xml`
- ☐ `git commit -am 'description of work you did since last commit'`
- ☐ `git push`

BEFORE THE DEADLINE:

- ☐ Update `meta/collaboration.txt`
- ☐ Update `meta/hours.txt`
- ☐ `git add meta`
- ☐ `git commit -m 'updated metadata for labX'`
- ☐ `git push`

0.11 How to do these labs

The lab manual will tell you what files you need to edit, what libraries you should use, and what tests to run. The lab manual will also explain what you need to do at each point in the code, and the background concepts you need to understand the lab. Most of the dagger (†) sections in the lab manual are to explain these background concepts. Every place in the skeleton code that you need to edit is marked by both a `Todo351Exception` and a `TODO` marker. We will discuss the next week's lab in class every Friday.

Despite this clear and explicit instruction, some students have difficulty getting started on these labs. Why? What is missing from the instructions described above? *The lab manual doesn't tell you the order in which you should edit the files.* There is an important reason for this: execution order is the wrong way to think about developing object-oriented software (most of the time). The right way to think about object-oriented software is to focus on developing cohesive and extendible modules (classes). In this way object-oriented programming (in the large) is mentally quite different from procedural programming (in the small). In procedural programming (in the small) one thinks primarily about the order in which the steps are performed.

Now don't misunderstand this to mean that execution order doesn't matter: it does. It's just that execution order is generally a separate design concern from modularity and extensibility. The best way to understand the execution order of a large object-oriented program is to run it and observe it in the debugger or via some other tracing mechanism (e.g., `printf`).

If you want to figure out what code you should edit first, run the test harness and see where an exception is thrown. Fix that. Then run again, get another exception, *etc.*. The 'fix' step might be quick and easy, or it might require reading several pages of the lab manual to understand what needs to be done at that particular point. Remember, the lab manual describes everything that needs to be done and where it needs to be done, it just doesn't describe the order in which you should do it. Aligning your editing order with the program execution order is one way to guide your work that will help you build an understanding of the code.

Thinking on your own to develop an understanding of the skeleton code is an important part of these labs. I promise you that it takes less time and effort to study 8,000+ lines of code than to write it from scratch.

In ECE250 you did object-oriented programming in the small. That is, you defined structures that spanned one or two classes and operations on those structures that spanned one or two methods. The programs you worked on were perhaps several hundred lines long.

In ECE351 you will work on a code base that is over 8,000 lines of code, of which you will write about 1000 lines, at an average rate of about 100 lines per week. The structures we will work with are defined across dozens of classes, and the operations on those structures are similarly defined across dozens of methods. This is the first time in the ECE curriculum that you are exposed to *programming in the large*. At this scale, modularity and code structure really matter.

By the standards of modern industrial software development, 8,000+ lines is just approaching medium sized. The code structuring ideas you will learn in these labs can take you up to maybe 100,000 lines: beyond that you will need new ideas and modularity mechanisms.

0.12 I think I found a bug in the skeleton code

You are welcome to do what you like with the skeleton code. Our recommendation if you want to make a change is the following:

- a. Report the change you want to make (in the forum, or to a staff member). Preferably by generating a patch.
- b. We tell you that the change is misguided and you really want to do something else.
- c. We say thanks and we patch the skeleton code with the change so everyone can use it. Potentially you earn bonus marks for participation.

If you make the change yourself without reporting it then you lose out on class participation points, and if someone else reports the change and we go to apply a patch the patch will fail on your code. This might or might not end up being a problem for you.

0.13 I want to change the skeleton code for my own usage

You may change the skeleton code for your own usage. Perhaps you have a better idea of how to write something. If you make such changes, be careful to not create more problems than you solve. Some things to look out for include:

- Breaking the JUnit test harnesses.
- Breaking some other code that depends on the code you are changing.
- Messing up one of the design patterns that are an explicit part of what you are supposed to be learning.
- Changing something from immutable to mutable because you do not want to learn to work with immutable data. First, this is depriving yourself of one of the important lessons of the labs. Second, you will likely be introducing some bugs that will be very difficult to fix later. Immutable data is a good engineering practice that helps you avoid many classes of difficult to diagnose and fix bugs.

Algorithmic changes are usually safe, because the algorithmic parts of this code are usually encapsulated.

0.14 Testing[†]

Just because your code passes all of the tests does not mean it is correct: there could be some as of yet unidentified test that it does not pass. Worse, that not-yet-identified test might occur during some lab later in the term.

This might be the first course in which you have to write a ‘real’ program: that is, a non-trivial program that you will have to depend on in the future. First, the labs are non-trivial: The total size of the code for this course is about 8500 lines, of which you will have to write about 1500, and roughly 7000 will be provided for you. Second, the labs are cumulative: Each week you will run some labs that you wrote in the past. At the end of the term you will run all of them together. So testing will be important.

TEST INPUTS MAY BE GENERATED either *manually* or automatically with a testing tool. There are two main approaches used by tools: *systematic* and *random*. Because the space of possible inputs is large, possibly infinite, systematic approaches tend to generate all small inputs and no big inputs. Random approaches will generate a mixture of small and large inputs. Programmers can use human insight to generate interesting inputs that might be missed by automated tools.

THERE ARE DIFFERENT STRATEGIES for evaluating whether a test passed or failed. One is to check the computed result against a known answer. In other words, the test suite consists of a set of known input/output pairs.

Another strategy is to check general properties of functions. For example, that a function never returns null. An advantage of this property-based approach is that one only needs the test inputs — the corresponding outputs do not need to be known. Some properties that we will be interested in for this course include:

<i>reflexive</i>	$x.\text{equals}(x)$
<i>symmetric</i>	$x.\text{equals}(y) \Rightarrow y.\text{equals}(x)$
<i>transitive</i>	$x.\text{equals}(y) \text{ and } y.\text{equals}(z) \Rightarrow x.\text{equals}(z)$
<i>antisymmetric</i>	$x \leq y \text{ and } y \leq x \Rightarrow x = y$
<i>total</i>	$x \leq y \text{ or } y \leq x$
<i>idempotent</i>	$f(x) = f(f(x))$
<i>invertible</i>	$f'(f(x)) = x$

Some functions do not have any of these properties. But if we expect a function to have one or more of these properties then it is a good idea to test for that property explicitly.

In particular, the first three properties define a mathematical *equivalence relation*. We expect the `equals()` method in Java to represent a mathematical equivalence relation.

A *total order* is defined to be transitive, antisymmetric, and total.

Program testing can be used to show the presence of bugs, but never to show their absence! — Edsger W. Dijkstra


The Korat tool generates test inputs systematically based on representation invariants.

The Randoop tool generates inputs randomly.

The integers, for example, are totally ordered. We expect the `compareTo()` method in Java to represent a total order.

A GOOD TEST SUITE will contain some of everything mentioned above: test inputs that are generated manually and automatically, both systematically and randomly; evaluations that look at specific input/output pairs and at general properties.

0.15 Phases of Compilation[†]



In a previous offering of this course the staff provided tests that only looked at general properties and none that examined specific input output pairs. It turned out that students could write code that had the general property without actually computing anything useful. They became unhappy when they discovered that their code that passed the staff-provided tests did not actually work when they wanted to run it on a future lab.

Figure 12: Phases of compilation and the labs in which we see them. The scanner, parser, and type checker are considered the *front end*, whereas the optimizer and code generator are considered the *back end* of the compiler.

0.16 Engineers and Other Educated Persons

There are a number of important general skills that all educated persons, including engineers, should possess. The first two years of engineering education need to focus on technical details in isolation in order for you to have the technical competency to tackle more interesting things. This focus on minutiae sometimes comes at the cost of limited growth in these larger skills. In this course you will not only learn lots of tiny technical details, but you will also need to develop some of the larger skills that all educated persons should possess.

THE ABILITY TO QUICKLY FIND RELEVANT INFORMATION. Information is more accessible now than at any previous point in human history. There are a wide variety of sources available to you: the course notes, the lab manual, old exams, the skeleton code, the recommended text books, other text books, lecture notes and slides and videos from other professors, *etc.* You should be facile in using all of these sources to find what you are looking for.

Books, in particular, have helpful features for navigating them, such as the table of contents, the index, and the preface or introduction. You should know how to use these features.

THE ABILITY TO QUICKLY ASSESS WHAT INFORMATION IS RELEVANT. As old-time engineers used to say, there is always more heat than light. Learn to see the light without being overwhelmed by the heat.

When doctors, lawyers, accountants, and other professionals assess a case study problem, the first order of business is to discern what the relevant facts and issues are. Those professions explicitly train their people to separate the wheat from the chaff. Engineering education, especially in the first two years, is often guilty of spoon-feeding its students with only and all of the relevant information, and hence developing a sense of intellectual complacency and entitlement in its students.

For example, you should be able to take a list of topics to be covered from the course outline and be able to use that to determine which sections of the text book are relevant and which questions from old exams are applicable.

Simplicity and elegance are unpopular because they require hard work and discipline to achieve and education to be appreciated.

— Edsger W. Dijkstra, 1997

Fools ignore complexity; pragmatists suffer it; experts avoid it; geniuses remove it.

— Alan Perlis

THE ABILITY TO ACCURATELY ASSESS THE CREDIBILITY OF A SOURCE.

Information comes from a variety of sources. Some of them are more credible than others. An educated person knows, for example, that a report from the Transportation Safety Board of Canada is more likely to be accurate than a newspaper report — more likely, but not infallibly so.

There are two ways of constructing a software design. One way is to make it so simple that there are obviously no deficiencies. And the other way is to make it so complicated that there are no obvious deficiencies.

— C.A.R. Hoare, 1982
Turing Award Speech

THE ABILITY TO MANAGE LARGE AMOUNTS OF COMPLEX, INTER-CONNECTED DETAILS. The human body is a complex system that doctors work with. The law is a complex system that lawyers work with. Engineers work with complex socio-technical systems.

In the first two years of engineering education you typically encounter only small text book problems in order to gain understanding of specific technical concepts in isolation. The labs for this course might be the first time in your engineering education where you have had to work with a non-toy system. Our skeleton code is still small by industrial standards, but it might be an order of magnitude greater than what you have worked with in your previous courses. Learning to manage this volume of complexity and interdependency is an essential part of your professional education.

The total code base that we work with in this course is about 8,000 lines. We provide about 7,000 lines of that to you — which you are expected to understand while developing the remaining 1,000 lines.

For reference, the source code for Microsoft's Windows operating system is somewhere around 25 *million* lines of code. The source code for ATI's video card driver is about 60 *million lines* — more than double the Windows operating system. Modern chip designs also get into millions of lines of code.

THE ABILITY TO THINK DEEP THOUGHTS. In elementary school you learned to multiply positive integers. In middle school you learned to multiply all integers (including negative ones). In high school you learned to multiply matrices, perhaps in two different ways: cross-product and dot-product. In first year you learned to write programs that multiply matrices. In this course you will learn a bit about how to design and implement programming languages in which someone might write a program to multiply matrices. You could go on in pure math and logic courses to learn more about multiplication. It's still the multiplication of your childhood, but your understanding of it is much deeper for the twenty years of study that you have devoted to it. Understanding one or a few things in some depth like this hopefully cultivates in you the ability to think deeper thoughts in other domains and circumstances.

Back where I come from we have universities — seats of great learning — where men go to become great thinkers. And when they come out, they think deep thoughts, and with no more brains than you have.

— The Wizard of Oz, 1939
Of course, in the modern world, women now go to (and graduate from) universities in greater numbers than men.

For example, Gödel's first incompleteness theorem shows that any formal logic that includes multiplication can express propositions that we know to be true but which cannot be proven within that logic. A result that rocked the world in 1931, and was an important intellectual precursor to Turing's proof of The Halting Problem in 1936.

Lab 1

Recursive Descent Parsing of \mathcal{W}

We consider that the input and outputs of a circuit are a set of waveforms. We use the \mathcal{W} language for expressing waveforms in text files. Figure 1.1 shows an example \mathcal{W} file and Figure 1.2 gives the grammar for \mathcal{W} .

```
A: 0 1 0 1 0 1 ;
B: 1 0 1 0 1 0 ;
OR: 1 1 1 1 1 1 ;
```

```
Program    → (Waveform)+
Waveform  → Id ':' Bits ';'
Id         → Char ( Char | Digit | '_' )*
Char       → [A-Za-z]
Digit      → [0-9]
Bits       → ('0' | '1')+
```

1.1 Grammars & EBNF [†]

Figure 1.2 lists the *grammar* for the language \mathcal{W} in *Extended Backus-Naur Form* (EBNF). A grammar is formal specification of the *syntax* of a language. A grammar tells us which *sentences* are included in the language and which are excluded. For certain classes of grammars, we can systematically derive a program to *recognize* and *parse* sentences written in the language.

The grammar in Figure 1.2 also includes the *lexical* specification of the language \mathcal{W} : that is, it also tells us what sequences of characters make legal *tokens* (words). For example, it tells us that an *identifier* (i.e., a name) must start with an alphabetic character, but may contain numerals or underscores in subsequent positions. Some times grammars do not explicitly include a lexical specification, and instead assume that characters not separated by *whitespace* form legal tokens.

Compiler Concepts: regular languages, regular expressions, EBNF, recognizers, parsing, recursive descent, lexing, pretty-printing, abstract syntax tree (AST)

Programming Concepts: classes, objects, variables, aliasing, immutability, test-driven development

Figure 1.1: Example waveform file for an OR gate. The input pins are named A and B, and the output pin is named OR. Our VHDL simulator will read a \mathcal{W} file with lines A and B and will produce a \mathcal{W} file with all three lines.

Figure 1.2: Grammar for \mathcal{W} in EBNF. By convention we will call the top production *Program*.

Crafting: §4.3

BNF was developed by John Backus while working on what became the Algol-60 language. Backus lead the FORTRAN compiler team at IBM in the 1950s, which was the first commercially successful compiler.

People do not usually have numerals in their names. A notable exception is former New York Times reporter *Jennifer 8. Lee*. She says that many computer programs prohibit her middle name, forcing her to write out 'Eight'. She chose the middle initial 8. because as a teenager she realized that there were about 10,000 other 'Jennifer Lee's in the United States.

There are a number of important concepts and notations for understanding grammars and EBNF:

Production Rules: Each line with an arrow (\rightarrow) represents a production rule. A rule has a *right-hand side* (RHS) and a *left-hand side* (LHS). We say the LHS *derives* the RHS, by which we mean that the RHS can be substituted for the LHS. The LHS will always be a single *non-terminal*, whereas the RHS can be some combination of *terminals* and non-terminals.

Terminals: A terminal is a symbol that cannot be derived any further.

Non-terminals: A non-terminal is a symbol that can be derived. Every non-terminal must appear on the LHS of at least one production.

Alternation: The bar ($|$) character indicates alternatives.

Repetition: The star ($*$) and plus ($+$) characters are used to indicate repetition. Star means zero or more, whereas plus means one or more. The plus, star and bar are part of *Extended Backus Naur Form* (EBNF), but are not part of regular Backus Naur Form (BNF). We will learn how to convert from EBNF to BNF in [N2.71].

Derivation: Consider the input \mathcal{W} program A: 1 0. Let's derive this string from the \mathcal{W} grammar in Figure 1.2.

<i>Program</i>	\rightarrow	$(\text{Waveform})^+$	top-level rule
	\rightarrow	$(\text{Id} \text{ ':' Bits ';' })^+$	derive <i>Waveform</i>
	\rightarrow	$(\text{'A' ':' Bits ';' })^+$	derive <i>Id</i>
	\rightarrow	$(\text{'A' ':' '1' '0' ';' })^+$	derive <i>Bits</i>
	\rightarrow	$\text{'A' ':' '1' '0' ';'}$	reached end of input

The *recursive descent* recognizer and parser that you will write in this lab will work in this way: it will descend from the top of the grammar, recursively if necessary (the grammar of \mathcal{W} is not recursive, but the grammar of \mathcal{F} in LAB3 is).

Examples are given with respect to the grammar for the \mathcal{W} language in Figure 1.2.

The name 'production rule' is sometimes shortened to just 'production' or 'rule'.

e.g., colon (:), semi-colon (;), zero (0), one (1), and underscore (_)

e.g., *Program*, *Waveform*, *Id*, *Char*, *Digit*, and *Bits*

e.g., $\text{Bit} \rightarrow \text{'0' } | \text{'1'}$

says that a *Bit* can be a zero or a one.

e.g., $\text{Program} \rightarrow \text{Waveform}^+$

says that a *Program* derives one or more *Waveforms*.

In this example we are doing a *leftmost* derivation: that is, we are expanding the leftmost non-terminal at each step.

If we expanded *Bits* before *Id* then it would be a *rightmost* derivation.

Pragmatics: §2.1.3, p.48

Crafting: §4.1.1, §4.1.2

1.1.1 Summary of Grammar Notation Conventions

When writing grammars in EBNF or BNF we will often follow these notational conventions for names.

ABC	non-terminals
abc	terminals
XYZ	nonterminals or terminals
xyz	token strings
$\alpha\beta\gamma$	strings of arbitrary symbols
LHS	Left Hand Side
RHS	Right Hand Side

This chart will be repeated elsewhere in the Lab Manual and Course Notes as convenient. The conventions are adopted from *Programming Language Pragmatics* by Michael L. Scott.

Pragmatics: §2

Crafting: §4.5.1

1.1.2 Derivations[†]

A *derivation* is how we show, on paper, that an input string can be recognized by a grammar. We start with the top production in the grammar and do substitutions until we derive the input string. In most textbooks this is done algebraically, but it can also be done graphically by drawing the partial AST at each step of the derivation. Here we will show both ways.

Pragmatics: §2.1.3, p.48 + J1

Tiger: J1


Crafting: §4.1.1, §4.1.2

Consider the input string '1+2+3' and the following grammar, with productions numbered for easy reference:


1. $S \rightarrow E$
2. $E \rightarrow E + E$
3. $E \rightarrow INT$

Here is a derivation of that input string with that grammar:


$S \rightarrow E$ start at the top




$\rightarrow (E + E)$ substitute by rule 2



$\rightarrow ((E + E) + E)$ substitute by rule 2




$\rightarrow ((1 + 2) + 3)$ substitute by rule 3 (3 times)




Here is another derivation of that input string with that grammar:

 $S \rightarrow E$


start at the top


 $\rightarrow (E + E)$


substitute by rule 2


 $\rightarrow (E + (E + E))$

substitute by rule 2


 $\rightarrow (1 + (2 + 3))$

substitute by rule 3 (3 times)



Since there are multiple derivations (*i.e.*, parse trees) for the same input string, we say that this grammar is *ambiguous*.

[N2.4]

1.2 Write a regular expression recognizer for \mathcal{W}

A *recognizer* is a program that *accepts* or *rejects* a string input based on whether that string is a valid sentence in some language. A recognizer for \mathcal{W} will accept or reject a file if that file is a legal \mathcal{W} ‘sentence’/‘program’.

\mathcal{W} is a *regular* language. *Regular* is a technical term here that describes the complexity of the grammar of \mathcal{W} . Regular languages are the simplest kind of languages that we will consider. The grammar of a regular language can be described by a *regular expression*.

Your first task is to write a regular expression describing the grammar of \mathcal{W} . You will do this in two steps, first editing `TestWRegexSimpleData` and running `TestWRegexSimple`, and second editing `TestWRegexAccept.REGEX` and running both `TestWRegexAccept` and `TestWRegexReject`. Your final answer should be stored in the field `TestWRegexAccept.REGEX`.

Sources:

`ece351.w.regex.TestWRegexSimple`
`ece351.w.regex.TestWRegexSimpleData`
`ece351.w.regex.TestWRegexAccept`
`ece351.w.regex.TestWRegexReject`

The website regexper.com will draw finite-state machine diagrams for your regular expressions. The website regexpal.com lets you test out your regular expressions. There are many other similar websites.

1.2.1 Test-Driven Development[†]

In traditional *waterfall* software development, one writes the code first and then the tests; moreover, the code is written in a *deductive* mental style, thinking about the problem generally.

Modern *agile* software development advocates a different approach: *Test-Driven Development (TDD)*. In TDD, one writes the tests first and then the code. The code is written in a more *inductive* mental style: what needs to be done to pass the next test case?

The first exercise is to follow a Test-Driven Development approach to building a regular expression recognizer for the \mathcal{W} language. There are test files `r1.wave` through `r8.wave`. You will start by running `TestWRegexSimple` and observing that the first test passes and the second test fails. The first test passes because the regular expression you are given, in `TestWRegexSimpleData`, matches the contents of `r1.wave`: this regular expression will work only for `r1.wave`.

Your job is to copy/paste/generalize that regular expression, making the minimal modifications necessary for it to accept both `r1.wave` and `r2.wave`. Then you will see that it fails `r3.wave`. Again, copy/paste/generalize, and you will see it fails `r4.wave`. Eventually you will build up a regular expression that will accept up to `r8.wave`.

A challenge that you will face is that the \mathcal{W} grammar in Figure 1.2 does not explicitly specify whitespace, whereas your regular expression will have to explicitly specify the whitespace. Grammars given in EBNF usually implicitly assume some *lexical* specification that describes how characters in the input string are to be grouped together to form *tokens*. Usually any amount of whitespace may occur between tokens.

1.2.2 Complete the regular expression recognizer for \mathcal{W}

Copy your final regular expression recognizer for \mathcal{W} from `TestWRegexSimpleData` to `TestWRegexAccept.REGEX`, and then run `TestWRegexAccept` and `TestWRegexReject`. You will see that some tests pass and some tests fail. Fix your regular expression until all of the tests pass.

1.3 Write a recursive descent recognizer for \mathcal{W} [†]

Recursive descent is a style of writing parsers (or recognizers) by hand (*i.e.*, without using a parser generator tool). In this style we make a method for each non-terminal in the grammar. For recognizers these methods return `void`.

The bodies of these non-terminal methods consume the input one token at a time. Kleene stars (*) or plusses (+) become loops.

Sources:

```
ece351.w.regex.TestWRegexSimple
ece351.w.regex.TestWRegexSimpleData
```

The danger of the traditional approach is that it produces general code that does not actually pass any specific tests. Obviously if the test-suite is inadequate then TDD will result in code that is not sufficiently general. But it will pass some test cases.

Sources:

```
ece351.w.regex.TestWRegexAccept
ece351.w.regex.TestWRegexReject
```

Sources:

```
ece351.w.rdescent.WRecursiveDescentRecognizer
```

Libraries:

```
ece351.util.Lexer
```

Tests:

```
ece351.w.rdescent.TestWRDRecognizerAccept
ece351.w.rdescent.TestWRDRecognizerReject
```

Alternation bars (`|`) become conditionals (if statements).

Execution of these non-terminal methods starts at the ‘top’ of the grammar. This is what the term *descent* in *recursive descent* refers to. The term *recursive* in *recursive descent* simply means that the various non-terminal methods may call each other.

By convention we will name the top production of our grammars *Program*, and so execution of our hand-written recursive descent parsers and recognizers will start in a method called `program`.

Your \mathcal{W} recognizer code will make use of the Lexer library class that we provide. This class *tokenizes* the input string: *i.e.*, strips out the whitespace and groups the input characters into chunks called *tokens*. The lexer provides a number of convenience methods, including methods for recognizing identifiers, so you won’t need to write an `id()` method in your recognizer.

The lexer has two main kinds of operations: *inspect* the next token to see what it looks like, and *consume* the next token. The *inspect* methods are commonly used in the tests of loops and conditionals. The *consume* methods are commonly used in regular statements.

1.4 Write a pretty-printer for \mathcal{W}

Pretty-printing is the opposite of parsing. Parsing is the process of constructing an *abstract syntax tree* (AST) from a string input. Pretty-printing is the process of producing a string from an AST.

For this step you will write the `toString()` methods for the \mathcal{W} AST classes. These methods will be tested in the next step. Note that these `toString()` methods just return a string; they do not actually print to the console nor to a file. *Pretty-printing* is the name for the inverse of parsing, and does not necessarily involve actually printing the result to an output stream. Whereas parsing constructs a tree from a string, pretty-printing produces a string from a tree.

1.5 Write a recursive descent parser for \mathcal{W}

A parser reads a string input and constructs a tree output. Specifically, an *abstract syntax tree* (AST).

To write your recursive descent parser for \mathcal{W} start by copying over the code from your recursive descent recognizer. The parser will be a superset of this code. The difference is that the recognizer discards the input whereas the parser will build up the AST based on the input. The AST classes are `WProgram` and `Waveform`.

The test performed here is to parse the input, pretty-print it, re-parse the pretty-printed output, and then compare the AST’s from the two parses to see that they are the same.

Sources:

`ece351.w.ast.WProgram`
`ece351.w.ast.Waveform`

Libraries:

`java.lang.String`
`System.lineSeparator()`
`org.parboiled.common.ImmutableList`

Once we have a function and its inverse we can test that $f'(f(x)) = x$ for any input x .

Sources:

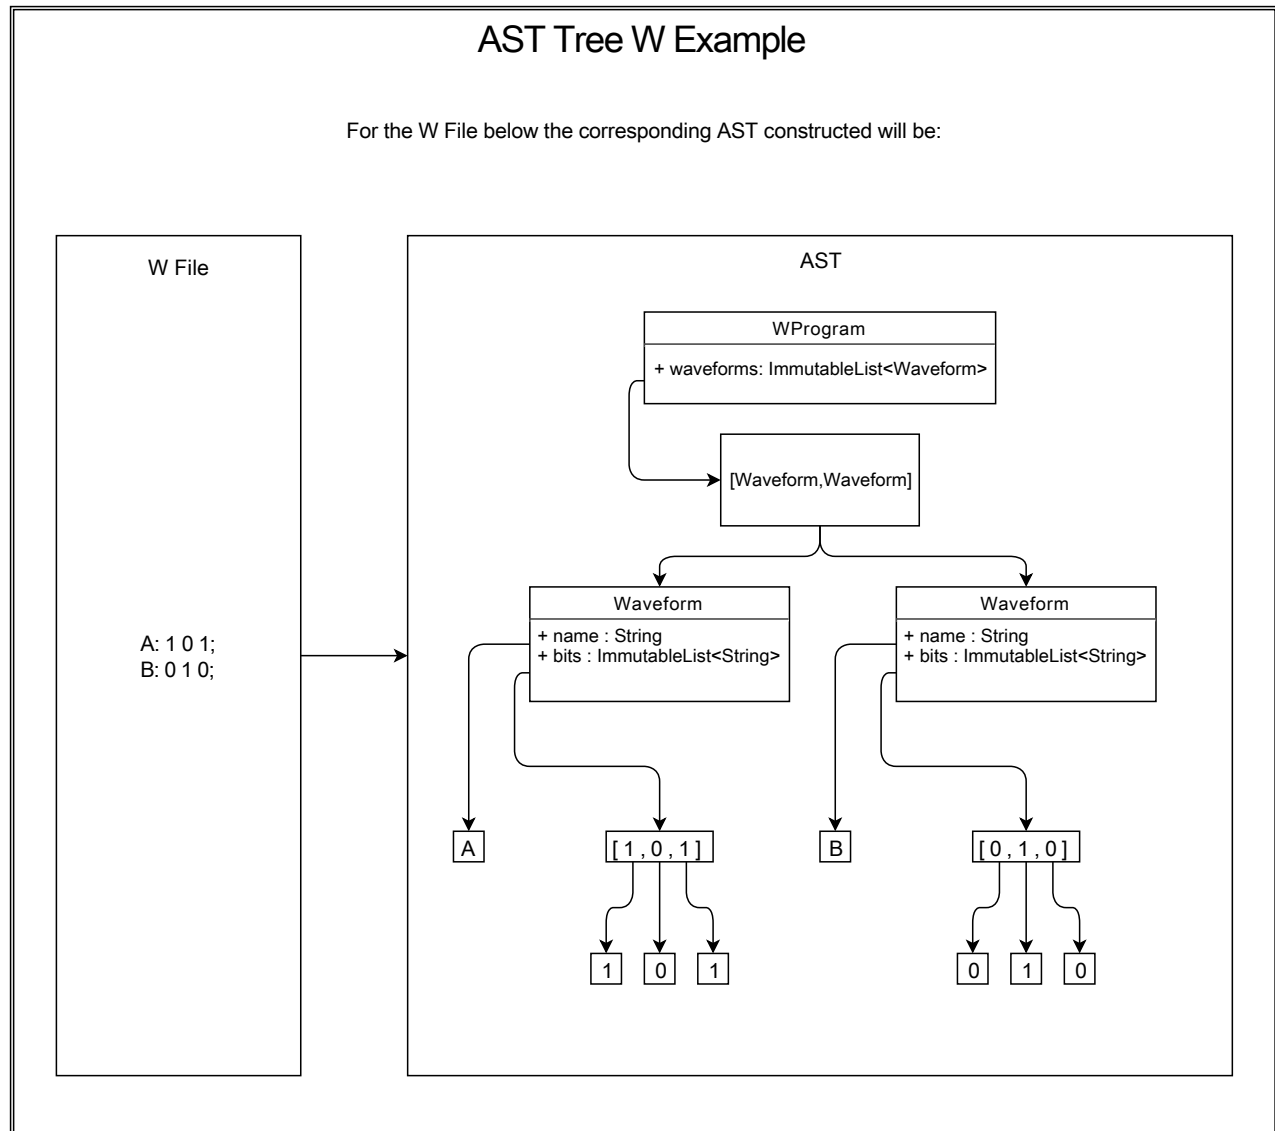
`ece351.w.rdescent.WRecursiveDescentParser`

Libraries:

`ece351.w.ast.WProgram`
`ece351.w.ast.Waveform`
`ece351.util.Lexer`
`org.parboiled.common.ImmutableList`

Tests:

`ece351.w.rdescent.TestWRDParserBasic`
`ece351.w.rdescent.TestWRDParserAccept`
`ece351.w.rdescent.TestWRDParserReject`

1.6 Object Diagram of Example \mathcal{W}_{AST} Figure 1.3: Object diagram for example \mathcal{W}_{AST}

1.7 Steps to success

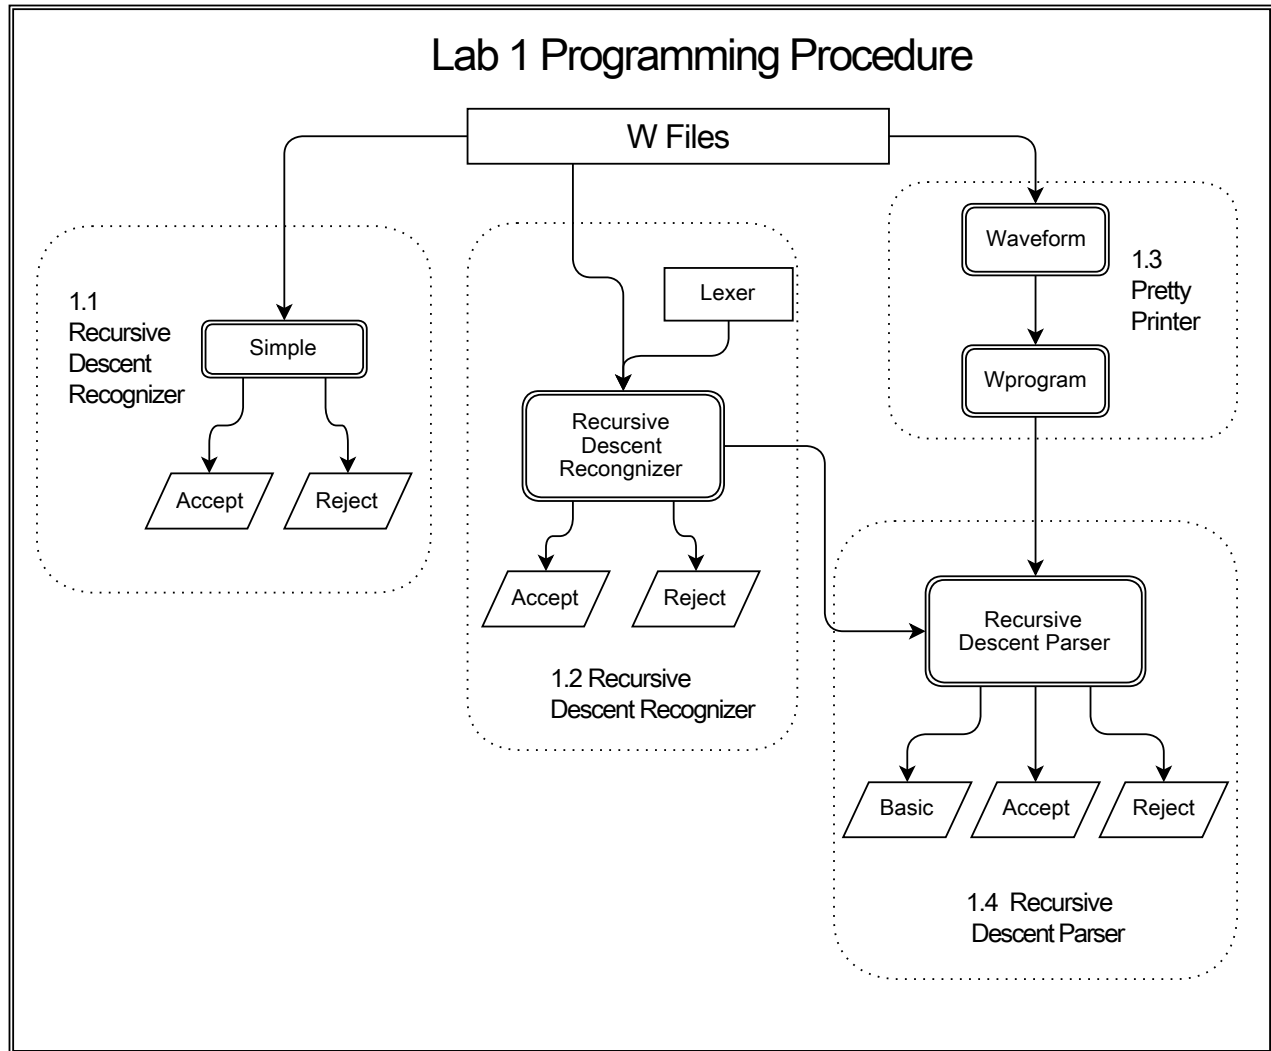


Figure 1.4: Steps to success for §1.
 Legend for icon meanings in Figure 1.5.

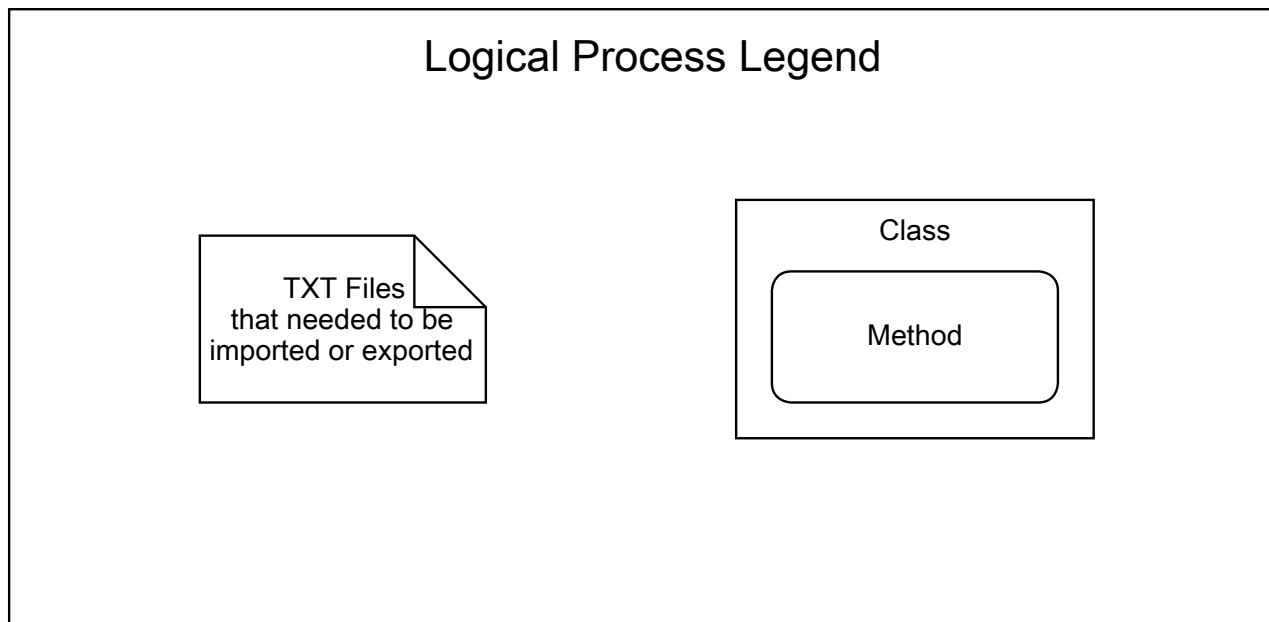
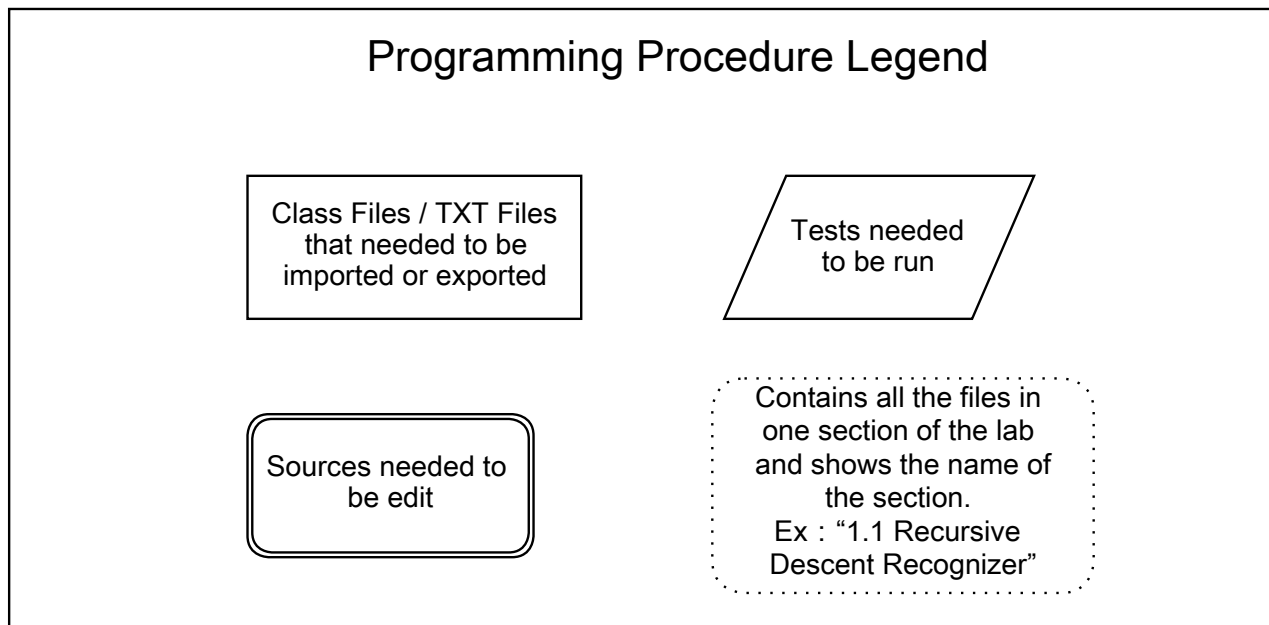


Figure 1.5: Legend for Figure 1.4

1.8 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

	Shared	Secret
TestWRegexAccept	10	5
TestWRegexReject	5	5
TestWRDRecognizerAccept	15	5
TestWRDRecognizerReject	5	5
TestWRDParserBasic	5	0
TestWRDParserAccept	20	10
TestWRDParserReject	5	5

Note that you don't earn any points for the rejection tests until some of the corresponding acceptance tests pass. Also, you don't earn any points for the parser until the TestWRDParserBasic tests pass.

1.9 Reading

1.9.1 Tiger Book¹

- 1 Introduction
- 2.0 Lexical Analysis
- 2.1 Lexical Tokens
- 2.2 Regular Expressions
- 3.0 Parsing
- 3.2.0 Predictive Parsing / Recursive Descent
 - skip First and Follow Sets for now
 - read Constructing a Predictive Parser on p.50
 - skip Eliminating Left Recursion and everything that comes after it — for now
- 4.1.1 Recursive Descent

¹ A. W. Appel and J. Palsberg. *Modern Compiler Implementation in Java*. Cambridge, 2004

1.9.2 Programming Language Pragmatics²

- 2.2 Scanning
- 2.3.1 Recursive Descent

² M. L. Scott. *Programming Language Pragmatics*. Morgan Kaufmann, 3 edition, 2009

1.9.3 Web Resources

*Thinking in Java*³

³ B. Eckel. *Thinking in Java*. Prentice-Hall, 2002. <http://www.mindview.net/Books/TIJ/>

Compiler Concepts: trees, transformations, XML, SVG

Programming Concepts: object contract, object equality, mathematical equivalence classes, DOM vs. SAX parser styles, call-backs, iterator design pattern


Lab 2

Transforming $\mathcal{W} \rightarrow \text{SVG}$ for Visualization

While our circuit simulator programs will read and write \mathcal{W} files, people often prefer to look at waveforms in a graphical format. In this lab we will translate \mathcal{W} files into SVG files for visualization.

An example \mathcal{W} file is shown in Figure 2.1. Figure 2.2 shows what the corresponding SVG looks like when visualized with a web browser such as Firefox or a vector graphics program such as Inkscape. Figure 2.3 shows the first few lines of the textual content of the SVG file.

```
A: 0 1 0 1 0 1 ;  
B: 1 0 1 0 1 0 ;  
OR: 1 1 1 1 1 1 ;
```



SVG is an XML-based graphics file format: in other words, it is structured text that describes some vectors to be rendered. Any modern web browser will render SVG.

Figure 2.1: Example waveform file for an OR gate. The input pins are named A and B, and the output pin is named OR. Our VHDL simulator will read a \mathcal{W} file with lines A and B and will produce a \mathcal{W} file with all three lines.

Figure 2.2: Rendered SVG of \mathcal{W} file from Figure 2.1

```

<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE svg PUBLIC "-//W3C//DTD SVG 1.1//EN" "http://www.w3.org/
  Graphics/SVG/1.1/DTD/svg11.dtd">
<svg xmlns="http://www.w3.org/2000/svg" width="100%" height="100%"
  version="1.1">
<style type="text/css"><![CDATA[|line{stroke:#006600;fill:#00 cc00;} text{font-
  size:"large";font-family:"sans-serif"}]]></style>

<text x="50" y="150">A</text>
<line x1="100" x2="100" y1="150" y2="200" />
<line x1="100" x2="200" y1="200" y2="200" />
<line x1="200" x2="200" y1="200" y2="100" />
<line x1="200" x2="300" y1="100" y2="100" />

```

2.1 Write a translator from $\mathcal{W} \rightarrow \text{SVG}$

The idea of the transformation is simple: for each bit (zero or one) in the input file, produce a vertical line and a horizontal line. To do this the transformer needs to remember three things: the current X and Y position of the (conceptual) cursor, and the YOffset so we can move the cursor down the page when we start drawing the next waveform.

Run your transformations and inspect the output both visually, in a program such as Firefox, and textually, using the text editor of your choice. Ensure that the output is sensible.

2.2 Write a translator from $\text{SVG} \rightarrow \mathcal{W}$

This translator is the inverse of the previous one. We write this inverse function as a way to test the previous translator. Now we can take advantage of the general property that $x = f'(f(x))$ or, in other words, $w = \text{toW}(\text{toSVG}(w))$.

This inverse translation is a bit tricky because the SVG files do not contain any explicit information about which line segments are associated with which waveform: the SVG file just contains a bunch of text labels and a bunch of line segments. We have to infer, from the y values, which line segments belong to which waveform/label. Then we use the x values to infer the ordering of the bits, and the y values to infer the bit values (0 or 1).

Many program analysis tasks used in optimizing compilers involve this kind of inference: recovering information that is implicit in a lower-level representation but was explicit in some higher-level representation.

Figure 2.3: First 10 lines of SVG text of \mathcal{W} file from Figure 2.1. The boilerplate at the top is already in W2SVG.java for your convenience.

Note that the origin of an SVG canvas is the top left corner, not the bottom left corner (as you might expect from math). Consequently, $Y=200$ is visually lower than $Y=100$, since Y counts down from the top.

Sources:

ece351.w.svg.TransformW2SVG

Libraries:

ece351.w.ast.WProgram

ece351.w.svg.Line

ece351.w.svg.Pin

Tests:

ece351.w.svg.TestLegalSVG

Exercise 2.1.1 Where is the origin on an SVG canvas? Which directions are positive and which are negative?

Sources:

ece351.w.svg.TransformSVG2W

Libraries:

ece351.w.ast.WProgram

ece351.w.svg.Line

ece351.w.svg.Pin

Tests:

ece351.w.svg.TestW2SVG2W

2.3 Introducing the Object Contract[†]

The *object contract* is a term sometimes used to describe properties that all objects should have. It is particularly concerned with the *equals* and *hashCode* methods. The *equals* method should represent a *mathematical equivalence relation* and it should be consistent with the *hashCode* method. A mathematical equivalence class has three properties:

reflexive `x.equals(x)`
symmetric `x.equals(y) ⇔ y.equals(x)`
transitive `x.equals(y) && y.equals(z) ⇒ x.equals(z)`

Additionally, by consistent with *hashCode* we mean the following:

consistent `x.equals(y) ⇒ x.hashCode() == y.hashCode()`

Finally, no object should be equal to null:

not null `!x.equals(null)`

The file `TestObjectContract` contains some code stubs for testing the object contract. In order to fill in those stubs correctly you might need to read and understand the code in `TestObjectContractBase`.

What is the default behaviour of the *equals* and *hashCode* methods? In other words, what happens if you do not implement those methods for your objects? Consider the code listings in Figure 2.4. What value is printed out by each statement?

Exercise 2.3.3

```
Double x = new Double(1.5);
Double y = new Double(1.5);
System.out.println(x.equals(y));
System.out.println(x == y);
```

Exercise 2.3.4

```
List a = new LinkedList();
List b = new LinkedList();
System.out.println(a.equals(b));
System.out.println(a == b);
```

Sources:

`ece351.objectcontract.TestObjectContract`

Libraries:

`ece351.objectcontract.TestObjectContractBase`

Tests:

`ece351.objectcontract.TestObjectContract`

The object contract looks simple (and it is), but even expert programmers have difficulty getting it right in all circumstances.

Exercise 2.3.1 What would happen if every *hashCode* method returned 42?

Exercise 2.3.2 What would happen if every *hashCode* method returned a random value each time it was called?

The `==` ('double equals' or 'equals equals') comparison operator compares the memory addresses of the objects referred to by the two variables.

Exercise 2.3.5

```
Object p = new Object();
Object q = new Object();
System.out.println(p.equals(q));
System.out.println(p == q);
```

2.3.1 Engineering Apparatus and Controls[†]

What is a kilogram? Since 1983, a metre has been defined as the length of the path travelled by light in vacuum during a time interval of $\frac{1}{299,792,458}$ seconds.¹ What is a second? Since 1967, a second has been defined as the duration of 9,192,631,770 periods of the radiation corresponding to the transition between the two hyperfine levels of the ground state of the caesium 133 atom.² So we have definitions of distance and time in terms of physical constants of nature.

The kilogram, by contrast, is defined by a cylinder (with nicely chamfered edges) of platinum-iridium alloy in a vault in Paris: the International Prototype Kilogram (IPK). It is the last remaining base unit of measure to be defined by a physical artefact rather than a property of nature (Figure 2.5).

Figure 2.4: Code comprehension exercises. Draw an object diagram for each statement in each listing. What will the result of the `println` statements be?

¹ 17th Conférence Générale des Poids et Mesures (CGPM) — Resolution 1 of the CGPM (1983): Definition of the metre. Bureau international des poids et mesures (BIPM). <http://www.bipm.org/en/CGPM/db/17/1/>


² Unit of time (second). SI Brochure. BIPM. <http://www.bipm.org/en/publications/si-brochure/second.html>

The kilogramme prototype

The kilogramme is officially defined by a lump of metal stored in a vault in France for more than 120 years under the International System of Units (SI)

The International prototype

- ▶ Manufactured in 1889
- ▶ Made up of 90% platinum, and 10% iridium
- ▶ Kept in a vault at the International Bureau of Weights and Measures (BIPM), along with six official copies
- ▶ The basis of more than 80 copies distributed around the world as national prototypes
- ▶ Steam-cleaned under strict guidelines on a regular basis
- ▶ The last remaining base unit measured against a specific material artefact



The chunk of metal is under triple lock-and-key in Sevres, France

AFP

Source: BIPM

Figure 2.5: Infographic and photo of the International Kilogram Prototype stored in Paris. From Phys.org <http://phys.org/news/2011-11-quandary-kilo-triggers-weighty-reflexion.html>

Suppose that we want to build a scale — an engineering apparatus to measure mass. The purpose of such a device, in principle, is to compare other masses to the IPK. This comparison has three possible outcomes: *less than*, *more than*, or *same*.

Many developed countries, including Canada, have an official replica of the IPK. One would expect that all of these official replicas have the same mass as the IPK. Suppose that we tested our scale with only these official replicas, and every time the scale reported *same*, as expected. Can we thus conclude that our scale is really accurate? No: perhaps it always returns *same*; maybe it never returns *more* or *less*.

We need some *control* objects that actually have more or less mass than the IPK, so that we can also check that our scale does return *more* or *less* when appropriate.

Part of the software exercise here is to build the apparatus (scale), and part of it is to build these control objects so that we know the apparatus works properly. Then, in the future, we can use the apparatus on unknown objects and trust the results it gives us.

In this analogy, it is obvious that the control objects are supposed to weigh more or less than the IPK. There are many objects in the world that weigh more or less than the IPK. In our software world, no object is ever supposed to violate the object contract, and in practice very few objects violate that contract. So it might feel weird to intentionally create control objects that violate the object contract. Nobody would ever want such an object in a regular program. But we need these control objects to assess the accuracy of our apparatus.

2.3.2 Learning to write professional code[†]

It is important, in general, for you to learn the Object Contract because it is a set of properties that must hold for all code you write in any object-oriented language (Java, C++, C#, Python, *etc.*). It is also specifically important that you learn the Object Contract to do the labs in this course, because checking the equality of objects forms the basis of the automated marking scripts. In this exercise you are also learning how to write professional code, with the Object Contract as an example. The steps that we are following here are:

- a. Identify the mathematical properties that the code should have.
(*Object contract: reflexive, symmetric, transitive, etc.*)
- b. Write methods that check these mathematical properties on some input object(s). (*e.g., checkEqualsIsReflexive(x)*) Call these the *property check methods*.
- c. Verify that the property check methods return true when expected by using third party inputs (*e.g., java.lang.String and java.lang.Integer*).
- d. Verify that the property check methods return false when expected by constructing pathological control objects with known deviant behaviour. (*e.g., constructAlwaysTrue(), constructAlwaysFalse(), constructToggle(), etc.*)
- e. Use the verified property check methods to assess real code.
(*We aren't doing this step in this exercise.*)

In future labs we will use staff-provided property check methods for Object Contract properties. The purpose of this exercise is for you to see what these property check methods look like, both so you understand the Object Contract and so you understand how your future labs will be graded.

2.3.3 Reading

[http://docs.oracle.com/javase/7/docs/api/java/lang/Object.html#equals\(java.lang.Object\)](http://docs.oracle.com/javase/7/docs/api/java/lang/Object.html#equals(java.lang.Object))

<http://www.artima.com/lejava/articles/equality.html>

<http://www.angelikalanger.com/Articles/JavaSolutions/SecretsOfEquals/Equals.html>

Implementation of java.util.AbstractList.equals()

*Effective Java*³

§0.14 of this lab manual on Testing

³ J. Bloch. *Effective Java*. Addison-Wesley, 2001

2.4 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

We provide you with about 30 \mathcal{W} files for these equations.

	<i>Current</i>	<i>New</i>	<i>Equation</i>
TestLegalSVG	20	10	$\text{legalSVG}(\text{TransformW2SVG}(w))$
TestW2SVG2W	40	10	$\text{SVG2W}(\text{W2SVG}(w)).\text{equivalent}(\text{SVG2W}(\text{staff.svg}))$
TestObjectContract	15	5	see §2.3 above

2.5 Steps to success

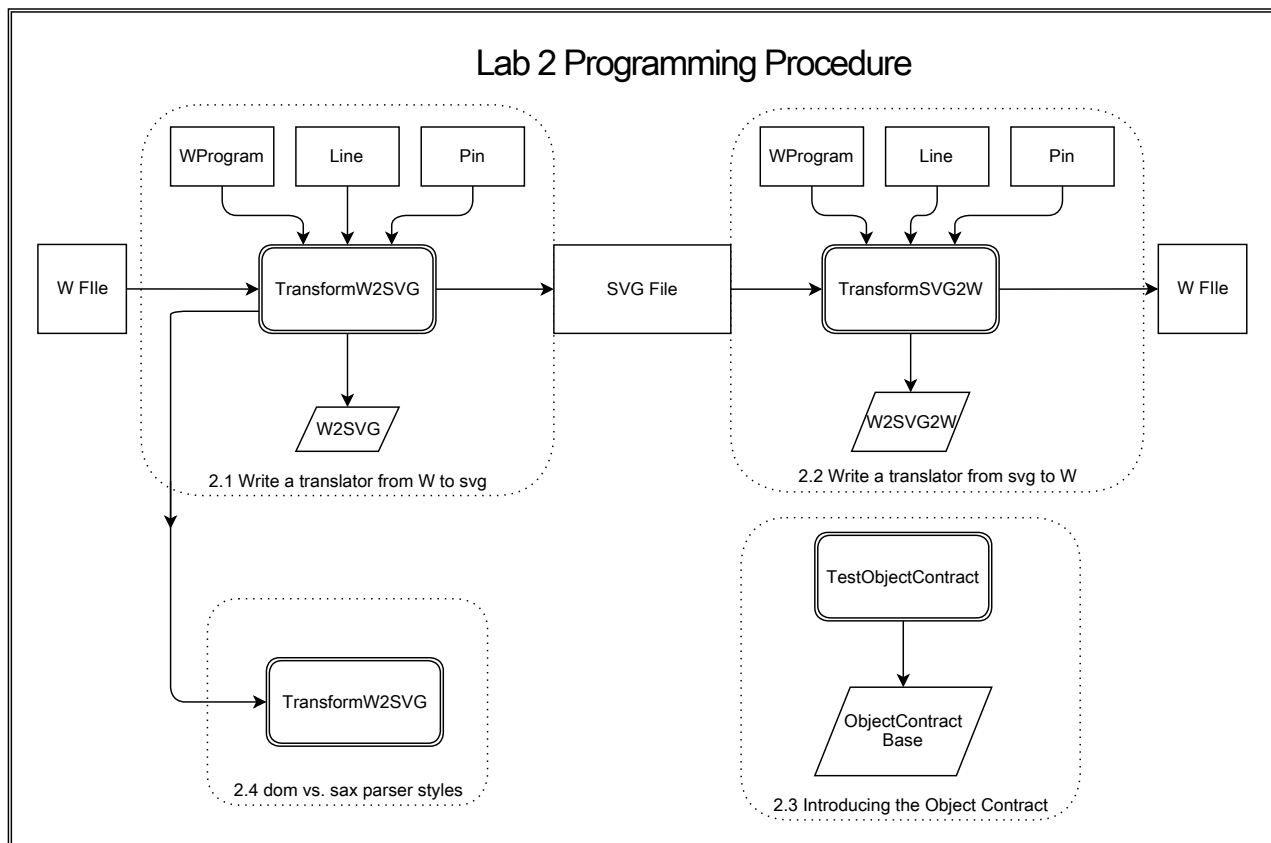


Figure 2.6: Steps to success for §2.
Legend for icon meanings in Figure 1.5.

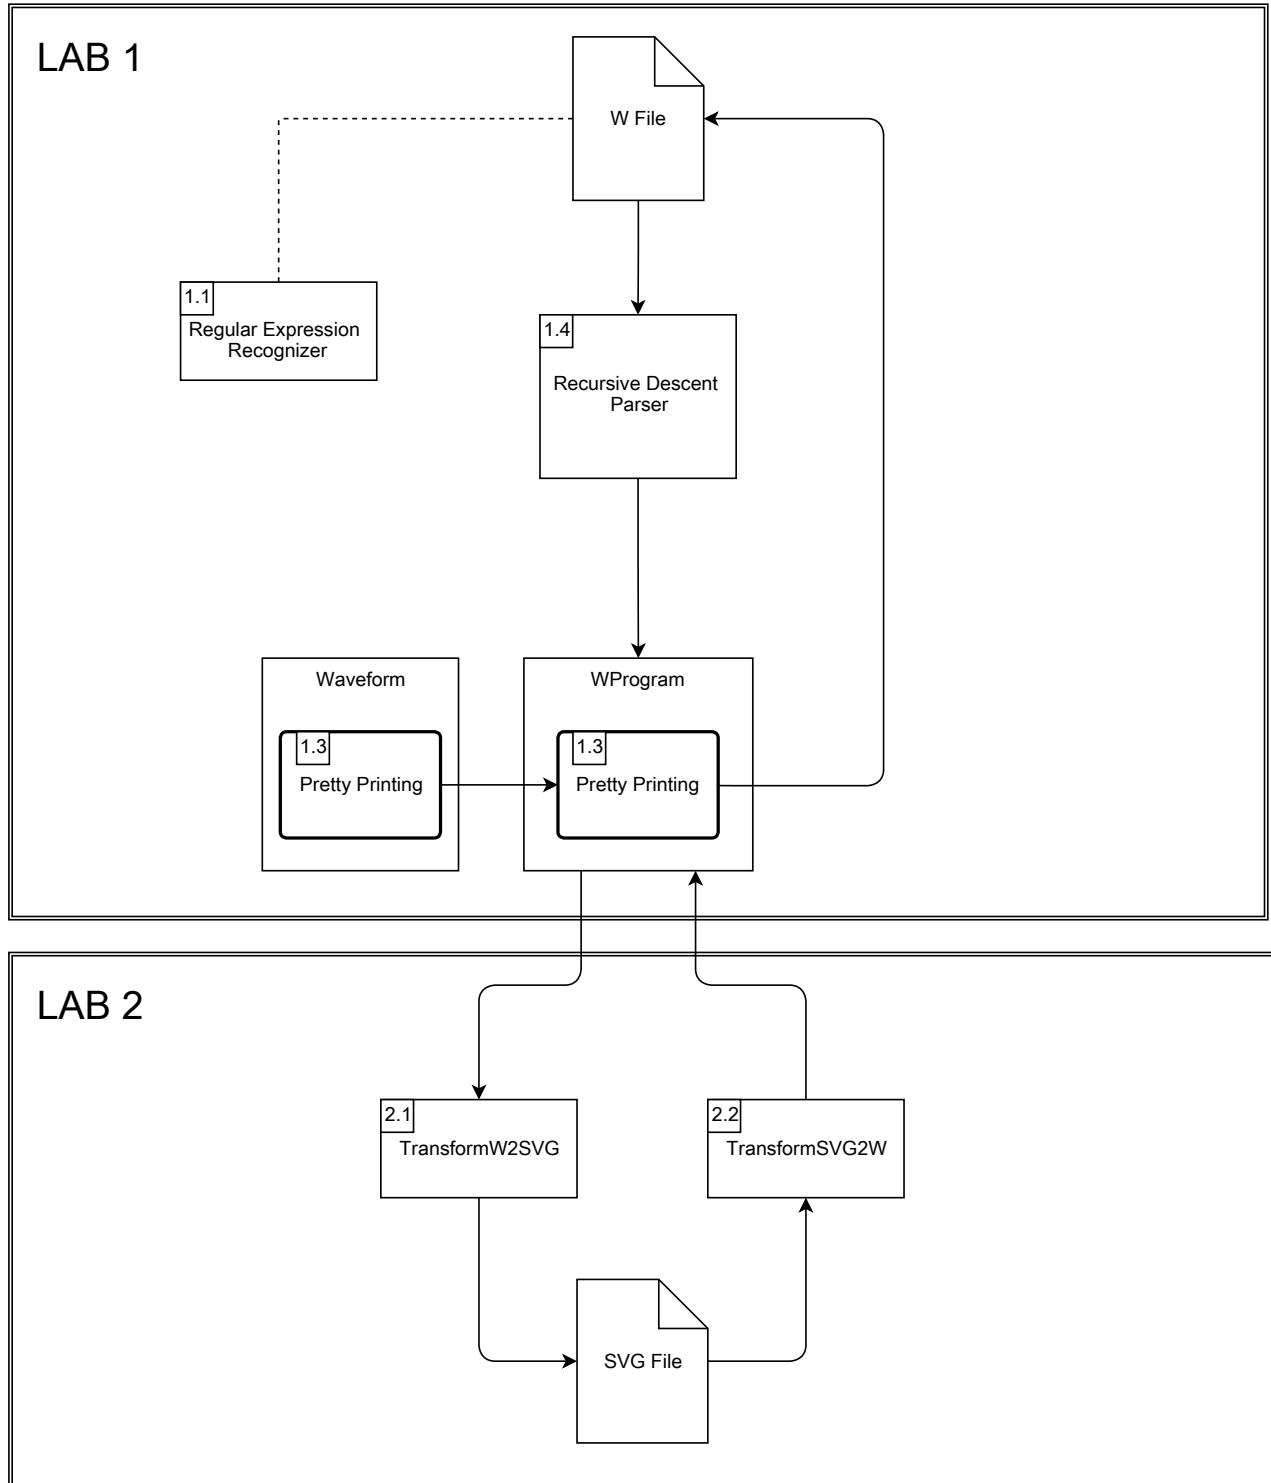


Figure 2.7: How §1 and §2 fit together

Lab 3

Recursive Descent Parsing of \mathcal{F}

Compiler Concepts: context-free grammars, LL(1) grammars, predict sets, parse trees, precedence, associativity, commutativity, program equivalence
Programming Concepts: inheritance, polymorphism, dynamic dispatch, type tests, casting, memory safety, composite design pattern, template design pattern, singleton design pattern, recursive functions, recursive structures, higher-order functions

This lab introduces formula language \mathcal{F} , which we will use as an intermediate language in our circuit synthesis and simulation tool. In a subsequent lab we will write a translator from \mathcal{V} to \mathcal{F} .

Compilers usually perform their optimizations on programs in intermediate forms. These intermediate forms are designed to be easier to work with mechanically, at the cost of being less pleasant for people to write large programs in. A program written in \mathcal{F} , for example, is just a list of boolean formulae. This is relatively easy to manipulate mechanically. \mathcal{V} , by contrast, has conditionals, module structure, *etc.*, which are all of great benefit to the \mathcal{V} programmer but are more work for the compiler writer to manipulate. In the next lab we will write a simplifier/optimizer for \mathcal{F} programs.

<i>Program</i>	\rightarrow	Formula ⁺ \$\$
<i>Formula</i>	\rightarrow	Var ' \leq ' Expr ';'
<i>Expr</i>	\rightarrow	Term ('or' Term)*
<i>Term</i>	\rightarrow	Factor ('and' Factor)*
<i>Factor</i>	\rightarrow	'not' Factor '(' Expr ')' Var Constant
<i>Constant</i>	\rightarrow	'0' '1'
<i>Var</i>	\rightarrow	id

Figure 3.1: LL(1) Grammar for \mathcal{F} . \mathcal{F} is a very simple subset of \mathcal{V} , which is in turn a subset of the real VHDL. \mathcal{F} includes only the concurrent assignment statement and the boolean operators conjunction (and), disjunction (or), and negation (not). Note that in the concrete syntax the constants '0' and '1' are surrounded by single quotes, whereas no other terminals are.

3.1 A Tale of Three Hierarchies[†]

There are (at least) three hierarchies involved in understanding this lab — and all of the future labs. In LAB1 we met the *abstract syntax tree* (AST) and the *parse tree* (the execution trace of a recursive descent recognizer/parser). In this lab we will also meet the *class hierarchy* of our Java code. In the past labs there wasn't much interesting in the class hierarchy, but now (and for the remainder of the term) it is a central concern.

3.1.1 Parse Tree

Consider the \mathcal{F} program $X \leq A \text{ or } B;$. When we execute the recursive descent recognizer that we are about to write its call tree will look something like this:

```

program()
  formula()
    var()
      id()
        'X'
      '<='
    expr()
      term()
        factor()
          var()
            id()
              'A'
          'or'
        term()
          factor()
            var()
              id()
                'B'
          ';;'
      EOI

```

3.1.2 AST

Recall that 'AST' stands for *abstract syntax tree*, and is the important information that we want to remember from the parse tree above. This important information is the structure of the tree and the interesting nodes. Things like term, factor, and parentheses are ways that the input string communicates its structure to the parser: the AST doesn't need to retain those things, just the structure itself.

FProgram

```

AssignmentStatement
  outputVar = X
  expr = OrExpr
    left = A
    right = B

```

Exercise 3.1.1 Try to draw a few parse trees on paper for other examples such as:

- $X \leq A \text{ and } B;$
- $X \leq A \text{ or } B \text{ and } C;$
- $X \leq A \text{ and } B \text{ or } C;$

We will use EOI (end of input) and EOF (end of file) and \$\$ interchangeably.

See the readings, especially Bruce Eckel's free online book *Thinking in Java*, for more background material on inheritance/sub-classing.

Exercise 3.1.2 For the examples above, also draw the corresponding AST. Both the call tree and the resulting AST depend on the input.

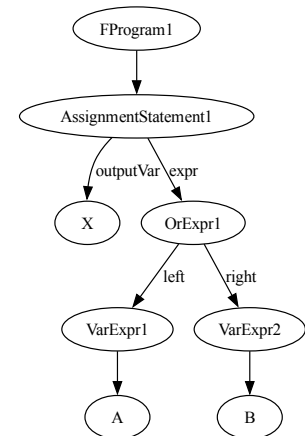


Figure 3.2: Object diagram of AST for \mathcal{F} program $X \leq A \text{ or } B;$

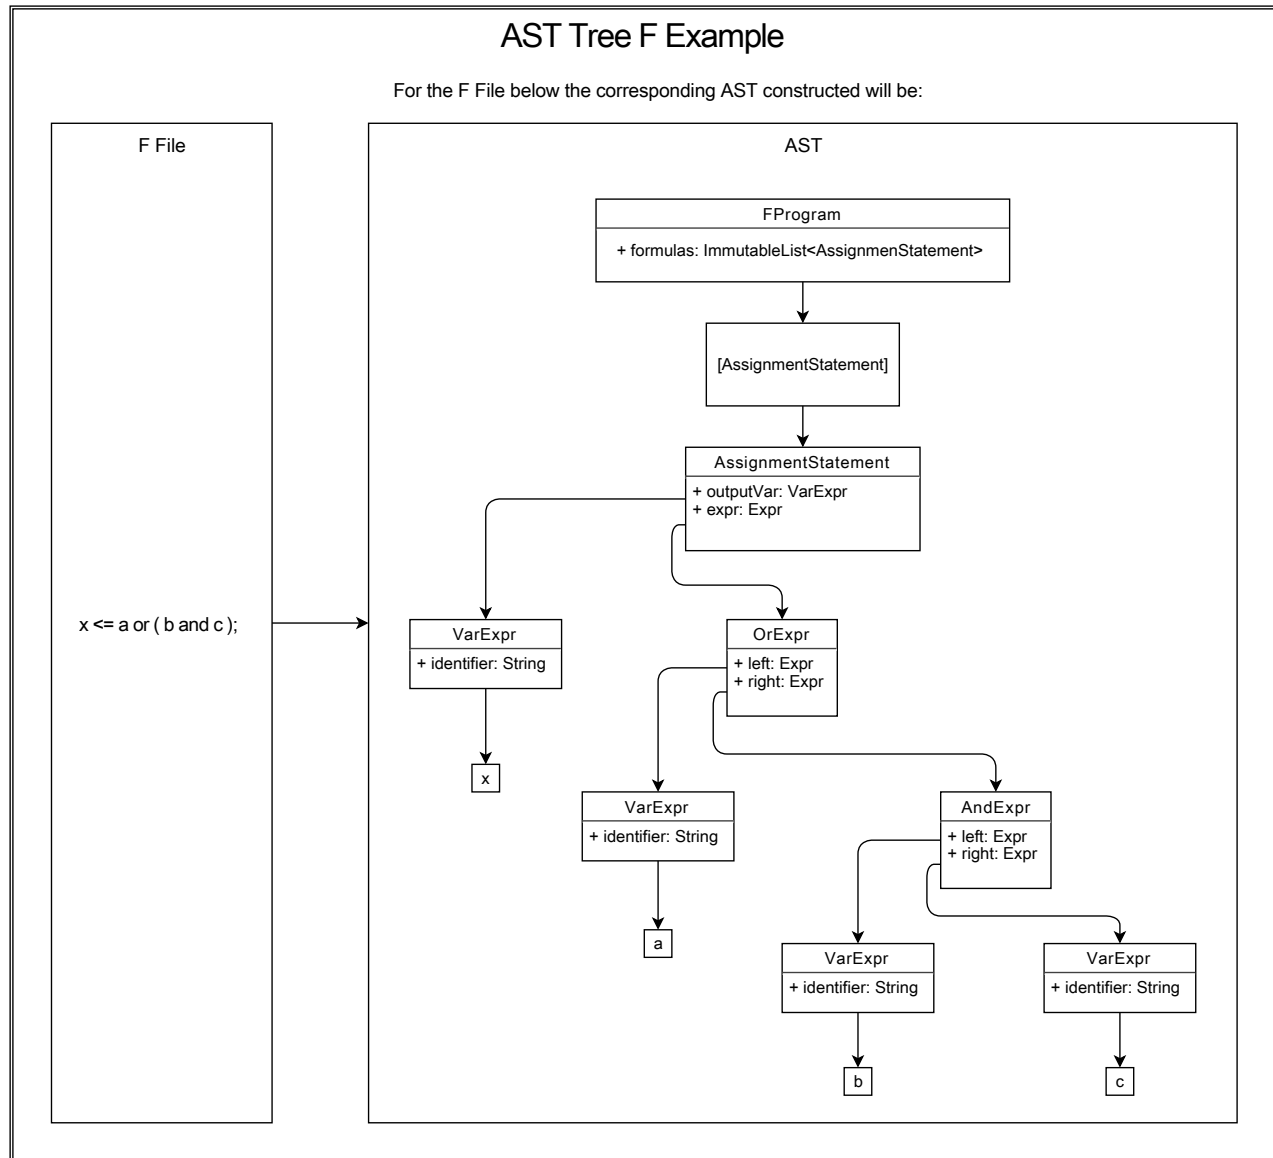


Figure 3.3: Object diagram for example \mathcal{F}_{AST}

3.1.3 A peak at the Expr class hierarchy

The class hierarchy, on the other hand, does not depend on the input: it is how the code is organized. The following listing highlights some features of the Expr class hierarchy and also demonstrates polymorphism/dynamic-dispatch.

Exercise 3.1.3 Find the class `common.ast.Expr` in Eclipse, right-click on it and select *Open Type Hierarchy*. This will give you an interactive view of the main class hierarchy.

Exercise 3.1.4 Draw this hierarchy in a UML class diagram, excluding the classes in `common.ast` that you do not need for this lab.

For this lab you will need the classes `AndExpr`, `OrExpr`, `NotExpr`, `VarExpr`, and `ConstantExpr`. You will also need the classes that those depend on, but you will not need classes corresponding to more esoteric logical operators such as exclusive-or. This lab also uses classes outside of the Expr class hierarchy, but you don't need to draw them on your UML class diagram.

Figure 3.4: Some highlights from the Expr class hierarchy

```

1  abstract class Expr {
2      abstract String operator();
3  }
4  abstract class BinaryExpr extends Expr {
5      final Expr left;
6      final Expr right;
7      abstract BinaryExpr newBinaryExpr(Expr l, Expr r);
8      public String toString() { return left + operator() + right; }
9  }
10 final class AndExpr extends BinaryExpr {
11     String operator() { return " and "; }
12     BinaryExpr newBinaryExpr(Expr l, Expr r) { return new AndExpr(l,r); }
13 }
14 final class OrExpr extends BinaryExpr {
15     String operator() { return " or "; }
16     BinaryExpr newBinaryExpr(Expr l, Expr r) { return new OrExpr(l,r); }
17 }
18 final class Main {
19     public static void main(String[] args) {
20         BinaryExpr b = new AndExpr(); // why isn't the type of e AndExpr?
21         Expr[] a = new Expr[3]; // an empty array of size 3
22         // what is the type of the object stored in each element of the array?
23         a[0] = b;
24         a[1] = new OrExpr();
25         a[2] = b.newBinaryExpr(null,null); // monomorphic call site
26         for (int i = 0; i < a.length; i++) {
27             Expr e = a[i];
28             System.out.println(e.operator()); // polymorphic call site
29             System.out.println(e.toString()); // mono or polymorphic?
30         }
31     }
32 }

```

Exercise 3.1.5 Draw a UML class diagram for this code.

Exercise 3.1.6 Draw an object diagram showing the relationship between the variables and the objects for the main method. (We have drawn this kind of diagram on the board in class, and these kinds of diagrams are also drawn by PythonTutor.com.)

Exercise 3.1.7 Annotate this diagram with the *static type* of each variable and the *dynamic type* of each object. In the past we only considered the case where the static type of the variable was the same as the dynamic type of the object referred to by that variable. Now we also consider the case where the dynamic type of the object is a subtype of the variable's static type.

Exercise 3.1.8 Is there any aliasing occurring in the main method? If so, what is it?

3.2 Polymorphism & Dynamic Dispatch[†]

A *monomorphic call site* is one where the dynamic dispatch will always resolve to the same target. A *polymorphic call site* is one where the dynamic dispatch might resolve to a different target each time the call is executed.

The Java compiler/runtime system inserts code like the following at each potentially polymorphic call site:

```
if (a[i] instanceof AndExpr) {
    return AndExpr::operator(); }
else if (a[i] instanceof OrExpr) {
    return OrExpr::operator(); }
```

This feature of implicit type tests to determine which method definition to execute is one of the key features of object-oriented programming. Many design patterns are about ways to organize code with this feature that make certain kinds of anticipated changes to the software modular.

3.3 Write a recursive-descent recognizer for \mathcal{F}

Figure 3.1 lists a grammar for \mathcal{F} . A recognizer merely computes whether a sentence is generated by a grammar: *i.e.*, its output is boolean. A parser, by contrast, also constructs an abstract syntax tree (AST) of the sentence that we can do things with. A recognizer is simpler and we will write one of them first.

The idea is simple: for each production in the grammar we make a function in the recognizer. These functions have no arguments and return void. All these functions do, from a computational standpoint, is examine the lexer's current token and then advance the token. If the recognizer manages to push the lexer to the end of the input without encountering an error then it declares success.

Exercise 3.3.1 Can you write a regular expression recognizer for \mathcal{F} ?

3.4 Write a pretty-printer for the AST

Pretty-printing is the inverse operation of parsing: given an AST, produce a string. (Parsing produces an AST from a string.) In this case the task is easy: implement the `toString()` methods of the AST classes. When the program invokes `toString()` on the root node of the AST the result should resemble the original input string.

3.5 Equals, Isomorphic, and Equivalent[†]

We were previously introduced to the *object contract* and the *equals* method, and learned that all of the equals methods together are supposed to define a *partitioning* of the objects, which has three properties: reflexivity, symmetry, and transitivity. But we did not talk about the specific semantics of the equals method. For example, always returning true from every equals method would define a partitioning: there would just be one partition and every object would be in it. That's not particularly useful.

In this lab we will give some more meaningful semantics to the equals method, and we will also define two other partitionings with different semantics: *isomorphic* and *equivalent*. Here is an example:

$$X \leq A \text{ or } !A; \xleftrightarrow{\text{equals}} X \leq A \text{ or } !A; \xleftrightarrow{\text{isomorphic}} X \leq !A \text{ or } A; \xleftrightarrow{\text{equivalent}} X \leq 1;$$

equals: Two objects are equals if any computation that uses either one will produce identical results.¹ This can only be true if the objects are *immutable* (*i.e.*, the values stored in their fields do not change).

isomorphic: We will say that two objects are *isomorphic* if they have the same elements and similar structures. For example, we will consider the expressions $X \text{ or } Y$ and $Y \text{ or } X$ to be isomorphic: they

Sources:

`ece351.f.rdescent.FRecursiveDescentRecognizer`

Libraries:

`ece351.util.Lexer`

`ece351.util.CommandLine`

Tests:

`ece351.f.rdescent.TestFRDRecognizerAccept`

`ece351.f.rdescent.TestFRDRecognizerReject`

Sources:

`ece351.f.ast.FProgram`

`ece351.common.ast.AssignmentStatement`

`ece351.common.ast.ConstantExpr`

`ece351.common.ast.UnaryExpr`

`ece351.common.ast.BinaryExpr`

Tests:

manual inspection of output

parser tests below

A *partitioning* is also known as a *mathematical equivalence class*. We'll try to move towards the term 'partitioning' and away from 'mathematical equivalence class', because we use the words 'equivalence' and 'class' in other programming contexts.

[†] B. Liskov and J. Guttag. *Program Development in Java: Abstraction, Specification, and Object-Oriented Design*. Addison-Wesley, 2001

are permutations of the same essential structure. Any two objects that are equals are also isomorphic, but isomorphic objects are not necessarily equals. For example, the expressions X or Y and Y or X are isomorphic but not equals.

equivalent: We will say that two objects are *equivalent* if they have the same meaning, but possibly different structures and possibly different elements. For example, the expression 1 is equivalent to the expression X or $!X$: they have the same meaning but totally different syntax. Any two objects that are isomorphic are also equivalent, but not necessarily *vice versa* (as in this example).

equals \Rightarrow *isomorphic*

You will not be implementing equivalent this term, but you will eventually (not for this lab) need to understand what it means and how we have implemented it for FProgram.

isomorphic \Rightarrow *equivalent*

3.6 Write Equals and Isomorphic for the \mathcal{F} AST classes

Start with VarExpr and ConstantExpr. You should be able to fill in these skeletons with the knowledge you have learned so far.

For the FProgram and AssignmentStatement classes you will need to make recursive function calls. Understanding why these calls will terminate requires understanding recursive object structures.

For the UnaryExpr and CommutativeBinaryExpr classes you will also need to understand inheritance and polymorphism.

Exercise 3.6.1 Why do we use getClass() for type tests instead of instanceof?

Exercise 3.6.2 Why do we cast to UnaryExpr instead of NotExpr?

Exercise 3.6.3 How does the operator() method work?

Sources:

```
ece351.common.ast.VarExpr
ece351.common.ast.ConstantExpr
ece351.common.ast.AssignmentStatement
ece351.common.ast.UnaryExpr
ece351.common.ast.CommutativeBinaryExpr
ece351.f.ast.FProgram
```

Libraries:

```
ece351.util.Examinable
ece351.util.Examiner
ece351.util.ExaminableProperties
```

Tests:

```
ece351.f.test.TestObjectContractF
```

3.7 Write a recursive-descent parser for \mathcal{F}

Our recursive-descent parser will follow the same structure as our recursive-descent recognizer. The steps to write the parser are the same as in the previous lab:

- Copy the procedures from the recognizer.
- For each procedure, change its return type from void to one of the AST classes.
- Modify each procedure to construct the appropriate AST object and return it.

Exercise 3.7.1 Is the result of pretty-printing an AST always character-by-character identical with the original input program? Why?

Exercise 3.7.2 Write a program in \mathcal{F} for which the result of pretty-printing the AST is not identical to the original input string.

Sources:

```
ece351.f.rdescent.FRecursiveDescentParser
```

Tests:

```
ece351.f.rdescent.TestFRDParserBasic
ece351.f.rdescent.TestFRDParser
```

It is important that you follow these steps. If you try to implement the parser for \mathcal{F} in one big method you will run into a lot of trouble.

3.8 Steps to success

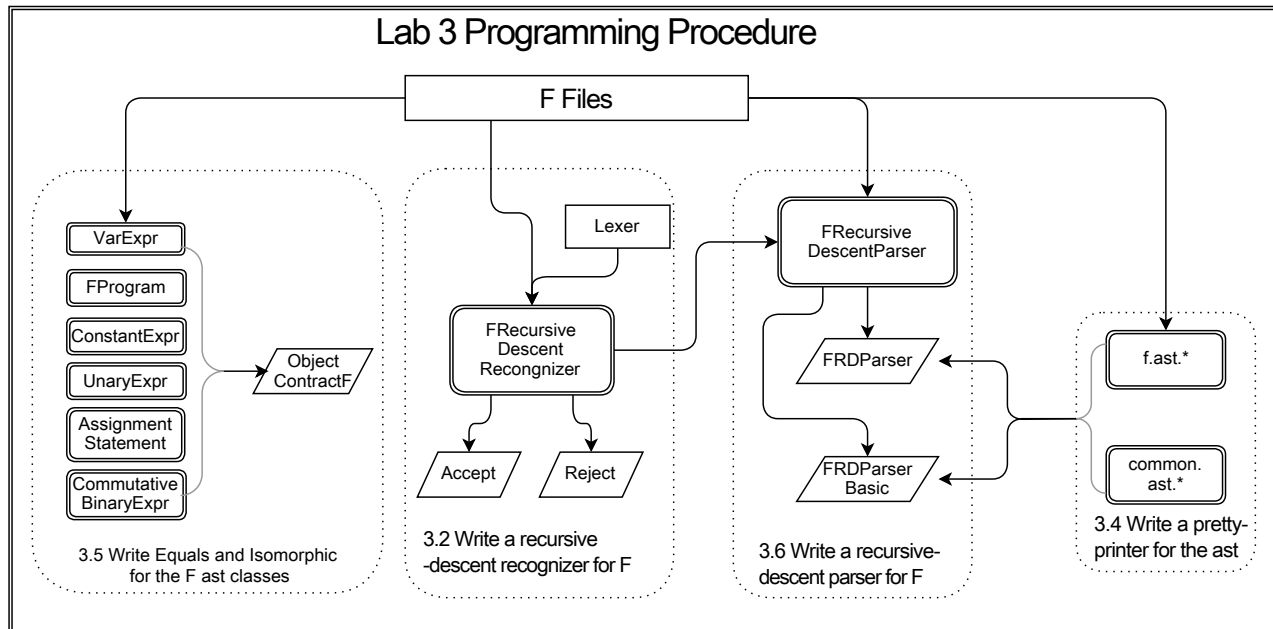


Figure 3.5: Steps to success for §3.
Legend for icon meanings in Figure 1.5.

3.9 Common Missteps

NOT TESTING YOUR PRETTY-PRINTER before starting to write the parser. There is no test suite for the pretty printer. You need to test it manually by looking at the output. If your pretty printer doesn't work correctly then you will get weird failures in the parser tests.

NOT FOLLOWING THE METHODOLOGY for writing recursive descent recognizers and parsers from the grammar. In LAB1 you could have got away with this because \mathcal{W} is a regular language and so there was no real recursion in the grammar. Because \mathcal{F} has nested expressions there is recursion in the grammar: `Expr` calls `Formula`, which in turn calls `Expr` back. If you do not follow the methodology, then you will probably get a parser that works for short formulas but not for long ones. And it will probably also have other bugs.

ASSOCIATIVITY of `BinaryExprs`. Since `AND` and `OR` are not associative, arithmetically you can parse them either way (left-associative or right-associative). For uniformity of testing, we require the parse to be left-associative: *i.e.*, `a or b or c` should parse as `(a or b) or c`.

`TestFRDParserBasic.testLeftAssociativeOr()`

3.10 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

The testing equation for this lab is:

$$\forall \text{AST} \mid \text{AST.equals}(\text{parse}(\text{prettyprint}(\text{AST})))$$

At present we have about 80 AST's to plug into this equation. We might explore some mechanical techniques to generate more AST's.

	Current	New
TestFRDRecognizerAccept	5	5
TestFRDRecognizerReject	5	5
TestFRDParserBasic	5	5
TestFRDParser	30	10
TestObjectContractF	20	10

One of the great benefits of having a testing equation is that we can separate out the correctness condition (the equation) from the inputs. Then we can look for other techniques to generate inputs. By contrast, if we test with just specific input/output pairs the correctness condition is that we get the output for that input.

3.11 Background & Reading

Tiger Book

3.0 Parsing

3.1 Context-Free Grammars

3.2.0 Predictive Parsing / Recursive Descent

- including First and Follow Sets
- read Constructing a Predictive Parser on p.50
- skip Eliminating Left Recursion and everything that comes after it — for now (we'll study this for exam, but not for the labs)

4.1.1 Recursive Descent

*Thinking in Java*² is a good resource for object-oriented programming in general and Java in particular.

² B. Eckel. *Thinking in Java*. Prentice-Hall, 2002. <http://www.mindview.net/Books/TIJ/>

1 Introduction to Objects

6 Reusing Classes

7 Polymorphism

Some topics you should be comfortable with include:

- inheritance / subtyping / subclassing
- polymorphism / dynamic dispatch
- objects *vs.* variables, dynamic *vs.* static type
- type tests, casting

If you are not already familiar with these topics then this lab is going to take you longer than five hours. Once you get through these background topics then the rest of the course material should be fairly straightforward.

ECE155 Lecture Notes In winter 2013 ECE155 switched to Java (from C#), and the third lecture covered some of the differences.

<http://patricklam.ca/ece155/lectures/pdf/L03.pdf>

Lab 4

Circuit Optimization: \mathcal{F} Simplifier

Compiler Concepts: intermediate languages, identity element, absorbing element, equivalence of logical formulas, term rewriting, termination, confluence, convergence

Programming Concepts: interpreter design pattern, template design pattern, representation invariants

\mathcal{F} is an intermediate language for our circuit synthesis and simulation tools, in between VHDL and the final output languages, as was depicted in Figure 1. In a subsequent lab we will write a translator from VHDL to \mathcal{F} .

Compilers usually perform their optimizations on programs in an intermediate language. These intermediate languages are designed to be easier to work with mechanically, at the cost of being less pleasant for people to write large programs in. A program written in \mathcal{F} , for example, is just a list of boolean formulas. This is relatively easy to manipulate mechanically. VHDL, by contrast, has conditionals, module structure, *etc.*, which are all of great benefit to the VHDL programmer but are more work for the compiler writer to manipulate.

The simplifier we will develop in this lab will work by *term-rewriting*. For example, when it sees a term of the form $x + 1$ it will rewrite it to 1. Figure 4.5 lists the algebraic identities that your simplifier should use.

Our simplifier works at a syntactic level: *i.e.*, it does not have a deep understanding of the formulas that it is manipulating. The testing framework to determine if the simplifier has produced correct output does, however, do a deep semantic analysis of the formulas, as discussed in §4.9.

You might have previously studied semantic techniques for boolean circuit simplification such as Karnaugh Maps and the Quine-McCluskey algorithm for computing prime implicants. By completing this lab you will have the necessary compiler-related knowledge to implement these more sophisticated optimization algorithms in the future.

4.1 Iteration to a Fixed Point and Termination[†]

A common technique in optimizing compilers is to *iterate to a fixed point*: that is, to keep applying the optimizations until the program being transformed stops changing. Figure 4.1 shows the code in `Expr.simplify()` that we will run in this lab to iterate to fixed point.

```
final public Expr simplify() {
    Expr e = this;
    while (true) { // loop forever?
        final Expr simplified = e.simplifyOnce();
        if (simplified.equals(e)) {
            // we're done: nothing changed
            return simplified;
        } else {
            // something changed: keep working
            e = simplified;
        }
    }
}
```

Notice the loop `while (true)`. Is this an infinite loop? Will this code ever terminate? Maybe. Inside the body of the loop there is a return statement that will exit the loop (and the method) when `simplified.equals(e)`. We refer to this test to decide whether to exit the loop as the *termination condition*. Will this termination condition ever be true? Will it become true for any possible \mathcal{F} program that we might optimize?

Yes, this termination condition will become true for any possible \mathcal{F} program that we optimize (unless there are other bugs in your code). How do we know this? Because our optimizer applies rewrite rules that make the formula smaller, and there is a limit to how small a formula can get. Therefore, we say that our term rewriting system is *terminating*: there is no input \mathcal{F} program on which it will get stuck in an infinite loop.

We will use this idea of iterating to a fixed point on paper in the course notes when analyzing a grammar to determine if it is $LL(1)$ and when doing dataflow analysis.

Figure 4.1: Code listing for `Expr.simplify()` showing iteration to a fixed point. This is the only implementation of `simplify()` in the project, and it is provided for you. You will be working on implementing `simplifyOnce()` for some of the AST classes.

From this termination condition we can reason that the `simplify` method is *idempotent*: that is, if we apply it twice we get the same result as if we apply it once. We first saw this property above in §0.14 in the form $f(x) = f(f(x))$. In our code here, `x.simplify().equals(x.simplify().simplify())`

Idempotence is an important general property that can be exploited in testing. It is one of the main general properties that must hold for data synchronizers: if a synchronization is performed, and no data changes on either side, then a second synchronization is performed, the second synchronization should not need to perturb the data. <http://en.wikipedia.org/wiki/Idempotence>

4.2 Confluence[†]

We now know that our term rewriting system is *terminating*: it will never get stuck in an infinite loop. But will it always compute the same result? What if we apply the rewrite rules in a different order? A term rewriting system that always reaches a unique result for any given input, regardless of the order in which the rules are applied, is called *confluent*.

Suppose, as a counter-example, that we consider a term rewriting system for the square root operator ($\sqrt{}$) that has two rewrite rules, one that gives the positive root and another that gives the negative root. This rewrite system is not confluent. For example, $\sqrt{4} \mapsto 2$ and $\sqrt{4} \mapsto -2$, and $2 \neq -2$. It is harder to test code that implements a non-confluent term rewriting system because there could be different outputs for the same input.

The term rewriting system that we are implementing in this lab is confluent. Consider the example in Figure 4.2. Whichever order we apply the rules in we eventually get to the same result.




Figure 4.2: Example of two confluent rewrites converging on a common solution

Proving that a term rewriting system is confluent is, in the general case, a fairly difficult problem. For our purposes, we will consider a rewrite system to be confluent if there are no two rules that have the same left-hand side. Our counter-example above had two rules with \sqrt{x} as the left-hand side. The rules that we are implementing for this lab all have different left-hand sides.

convergent = confluent + terminating

There are a variety of sources online where you can read more if you are interested. One good one is by Paul Klint: <http://www.meta-environment.org/doc/books/extraction-transformation/term-rewriting/term-rewriting.html>

4.3 Mathematical Properties of Binary Operators[†]

These are things you’ve probably learned and forgotten many times. Turns out they are actually important in this course.

The term *binary operators* refers to the grammatical fact that these operators take two operands (arguments) — not to the semantic issue that some operators apply to boolean (binary) values. For example, conjunction (logical and) is binary because it takes two operands, not because it operates on true/false values.

4.3.1 Commutativity

Changing the order of the arguments doesn’t matter, *e.g.*:

$$x + y = y + x$$

Addition, multiplication, conjunction (and), disjunction (or) are all commutative. Subtraction and division are not commutative: the order of the operands (arguments) matters.

4.3.2 Associativity

Wikipedia says it nicely:¹

¹ http://en.wikipedia.org/wiki/Associative_property

Within an expression containing two or more occurrences in a row of the same associative operator, the order in which the operations are performed does not matter as long as the sequence of the operands is not changed. That is, rearranging the parentheses in such an expression will not change its value. Consider, for instance, the following equations:

$$1 + (2 + 3) = (1 + 2) + 3$$

Addition, multiplication, conjunction (logical and), disjunction (logical or) are all associative.

Subtraction, division, exponentiation, and vector cross-product are not associative. The terminology can be a little bit confusing here. On the one hand we say that subtraction is *not associative*, and on the other hand we say it is *left associative* because:

$$x - y - z = (x - y) - z$$

Exponentiation is *right associative*:

$$x^{y^z} = x^{(y^z)}$$

Saying that an operator is *not associative* means that it is either *left associative* or *right associative*.

4.4 Transforming *BinaryExprs* to *NaryExprs*

To test your simplifier we'll need to compare the output it computes with the output that the staff simplifier computes. In a previous lab we wrote an *isomorphic* method that accounted for the commutativity property of many binary operations: *e.g.*, it would consider $x + y$ and $y + x$ to be isomorphic. That's a good start, but we'll need more than this to really evaluate your simplifier mechanically.

Many of the binary operations we consider are also *associative*. That is, $x + (y + z)$ is equivalent to $(x + y) + z$. Disjunction and conjunction are both associative, as are addition and multiplication.

Detecting the equivalence of $x + (y + z)$ and $(x + y) + z$ is tricky because these parse trees are not structurally similar (recall that 'structurally similar' is what *isomorphic* means). We could try to implement a really clever *equivalent* method that detects this equivalence, or we could go all out and compute the truth tables for each expression and compare them, or we could try something else: transforming these trees into a standardized form for which it is easy to check isomorphism.

Figure 4.3 shows four different trees that all represent the same logical expression. The first three trees are binary. Comparing them for equivalence is difficult. However, all three binary trees can be transformed into the sorted n-ary tree fairly easily, and these sorted n-ary trees can be easily compared to each other for isomorphism.

Note that this transformation requires that the operator (*e.g.*, '+') be both associative and commutative. Fortunately, both conjunction and disjunction are associative and commutative, and these are the only operators that we are applying this transformation to.


Where to implement this idea will be described below.

We say that an operator is *binary* because it has two operands. It is just a coincidence that our binary operators operate on boolean values. For example, arithmetic addition is also a binary operator, even though it operates on numbers. We say that an operator is *n-ary* if it has a variable number of operands, possibly more than two.


conjunction means 'logical and'

disjunction means 'logical or'


right-associative parse
of $x + y + z$




left-associative parse
of $x + y + z$



right-associative parse
of $y + x + z$



sorted n-ary tree
of $x + y + z$



Where to implement this transformation will be described below.

Figure 4.3: Four trees that represent logically equivalent circuits. The sorted n-ary representation makes equivalence comparisons easier.

4.5 Identity Element & Absorbing Element[†]

Let I represent the *identity element* of a binary operation \otimes on elements of set S , and let A represent the *absorbing element* of that operation on that set. Then the following equations hold:

$$\begin{aligned}\forall x \in S \mid x &= x \otimes I \\ \forall x \in S \mid A &= x \otimes A\end{aligned}$$

What are the absorbing and identity elements for conjunction and disjunction in boolean logic?

4.6 Simplify Once

In §4.1 above we discussed how the `simplify()` method keeps calling `simplifyOnce()` until there are no changes (*i.e.*, it iterates to a fixed point). Now we turn our attention to `simplifyOnce()`.

Figure 4.5 lists all of the transformations to be implemented within the `simplifyOnce()` methods. By far the most interesting case is `NaryExpr.simplifyOnce()`, which is broken down into a number of sub-cases, as shown in Figures 4.4 and 4.5.

```
@Override
protected final Expr simplifyOnce() {
    assert repOk();
    final Expr result =
        simplifyChildren().
        mergeGrandchildren().
        foldIdentityElements().
        foldAbsorbingElements().
        foldComplements().
        removeDuplicates().
        simpleAbsorption().
        subsetAbsorption().
        singletonify();
    assert result.repOk();
    return result;
}
```

For example, with integer addition and multiplication we have:

$$\begin{aligned}\forall x \in \mathbb{Z} \mid x &= x + 1 \\ \forall x \in \mathbb{Z} \mid 0 &= x + 0 \\ \forall x \in \mathbb{Z} \mid x &= x \times 0\end{aligned}$$

Sources:

```
ece351.common.ast.AndExpr
ece351.common.ast.OrExpr
ece351.common.ast.NotExpr
ece351.common.ast.NaryExpr
....NaryOrExpr.getIdentityElement()
....NaryOrExpr.getAbsorbingElement()
....NaryAndExpr.getIdentityElement()
....NaryAndExpr.getAbsorbingElement()
```

Libraries:

```
NaryExpr.filter()
NaryExpr.removeAll()
NaryExpr.contains()
NaryExpr.getThatClass()
```

Tests:

```
TestSimplifierEquivalence
TestSimplifier2
```

Figure 4.4: The steps of `NaryExpr.simplifyOnce()`.

You are not obliged to use exactly these steps.

These methods have been stubbed with `return this`, so that the code will run — it just won't do any transformations until you implement them.

All of these methods transform an `NaryExpr` to a new `NaryExpr` — except `singletonify()`, which might return a different kind of `Expr`. Why? Most of these transformations reduce the number of nodes in the `AST`. That might result in an `NaryExpr` with just one child, which doesn't make sense. So `singletonify()` will replace this malformed `NaryExpr` with its lone child. Implement `singletonify()` early.


Level	Description	Code & Transformation	Notes
0	Convert binary to n-ary	<code>OrExpr.simplifyOnce()</code> <code>AndExpr.simplifyOnce()</code> <code>NaryExpr.mergeGrandchildren()</code>	
1	Fold Identity Elements	 <code>NaryExpr.foldIdentityElements()</code> $x \cdot 1 \mapsto x$ $x + 0 \mapsto x$	Implement <code>NaryExpr.singletonify()</code> before implementing folds. Why?
2	Fold Absorbing Elements	<code>NaryExpr.foldAbsorbingElements()</code> $x \cdot 0 \mapsto 0$ $x + 1 \mapsto 1$	Write more concise code by using <code>getIdentityElement()</code> and <code>getAbsorbingElement()</code>
3	Fold Negation	<code>NotExpr.simplifyOnce()</code> $!0 \mapsto 1$ $!1 \mapsto 0$ $!!x \mapsto x$	
4	Fold Complements	<code>NaryExpr.foldComplements()</code> $x \cdot !x \mapsto 0$ $x + !x \mapsto 1$	
5	Deduplication	<code>NaryExpr.removeDuplicates()</code> $x + x \mapsto x$ $x \cdot x \mapsto x$	
6	Simple Absorption	<code>NaryExpr.simpleAbsorption()</code> $x + (x \cdot y) \mapsto x$ $x \cdot (x + y) \mapsto x$	Is x a <code>VarExpr</code> or an <code>NaryExpr</code> ? The latter case is harder.
7	Subset Absorption	<code>NaryExpr.subsetAbsorption()</code> $(a + b) + ((a + b) \cdot y) \mapsto a + b$	

Figure 4.5: Simplifications for \mathcal{F} programs

4.7 Object sharing in the simplifier

Figure 4.6 was drawn by some students to help you visualize which objects are shared between the original AST and the simplified AST. The general pattern is that the leaves of the tree will be re-used (shared), whereas the interior nodes will be replaced.

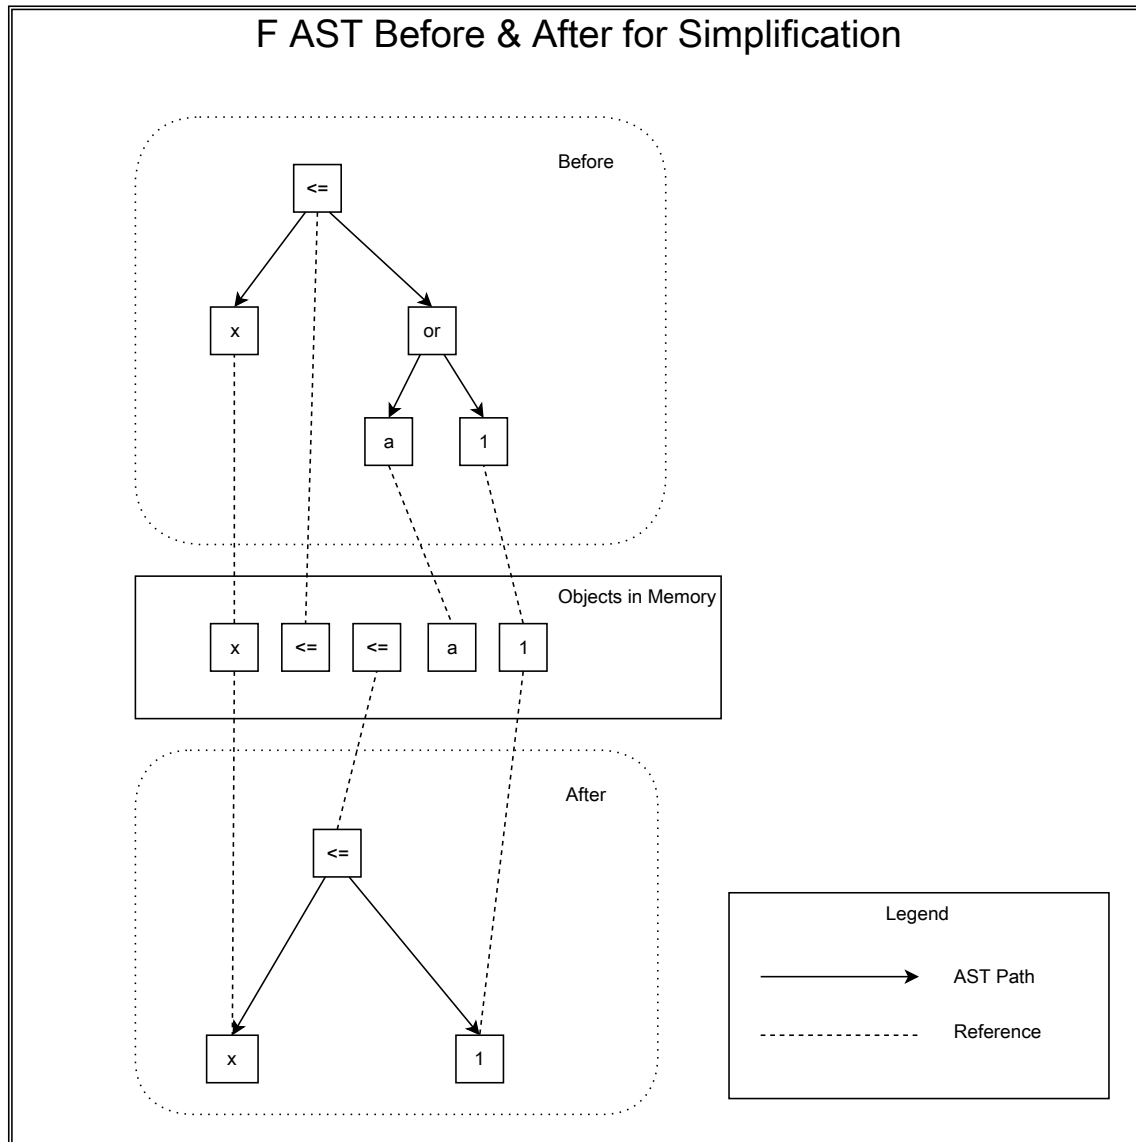


Figure 4.6: Object sharing in the \mathcal{F} simplifier. The original AST and the simplified AST will share some nodes. This is permissible because all of our AST nodes are immutable.

4.8 Class Representation Invariants[†]

The *invariants* of a class are conditions that are expected to hold for all legal objects of that class. For example, a legal FProgram will contain at least one formula and will have exactly one formula for each output pin (*i.e.*, no output pin will be computed by two different formulas). A legal NaryExpr will have more than one child, its children must not be null, and its children must be sorted.

For mutable objects, invariants are usually expected to hold at all public method boundaries (before the method starts and finishes). All of our AST objects are immutable. We take the position that it is possible to construct illegal AST objects (where the invariants do not hold) as temporary values. Consider the code for NaryExpr.simplifyOnce() in Figure 4.4: each of the helper methods (except singletonify) is allowed to construct and return an illegal NaryExpr, but at the end we expect to have a legal Expr object.

http://en.wikipedia.org/wiki/Class_invariant

By convention, the class representation invariants are checked in a method called repOk().

```
public boolean repOk() {
    // programming sanity
    assert this.children != null;
    // should not have a single child: indicates a bug in simplification
    assert this.children.size() > 1 : "should have more than one child, probably a bug in simplification";
    // check that children is sorted
    int i = 0;
    for (int j = 1; j < this.children.size(); i++, j++) {
        final Expr x = this.children.get(i);
        assert x != null : "null children not allowed in NaryExpr";
        final Expr y = this.children.get(j);
        assert y != null : "null children not allowed in NaryExpr";
        assert x.compareTo(y) <= 0 : "NaryExpr.children must be sorted";
    }
    // Note: children might contain duplicates ---- not checking for that
    // ... maybe should check for duplicate children ...
    // no problems found
}
```

Figure 4.7: NaryExpr.repOk() class representation invariants for NaryExpr; *i.e.*, the rules that define what is a well-formed NaryExpr object.

4.9 Logical equivalence for boolean formulas[†]

\mathcal{F} is a language of boolean formulas. Checking equivalence of boolean formulas is an NP-complete problem.

To see why checking equivalence of boolean formulas is NP-complete consider the example of comparing $(x + y) + z$ and $x + (y + z)$ and $!(\neg x \cdot \neg y \cdot \neg z)$ by constructing their truth tables, where f names the output:

$(x + y) + z$	$= f$	$x + (y + z)$	$= f$	$!(\neg x \cdot \neg y \cdot \neg z)$	$= f$
0	0	0	0	0	0
0	1	0	1	0	1
0	0	1	0	0	0
0	1	1	1	0	1
1	0	0	1	1	0
1	0	1	1	1	0
1	1	0	1	1	1
1	1	1	1	1	1

We can see from examining the three truth tables that these three formulas are equivalent. Great. But the number of rows in each truth table is an exponential function of the number of variables in the formula. For example, these formulas have three variables and eight rows in their truth tables: $2^3 = 8$.

What to do? This is a hard problem. As discussed in the *Course Notes* §0 we have three main options: implement an NP-complete algorithm ourselves; use a polynomial time approximation; or convert the problem to SAT and ask a SAT-solver for the answer. If you look at the skeleton code in `FProgram.equivalent()` you will see that it uses this last approach, and this is the approach that we generally recommend you follow in your future career (unless there are well known and accepted polynomial time approximations for the problem you are trying to solve).

For this particular problem of computing the equivalence of boolean formulas there is a fourth option: translate the formulas into a standardized form, and then compare that standardized form. This is the approach that we took above when we converted the binary expressions to sorted n-ary expressions. In the case of boolean formulas, the most common standardized form is *reduced, ordered binary decision diagrams* (or BDD for short).² A BDD is essentially a compressed form of a truth table.

Two equivalent boolean formulas will have the exact same BDD representation. Constructing a BDD for a boolean formula is, in the worst case, an exponential problem, but in practice usually runs fairly quickly. BDDs are commonly used for hardware verification and other tasks that require working with boolean formulas.

The correctness of your simplifier is determined by checking whether the output is logically equivalent to, yet structurally smaller (or no larger) than the input.

Is the termination condition in `Expr.simplify()` (Figure 4.1) written in terms of logical equivalence? Why?

Figure 4.8: Equivalent truth tables

Our translation to SAT occurs via an intermediate language named *Alloy* (<http://alloy.mit.edu>), which is in turn translated to SAT. We found it convenient to work this way, but it would also be easy to translate this problem to SAT directly.

The SAT solver that Alloy is configured to use in this case is SAT4J, which is open-source, written in Java, and used by Eclipse to compute plugin dependencies.

² R. E. Bryant. Graph-based algorithms for boolean function manipulation. *IEEE Transactions on Computers*, C-35(8):677–691, Aug. 1986

There will be no exam questions on BDDs. This material is just here for your interest. Boolean formulas are of fundamental and intrinsic interest both in theory and in practice.

4.10 *BinaryExpr.examine()* is a Higher-Order Function[†]

A higher-order function is a function that takes another function as an argument. `BinaryExpr.examine()` is an example: the first argument, the Examiner, is essentially another function. The Examiner object has no state: it is just used to call the `examine()` method.

```
private boolean examine(final Examiner e, final Object obj) {
    // basics
    if (obj == null) return false;
    if (!this.getClass().equals(obj.getClass())) return false;
    final BinaryExpr be = (BinaryExpr) obj;

    // compare field values
    if (le.examine(left, be.left)) return false;
    if (le.examine(right, be.right)) return false;

    // no differences
    return true;
}
```

Higher-order functions were briefly introduced in [No.6]

Figure 4.9: `BinaryExpr.examine()` is a higher-order function

Exercise 4.10.1 How many Examiner objects are there in this codebase? What are their names?

Exercise 4.10.2 Is it possible to create any more Examiner objects?

4.11 *Choice of Data Structures for NaryExpr.children*[†]

`NaryExpr` uses an `ImmutableList` to store its children. Some points of discussion:

- `ImmutableList` is used widely in this code, so it is familiar.
- We must use linear search with `ImmutableList`. Perhaps some other data structure would offer better asymptotic complexity. For example, a hash structure would give us constant time lookup. But, we are engineers: we care about constant factors. Linear search in a small list can be faster, in practice, than looking up in a hash structure — depending on what the constant factors are.
- A list is the most compact storage for this purpose. The children of `NaryExpr` need to be ordered. So a structure like `HashMap` would not be adequate: `LinkedHashMap` would be necessary. `LinkedHashMap` maintains ordering via an internal list, in addition to the space allocated for the hash structure.
- `ImmutableList.remove()` will throw exceptions because `ImmutableList` is immutable.

`NaryExpr.contains()`

4.12 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below. The testing equations are as follows. First, the simplifier should always produce an equivalent \mathcal{F} program, for all inputs:

```
originalAST.simplify().equivalent(originalAST)
```

Second, the simplifier should be idempotent:

```
originalAST.simplify().equals(originalAST.simplify().simplify())
```

Finally, for a select set of \mathcal{F} programs (`opt*.f`), the simplifier should perform some specific transformations:

```
originalAST.simplify().isomorphic(staffSimplifiedAST)
```

	Shared	Secret
TestSimplifierEquivalence	30	10
TestSimplifier2	40	20

In the unlikely, but not impossible, event that you implement a better simplifier than the staff has then your simplifier might produce spurious failures while evaluating the correctness equation. Talk to us to rectify the situation and get some bonus marks.

Note that you do not earn any marks for this lab until `TestObjectContractF` from `LAB3` passes. As you can see, the testing equations for this lab depend on having the object contract correctly implemented for the \mathcal{F} AST classes. If your object contract implementation is buggy then we cannot trust the results of the testing equations.

Lab 5

Parsing \mathcal{W} with Parboiled

In this lab you will write a new parser for \mathcal{W} . In lab1 you wrote a parser for \mathcal{W} by hand. In this lab you will write a parser for \mathcal{W} using a parser generator tool named *Parboiled*. A parser generator is a tool that takes a description of a grammar and generates code that recognizes whether input strings conform to that grammar. Many developers in practice choose to use a parser generator rather than write parsers by hand.

There are many different parser generator tools, and there are a number of dimensions in which they differ, the two most important of which are the theory they are based on and whether they require the programmer to work in an *internal* or *external* DSL. The theory behind Parboiled is called *Parsing Expression Grammars* (PEG). Other common theories are LL (e.g., Antlr) and LALR (e.g., JavaCup).

A DSL is often used in combination with a general purpose programming language, which is sometimes called the *host* language. For example, you might write an SQL query in a string inside a Java program. In this example Java is the host language and SQL is an external DSL for database queries. Whether the DSL is internal or external is determined by whether it shares the grammar of the host language, not by the physical location of snippets written in the DSL. In this example the SQL snippet is written inline in the Java program, but SQL has a different grammar than Java, and so is considered to be an external DSL.

An internal DSL uses the grammar of the host language. A common way to implement an internal DSL is as a library written in the host language. Not all libraries represent DSL's, but almost all internal DSL's are implemented as libraries.

Most parser generators require the programmer to work in an external DSL. That is, they require the programmer to learn a new language for specifying a grammar. Parboiled, by contrast, provides an internal DSL. I think that it is easier to learn an internal DSL, and this is why we have chosen to use Parboiled.

Compiler Concepts: parser generators, Parsing Expression Grammars (PEG), push-down automata
Programming Concepts: domain specific languages (DSL): internal *vs.* external, debugging generated code, stacks

Files:

ece351.w.parboiled.WParboiledRecognizer
ece351.w.parboiled.WParboiledParser

Tests:

ece351.w.parboiled.TestWParboiledRecognizerAccept
ece351.w.parboiled.TestWParboiledRecognizerReject
ece351.w.parboiled.TestWParboiledParserAccept
ece351.w.parboiled.TestWParboiledParserReject

Tool Documentation:

<http://parboiled.org>

DSL = Domain Specific Language. This terminology is in contrast to a *general purpose* programming language, such as Java/C/etc.

Exercise 5.0.1 Name a few more common DSLs

The Tiger Book §3.4 discusses two other parser generator tools, JavaCC and SableCC. You can look at that section to get a sense for what their external DSL's look like, and why it might be easier to learn Parboiled first. Parboiled did not exist when the Tiger book was written.

The purpose of this lab is for you to learn how to use Parboiled. This lab is specifically structured to facilitate this learning. You are already familiar with the input language \mathcal{W} and the host language Java. \mathcal{W} is an even simpler language than is used to teach Parboiled on its website. The only new thing in this lab is Parboiled.

Once you learn to use one parser generator tool it is not too hard to learn another one. Learning your first one is the most challenging.

The other reason why we are using Parboiled is that it is like a pushdown automata because its main storage mechanism is a stack, and so it reinforces that theoretical concept.

5.1 Introducing Parboiled

To write a recognizer or a parser with Parboiled we extend the `BaseParser` class, which either defines or inherits almost all of the methods that we will use from Parboiled. For our labs we will actually extend `BaseParser351`, which in turn extends `BaseParser` and adds some additional utility methods.

We can divide the methods available in our recognizer/parser into a number of groups:

Rule Constructors. These methods are used to describe both the grammar and the lexical specification of the language that we wish to recognize or parse, and we can subdivide this group this way:

EBNF	Parboiled	
*	<code>ZeroOrMore()</code>	
+	<code>OneOrMore()</code>	
?	<code>Optional()</code>	
	<code>FirstOf()</code>	similar but importantly different
	<code>Sequence()</code>	no explicit character in EBNF
Regex	Parboiled	
[ab]	<code>AnyOf("ab")</code>	
[^ab]	<code>NoneOf("ab")</code>	
a	<code>Ch('a')</code> or just <code>'a'</code>	
[a-z]	<code>CharRange('a', 'z')</code>	
	<code>IgnoreCase()</code>	no regex equivalent
	<code>EOI</code>	special char for end of input
	<code>Wo()</code>	optional whitespace (zero or more)
	<code>W1()</code>	mandatory whitespace (at least one)

Access to input. A recognizer doesn't need to store any of the input that it has already examined. A parser, however, often saves substrings of the input into an AST. The `match()` method returns the substring that matched the most recently triggered rule.

Learning progression: repetition with increasing complexity. In a traditional compilers project, without learning progression, you would have to learn everything simultaneously in one big lab: the source language (\mathcal{W}), the host language (Java), and the parser generator's external DSL. Moreover, the source language would be more like our toy subset of VHDL, which we won't see until after midterm week, rather than a simple regular language like \mathcal{W} . Even the pedagogical example on the Parboiled website is as complicated as \mathcal{F} (a simple context-free language).

Recall that in previous labs there was a separate `Lexer` class that encoded the lexical specification for \mathcal{W} (and \mathcal{F}). Some parser generator tools have separate DSL's for the lexical and syntactic specifications of the input language. In Parboiled, by contrast, we specify the tokenization as part of the grammar.

EBNF = Extended Backus Naur Form. This is the name of the notation used to specify the grammars for \mathcal{W} (Figure 1.2) and \mathcal{F} (Figure 3.1).

Regular expressions are often used for lexical specifications.

For this lab we will specify whitespace explicitly. In the next lab we will learn how to specify the whitespace implicitly, which makes the recognizer rules look a bit less cluttered.

The `match()` method is in Parboiled's `BaseActions` class, which is the superclass of `BaseParser`.

Exercise 5.1.1 Draw a UML class diagram for `WParboiledRecognizer` and `WParboiledParser`.

Stack manipulation. A parser builds up an AST. Where do the fragments of this AST get stored while the parser is processing the input? When we wrote the parsers for \mathcal{W} and \mathcal{F} by hand in the previous labs we used a combination of fields, local variables, and method return values to store AST fragments. None of these storage mechanisms are available to us within Parboiled's DSL. Parboiled gives us one and only one place to store AST fragments: the Parboiled value stack. In this sense, working with Parboiled is very much like programming a pushdown automata. We can manipulate this stack with the standard operations, including: `push()`, `pop()`, `peek()`, `swap()`, and `dup()`. When the parser has processed all of the input we expect to find the completed AST on the top of the stack.

Hewlett-Packard calculators also famously have a value stack for intermediate results.

These stack operations are in Parboiled's `BaseActions` class, which is the superclass of `BaseParser`.

Grammar of the input language. We will add a method to our recognizer/parser for each of the non-terminals in the grammar of the input language. These methods will have return type `Rule` and will comprise just one statement: a return of the result of some Parboiled rule constructor.

We might also add some methods for parts of the lexical specification of the input language.

See rule constructors above.

5.2 Write a recognizer for \mathcal{W} using Parboiled

Suppose we want to process files that contain a list of names, such as 'Larry Curly Moe'. A recognizer for such a language might include a snippet as in Figure 5.1.

```
public Rule List() { return ZeroOrMore(Sequence(Name(), W0())); }
public Rule Name() { return OneOrMore(Letter()); }
public Rule Letter() { return FirstOf(CharRange('A', 'Z'), CharRange('a', 'z')); }
```

Figure 5.1: Snippet of a recognizer written in Parboiled's DSL.

Exercise 5.2.1 Write the EBNF grammar that corresponds to this Parboiled code.

5.3 Add actions to recognizer to make a parser

Let's add actions to our recognizer from Figure 5.1 to make a parser. Figure 5.2 lists code for the actions. The general idea is that we wrap every recognizer rule in a `Sequence` constructor, and then add new clauses to manipulate the stack: *i.e.*, a clause with one of the `push`, `pop`, `peek`, `swap`, `dup`, *etc.*, commands.

Figure 5.5 augments the listing of Figure 5.2 with some debugging clauses using the `debugmsg()` and `checkType()` methods. If you want to inspect memory while your recognizer or parser is executing then use one of the `debugmsg()` or `checkType()` methods and set a breakpoint inside that method. *Setting a breakpoint inside a rule constructor will not do what you want.*

`debugmsg()` and `checkType()` are defined in `BaseParser351`, which is a superclass of our recognizer/parser classes and a subclass of Parboiled's `BaseParser`.

Rule constructors are executed once in a grammar analysis phase in order to generate code that will actually process input strings.

```

public Rule List() {
    return Sequence(                                     // wrap everything in a Sequence
        push(ImmutableList.of()),                       // push empty list on to stack
        ZeroOrMore(Sequence(Name(), W0()))              // rule from recognizer
    );
}

public Rule Name() {
    return Sequence(                                     // wrap everything in a Sequence
        OneOrMore(Letter()),                            // consume a Name, which can be retrieved later by match()
        push(match()),                                  // push the Name on top of the stack
        swap(),                                          // swap the Name and the List on the stack
        push( ((ImmutableList)pop()).append(pop()) )    // make a new list by appending the new Name to the old List
    );
}

public Rule Letter() { return FirstOf(CharRange('A', 'Z'), CharRange('a', 'z')); } // rule from recognizer

```

When the parser reaches the end of the input then we expect to find a list object on the top of the stack, and we expect that list object to contain all of the names given in the input. For example, for the input string 'Ren Stimpy' we would expect the result of the parse to be the list ['Ren', 'Stimpy']. Similarly, for the input string 'Larry Curly Moe' we would expect the result of the parse to be the list ['Larry', 'Curly', 'Moe']. The result of a parse is the object on the top of the stack when the parser reaches the end of the input. Figure 5.3 illustrates the state of the stack as this example parser processes the input 'Ren Stimpy'.

Figure 5.2: Snippet of a parser written in Parboiled's DSL, corresponding to the recognizer snippet in Figure 5.1. See this snippet extended with some debugging commands in Figure 5.5.

http://www.youtube.com/results?search_query=ren+stimpy




Figure 5.3: Stack of Parboiled parser from Figure 5.2 while processing input 'Ren Stimpy'. Time proceeds left to right.

5.3.1 An alternative Name() rule

Figure 5.4 shows an alternative formulation of the Name() rule from Figure 5.2.

```

public Rule Name() {
    return Sequence(
        OneOrMore(Letter()),
        push( ((ImmutableList)pop()).append(match()) )
    );
}

```

Exercise 5.3.1 Which state of the stack corresponds to which line of the parser source code?

Figure 5.4: An alternative definition of the Name() rule in Figure 5.2

Exercise 5.3.2 Draw out the stack that this alternative formulation creates (*i.e.*, the analogue of Figure 5.3).

```

public Rule List() {
    return Sequence(                                     // wrap everything in a Sequence
        push(ImmutableList.of()),                       // push empty list on to stack
        ZeroOrMore(Sequence(Name(), W1()))             // rule from recognizer
        checkType(peek(), List.class)                  // expect a list on the top of the stack
    );
}
public Rule Name() {
    return Sequence(                                     // wrap everything in a Sequence
        OneOrMore(Letter()),                           // rule from recognizer
        push(match()),                                 // push a single name, so stack is: [List, String]
        debugmsg(peek()),                             // print the matched name to System.err
        checkType(peek(), String.class)                // match always returns a String
        swap(),                                         // swap top two elements on stack: [String, List]
        checkType(peek(), List.class)                  // confirm that the list is on top
        push( ((ImmutableList)pop()).append(pop()) )   // construct and push a new list that contains
    );                                                  // all of the names on the old list plus this new name
}
public Rule Letter() { return FirstOf(CharRange('A', 'Z'), CharRange('a', 'z')); } // rule from recognizer

```

Figure 5.5: Snippet of a parser written in Parboiled's DSL (Figure 5.2). Extended with debugging commands. Corresponding to the recognizer snippet in Figure 5.1.

5.3.2 The stack for our \mathcal{W} parser will have two objects

The stack for our \mathcal{W} programs will have two objects once it reaches steady state: a WProgram object and a Waveform object. Periodically the Waveform object will be appended with the WProgram object and a new Waveform object will be instantiated. The WProgram object will usually be closer to the bottom of the stack and the Waveform object will usually be closer to the top.

Exercise 5.3.3 Draw out a picture of what you expect the stack to look like at each point in time of the parser execution *before you start programming it*. See Figure 5.3 as an example.

Exercise 5.3.4 What is the maximum size that the stack will grow to?

5.4 Parsing Expression Grammars (PEGS)[†]

Parboiled works with *Parsing Expression Grammars* (PEGS). PEGS are slightly different than *Context Free Grammars* (CFGs). The important difference is that PEGS cannot represent ambiguity, whereas CFGs can. We never want ambiguity in programming languages, so this 'feature' of CFGs is not much of a feature in our context.

CFGs were originally developed by linguists to model natural languages, which often contain grammatical ambiguities. Only after linguists had successfully applied the idea of CFGs in their natural language context did computer scientists try using CFGs for programming languages.

Exercise 5.4.1 What language feature can CFGs represent that PEGS cannot?

http://en.wikipedia.org/wiki/Parsing_expression_grammar
http://en.wikipedia.org/wiki/Context-free_grammar

Where does ambiguity arise in a CFG? From the choice (alternation) operator: ‘|’ (the bar). An ambiguity is when two (or more) different alternatives might apply. PEGs remove this possibility for ambiguity by treating alternatives in priority order: the first alternative to match is taken as correct; if a subsequent alternative would also have matched it is just ignored. So Parboiled does not have a bar operator, it has the `FirstOf` method instead, which says ‘use the alternative in this list that matches first.’

CFGs and PEGs recognize almost the same set of grammars, but there are a few that each can do that the other cannot. For example, a CFG can recognize the following, while a PEG cannot:

$$S \rightarrow 'x' S 'x' \mid 'x'$$

Similarly, there are some grammars that PEGs can recognize that CFGs cannot, such as counting three things:

$$\{a^n b^n c^n \mid n \geq 0\}$$

5.5 Expressions, Statements, and Side Effects[†]

EXPRESSIONS are *evaluated* to a value. From an idealistic perspective, expressions are mathematical: they comprise variables and operators; they do not change the state of the machine.

STATEMENTS are *executed* to do something (*i.e.*, a *side-effect*). Statements end with a semi-colon (;) in C-like languages. They change the state of the machine (their *side-effect*). For example, an *assignment statement* changes the value bound to a variable; a *goto statement* changes the program counter (the statement which is currently being executed). The Wikipedia definition of statement:

In computer programming, a statement is the smallest standalone element of an imperative programming language that expresses some action to be carried out. It is an instruction written in a high-level language that commands the computer to perform a specified action. A program written in such a language is formed by a sequence of one or more statements. A statement may have internal components (e.g., expressions).

https://en.wikipedia.org/wiki/Statement_%28computer_science%29

THE LINE BETWEEN expressions and statements can be a bit fuzzy in modern languages (most languages invented after around 1970; C was one of the first to blurr this line a bit). There are some nice articles on StackOverflow that discuss this bluriness in C# and Python. That discussion is beyond the scope of this course, but might be of interest to you.

C#: <http://stackoverflow.com/questions/19132/expression-versus-statement>

Python: <http://stackoverflow.com/questions/4728073/what-is-the-difference-between-an-expression-and-a-s>

5.6 A few debugging tips

You might get a weird error message that Parboiled has difficulty creating the parser class. If so, see if your rule constructors are throwing exceptions. For example, the skeleton code ships with stubs like this:

```
public Rule Program() {
    // TODO: 1 lines snipped
    throw new ece351.util.TODO351Exception();
}
```

You need to get rid of all of these stubs and replace them with code that does something sensible — or just return null. If you leave the exception throwing in, even in a rule constructor you don't call, Parboiled will complain.

Also, you should explicitly include EOI in your grammar.

5.7 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

	Current	New
TestWParboiledRecognizerAccept	20	5
TestWParboiledRecognizerReject	5	5
TestWParboiledParserBasic	5	0
TestWParboiledParserAccept	30	15
TestWParboiledParserReject	10	5

You do not earn any marks for the rejection tests until some acceptance tests pass. You do not earn any marks for the main parser tests until the basic parser test passes.

The recognizer tests just run the recognizer to see whether it rejects or accepts. The parser rejection tests work similarly. The parser acceptance tests check two testing equations. Let w name the input \mathcal{W} file. Let $x = \text{ParboiledParse}(w)$. The two testing equations are:

```
x.isomorphic(ParboiledParse(PrettyPrint(x)))
x.isomorphic(RecursiveDescentParse(w))
```

5.8 Reading

Parboiled & PEGs:

<http://parboiled.org>
<https://github.com/sirthias/parboiled/wiki/Grammar-and-Parser-Debugging>
<https://github.com/sirthias/parboiled/wiki/Handling-Whitespace>
<http://www.decodified.com/parboiled/api/java/org/parboiled/BaseParser.html>
http://en.wikipedia.org/wiki/Parsing_expression_grammar

JavaCC & SableCC: (two other parser generators)

Tiger Book §3.4

Note that both JavaCC and SableCC require the grammar to be defined in an external DSL (*i.e.*, not in Java).

You will not be tested on specifics of JavaCC and SableCC, but you should know of their existence and that they use an external DSL for specifying the grammar.

Lab 6

Parsing \mathcal{F} with Parboiled

This lab is similar to some previous labs: we will write a recognizer and parser for \mathcal{F} , this time using Parboiled.

6.1 Write a recognizer for \mathcal{F} using Parboiled

<i>Program</i>	→	Formula ⁺ \$\$
<i>Fomula</i>	→	Var '<=' Expr ';'
<i>Expr</i>	→	Term ('or' Term)*
<i>Term</i>	→	Factor ('and' Factor)*
<i>Factor</i>	→	'not' Factor '(' Expr ')' Var Constant
<i>Constant</i>	→	"0" "1"
<i>Var</i>	→	id

INSPECT TOKEN WITHOUT CONSUMING IT. Parboiled has two rule constructors that allow you to inspect the next token without consuming it: `Test()` and `TestNot`. For example, in a rule for `Var` one might include `TestNot(Keyword())` to ensure that a keyword (e.g., 'and', 'or') is not considered as a variable name.

IMPLICIT WHITESPACE HANDLING. As we saw in a previous lab, Parboiled incorporates both syntactic and lexical specifications. One of the practical consequences of this is that we cannot delegate whitespace handling to a separate lexer (as we did when writing recursive descent parsers by hand). The explicit mention of whitespace after every literal can make the recognizer/parser code look cluttered. There is a standard trick in Parboiled that if we add a single trailing space to each literal then the match for that literal will include all trailing whitespace. For example, in the rule constructor for `Constant` we would write: `FirstOf("0 ", "1 ")` (notice the trailing space after the zero and the the one). If you look at the method `ConstantExpr.make(s)` you will see that it looks only at the first character of its argument `s`, and so thereby ignores any trailing whitespace. When you call `match()` the result will contain the trailing whitespace.

Compiler Concepts: parser generators, Parsing Expression Grammars (PEG), push-down automata

Programming Concepts: domain specific languages (DSL): internal *vs.* external, debugging generated code, stacks

Files:

`ece351.f.parboiled.FParboiledRecognizer.java`
`ece351.f.parboiled.FParboiledParser.java`

Tests:

`ece351.f.parboiled.TestFParboiledRecognizerAccept`
`ece351.f.parboiled.TestFParboiledRecognizerReject`
`ece351.f.parboiled.TestFParboiledParserBasic`
`ece351.f.parboiled.TestFParboiledParser`
`ece351.f.parboiled.TestFParserComparison`

Figure 6.1: LL(1) Grammar for \mathcal{F} (reproduced from Figure 3.1)

The staff solution code does not use the `Test()` method. The `FirstOf()` method provides adequate look-ahead functionality for the \mathcal{F} grammar.

If you are interested in how this trick works there is a page devoted to it on the Parboiled website. This is just a trick of the tool and not part of the intellectual content of this course, so you are not expected to understand how it works.

6.2 Add actions to recognizer to make a parser

Following our standard design method, once your recognizer is working, copy it to start making the parser. As before, you will introduce a new Sequence at the top level of each rule constructor, and the parser actions will be additional arguments to this Sequence.

Your Parboiled \mathcal{F} parser should construct BinaryExprs, just as your recursive descent \mathcal{F} parser did. We can apply the same set of transformations, developed in a previous lab, to produce NaryExprs.

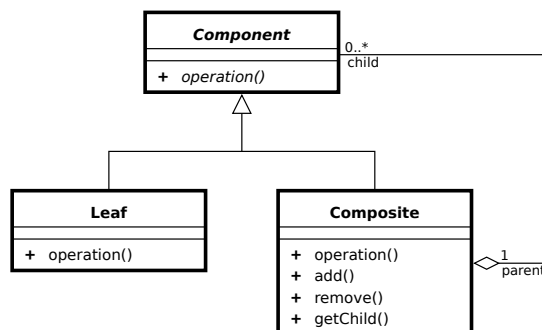
Exercise 6.2.1 Draw the stack as it evolves through a simple parse.

Exercise 6.2.2 How large can the stack grow while parsing \mathcal{F} ?

6.3 Composite Design Pattern[†]

The Composite design pattern is commonly used when we need to construct trees of objects, and we want to work with the leaf nodes and the interior nodes ('composite' nodes) via a common interface ('component'). Figure 6.2 illustrates this idea.

The main place we've seen the composite design pattern in the labs is the \mathcal{F} AST classes. The classes representing the leaves of an AST are ConstantExpr and VarExpr. The classes representing the interior ('composite') nodes of the AST are AndExpr, OrExpr, NotExpr, etc.. The shared interface is Expr ('component').



Why didn't we use the Composite design pattern for the \mathcal{W} AST classes? Because \mathcal{W} is a regular language, and so does not have nested expressions, it will always result in AST's of a known fixed height. \mathcal{F} , on the other hand, is a context-free language with nested expressions, so the AST depth will vary from input to input. The Composite design pattern lets us work with these AST's of varying depth in a uniform way.

See Figure 5.3 for an example of a sketch of the state of the stack.

We will reuse the \mathcal{F} AST classes from the previous lab.

What about classes like BinaryExpr and UnaryExpr?

Figure 6.2: UML class diagram for Composite Design Pattern.

From [http://en.wikipedia.org/wiki/File:Composite_UML_class_diagram_\(fixed\).svg](http://en.wikipedia.org/wiki/File:Composite_UML_class_diagram_(fixed).svg), where it is released in the public domain (i.e., free for reuse).

6.4 Template Method Design Pattern[†]

A *Template Method* defines an algorithm in outline, but leaves some primitive operations to be defined by the particular datatypes the algorithm will operate on. For example, most sorting algorithms depend on a comparison operation that would be defined differently for strings or integers.

Typically the template algorithm is implemented in an abstract class that declares abstract method signatures for the primitive operations. The concrete subclasses of that abstract class provide definitions of the primitive operations for their respective data values.

The abstract class does not know the name of its subclasses, nor does it perform any explicit type tests.[†] The type tests are implicit in the dynamic dispatch performed when the primitive operations are called.

http://en.wikipedia.org/wiki/Template_method_pattern

http://sourcemaking.com/design-patterns/template_method

<http://www.oodeesign.com/template-method-pattern.html>

[†] An explicit type tests is performed by the instanceof keyword or the getClass() method.

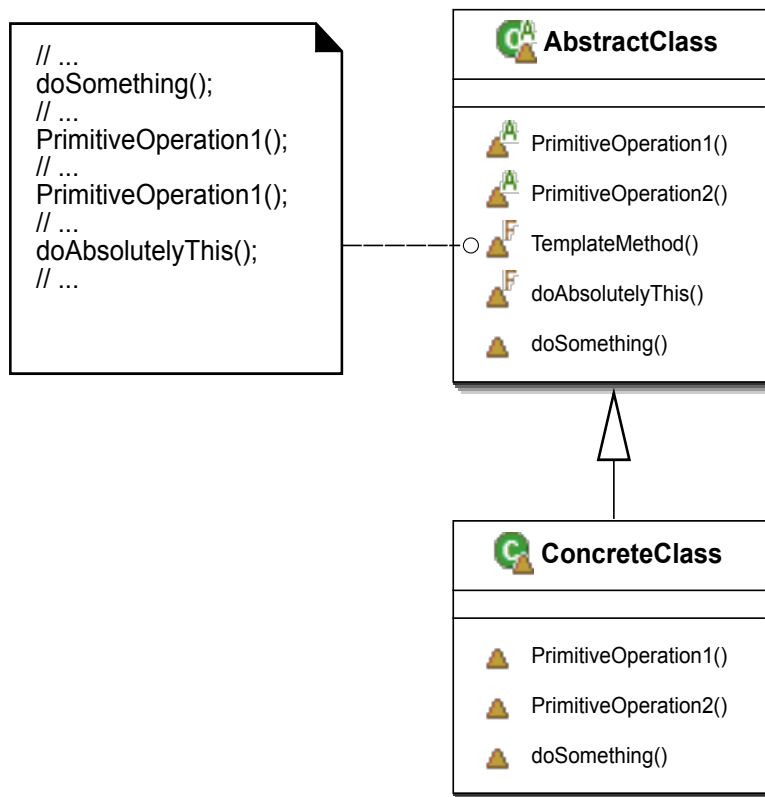


Figure 6.3: Template Design Pattern illustrated by UML Class Diagram. Image from http://en.wikipedia.org/wiki/File:Template_Method_UML.svg. Licensed under GNU Free Documentation Licence.

Expr:

<i>Primitive Op</i>	<i>Implemented By</i>	<i>Template Method</i>
operator()	every subclass of Expr	toString()
simplifyOnce()	NaryExpr, AndExpr, OrExpr, NotExpr	simplify()

NaryExpr:

<i>Primitive Op</i>	<i>Implemented By</i>	<i>Template Method</i>
getAbsorbingElement()	NaryAndExpr, NaryOrExpr	simplifyOnce() helpers
getIdentityElement()	NaryAndExpr, NaryOrExpr	simplifyOnce() helpers
getThatClass()	NaryAndExpr, NaryOrExpr	simplifyOnce() helpers

WHY?

- No superclass (*e.g.*, NaryExpr) should know what its subclasses (*e.g.*, NaryAndExpr, NaryOrExpr) are.
- Should be able to add a new subclass without perturbing parent (*e.g.*, NaryExpr) and siblings (*e.g.*, NaryAndExpr, NaryOrExpr).
- Explicit type tests (*e.g.*, instanceof, getClass()) usually make for fragile code — unless comparing against own type for purpose of some equality-like operation.
- Use dynamic (polymorphic) dispatch to perform type tests in a modular manner.
- Elegant and general implementation of overall transformations.
- Prevent subclasses from making radical departures from general algorithm.
- Reduce code duplication.

6.5 Singleton Design Pattern[†]

The singleton design pattern is used when the number of objects to be instantiated for a given class is independent of the program's input. If the singleton objects are immutable then the primary motivation is to conserve memory by not creating duplicate objects. If the singleton objects are mutable then the primary motivation is to create a set of global variables (which is a controversial decision²).

The main use of the singleton pattern in our labs is the `ConstantExpr` class, which is immutable, so the motivation is to conserve memory by instantiating just one object for `True` and one for `False`, regardless of how many times '1' and '0' appear in the \mathcal{F} programs manipulated by the compiler. Figure 6.4 shows the creation of the two instances of class `ConstantExpr`.

```
public final class ConstantExpr extends Expr {
    public final Boolean b;

    /** The one true instance. To be shared/aliased wherever necessary. */
    public final static ConstantExpr TrueExpr = new ConstantExpr(true);
    /** The one false instance. To be shared/aliased wherever necessary. */
    public final static ConstantExpr FalseExpr = new ConstantExpr(false);
```

While Figure 6.4 shows the creation of two objects that have global names so they can be used by anyone, it does not prevent anyone from creating more objects of class `ConstantExpr`. To do that we make the constructor private, and then provide an alternative method to return a reference to one of the existing objects, as listed in Figure 6.5.

```
/** Private constructor prevents clients from instantiating. */
private ConstantExpr(final Boolean b) { this.b = b; }

/** To be used by clients instead of the constructor.
 * Returns a reference to one of the shared objects. */
public static ConstantExpr make(final Boolean b) {
    if (b) { return TrueExpr; } else { return FalseExpr; }
}
```

http://en.wikipedia.org/wiki/Singleton_pattern

² W. Wulf and M. Shaw. Global variables considered harmful. *ACM SIGPLAN Notices*, 8(2):80–86, Feb. 1973
<http://c2.com/cgi/wiki?GlobalVariablesConsideredHarmful>

Figure 6.4: The only two instances of class `ConstantExpr`

Figure 6.5: Preventing clients from instantiating class `ConstantExpr`

6.6 Evaluation

We evaluate the code that you have committed and pushed before the deadline. The test harnesses are run with the \mathcal{F} programs we have released to you and a secret set of staff \mathcal{F} programs. The weights for the test harnesses are as follows:

	Current	New
TestFParboiledRecognizerAccept	10	5
TestFParboiledRecognizerReject	5	0
TestFParboiledParserBasic	5	0
TestFParboiledParser	40	10
TestFParserComparison	20	5

You do not get any marks for any other parser tests until TestFParboiledParserBasic and TestObjectContractF pass everything.

TestFParboiledRecognizer just runs the recognizer to see if it crashes. TestFParboiledParser checks the following equation:

$$\forall \text{AST} \mid \text{AST.equals}(\text{parse}(\text{prettyprint}(\text{AST})))$$

TestFParserComparison checks that your recursive descent and Parboiled parsers give isomorphic AST's when given the same input file.

Lab 7

Technology Mapping: $\mathcal{F} \rightarrow \text{Graphviz}$

In this lab we will produce gate diagrams of our \mathcal{F} programs. We will do this by translating \mathcal{F} programs into the input language of AT&T GraphViz. Graphviz is an automatic graph layout tool: it reads (one dimensional) textual descriptions of graphs and produces (two dimensional) pictures of those graphs. Graphviz, which is primarily a visualization tool, attempts to place nodes and edges to minimize edge crossings. Tools that are intended to do digital circuit layout optimize the placement of components based on other criteria, but the fundamental idea is the same: transform a (one dimensional) description of the logic into a (two dimensional) plan that can be visualized or fabricated.

Figure 7.1 lists a sample Graphviz input file and its rendered output for the simple \mathcal{F} program $X \Leftarrow A \text{ or } B$. The input pins A and B are on the left side of the diagram and the output pin X is on the right hand side of the diagram.

```
digraph g {
    // header
    rankdir=LR;
    margin=0.01;
    node [shape="plaintext"];
    edge [arrowhead="diamond"];
    // circuit
    or12 [label="or12", image="../../gates/or_noleads.png"];
    var0[label="x"];
    var1[label="a"];
    var2[label="b"];
    var1 -> or12 ;
    var2 -> or12 ;
    or12 -> var0 ;
}
```



Compiler Concepts: common subexpression elimination

Programming Concepts: hash structures, iteration order, object identity, non-determinism, Visitor design pattern, tree traversals: preorder, postorder, inorder

Files:

ece351.f.techmapper.TechnologyMapper

Libraries:

ece351.f.techmapper.GraphvizToF

ece351.f.analysis.ExtractAllExprs

ece351.common.visitor.PostOrderExprVisitor

Tests:

ece351.f.techmapper.TestTechnologyMapper

<http://graphviz.org>

Graphviz is a widely used graph layout tool. Circuit layout is often done by specialized tools, but the basic idea is the same as Graphviz: algorithmic placement of interconnected components. What differs is the criteria used for placement and the ways in which things are connected. Graphical layout tools such as Graphviz try to minimize edge crossings, for example, which is important for people looking at pictures but may be less relevant for circuit boards. Graphviz also draws nice Bezier curve lines that are pleasing to look at, whereas circuit boards are typically laid out with orthogonal lines. Technology mapping and circuit layout are studied in ECE647.

Figure 7.1: Example Graphviz input file and rendered output for \mathcal{F} program $X \Leftarrow A \text{ or } B$;

The input file in Figure 7.1 contains a header section that will be common to all of the Graphviz files that we generate. This header says that the graph should be rendered left to right (instead of the default top-down), that the whitespace margin around the diagram should be 0.01 inches, and that the default look for nodes and edges should be plain.

The main thing to notice in the lines of the input file in Figure 7.1 that describe the formula is that visual nodes are created for both the pins (A, B, X) and the gates (there is a single OR gate in this example).

7.1 *Interpreters vs. Compilers*[†]

An *programming language interpreter* is a special kind of compiler that executes a program directly, instead of translating the program to another language that is then executed. Executing a program through a language interpreter is usually slower than executing the compiled form of that program. It can take less engineering effort to implement a language interpreter than a compiler, so language interpreters are most commonly used when engineering time is more valuable than machine time. Compilers are used when the converse is true: *i.e.*, it is worth the engineering effort to make the code run faster.

7.2 *Introducing the Interpreter Design Pattern*[†]

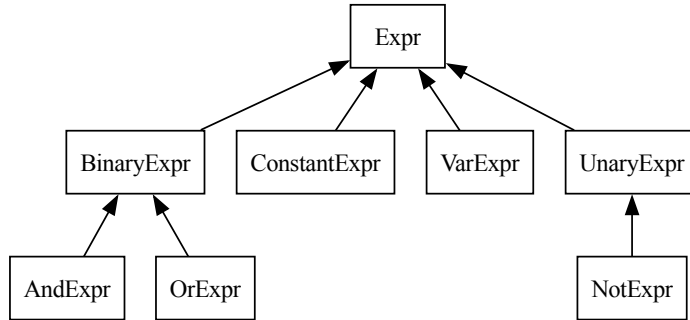
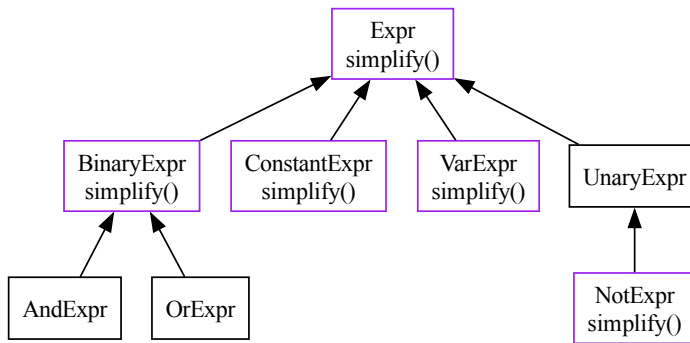
The *Interpreter design pattern*¹ names a particular way to structure code, that we have been using in previous labs but have not yet named. Programming language interpreters are often implemented using the interpreter design pattern, which is how the design pattern got its name. When learning about both concepts this name collision can be a bit confusing.

Figure 7.2 shows a UML class diagram for a subset of the \mathcal{F} expression class hierarchy that we've been working with recently.

In §4 we wrote a simplifier/optimizer for \mathcal{F} programs. How did we structure that code? We added a method called `simplify()` to almost every class in this hierarchy, similar to what is depicted in Figure 7.3. This is a non-modular change: the simplifier functionality is scattered across the Expr class hierarchy.

Is it possible to structure this code in another way? Could we add the simplifier functionality to our compiler in a modular way?

[†] E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1995

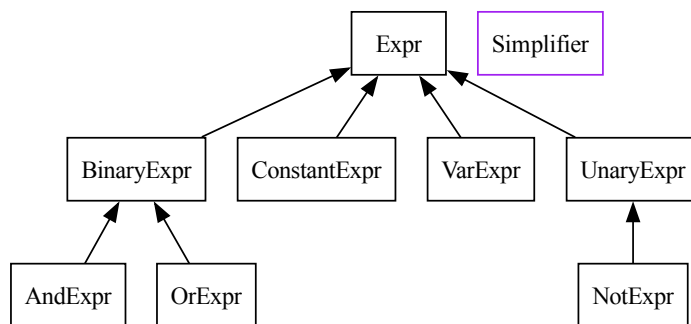
Figure 7.2: A UML class diagram for a subset of \mathcal{F} expression class hierarchyFigure 7.3: A UML class diagram for a subset of \mathcal{F} expression class hierarchy, showing the simplifier functionality scattered across the hierarchy by use of the Interpreter design pattern. Classes modified to add the simplifier functionality are highlighted in colour.

AST Class	Operation		
	simplify	pretty-printing	conversion to gates
AndExpr	AndExpr (BinaryExpr)	AndExpr (BinaryExpr)	AndExpr (BinaryExpr)
OrExpr	OrExpr (BinaryExpr)	OrExpr (BinaryExpr)	OrExpr (BinaryExpr)
VarExpr	VarExpr	VarExpr	VarExpr
ConstantExpr	ConstantExpr	ConstantExpr	ConstantExpr
NotExpr	NotExpr	NotExpr	NotExpr
NaryAndExpr	NaryAndExpr (NaryExpr)	NaryAndExpr (NaryExpr)	NaryAndExpr (NaryExpr)
NaryOrExpr	NaryOrExpr (NaryExpr)	NaryOrExpr (NaryExpr)	NaryOrExpr (NaryExpr)

Figure 7.4: Matrix view of Interpreter Design Pattern. Rows for AST leaf classes. Columns for operations to be performed on the AST. Cell values indicate where the code for that class/op pair is located. In some cases we place the operation in an abstract superclass, e.g., BinaryExpr, rather than directly in the leaf class. This location is indicated with the concrete leaf class name followed by the abstract superclass name in parenthesis. Notice that the pattern here is that the same value is repeated across each row.

7.3 Introducing the Visitor Design Pattern[†]

The *Visitor* design pattern² is an alternative way to structure code from the interpreter design pattern. The Visitor pattern is also widely used in compilers, and is how we will structure our code in this lab and future labs. If we had used the Visitor design pattern for §4 then we could have added the simplifier functionality in a more modular manner, as depicted in Figure 7.5.



Tiger Book §4.3

² E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1995

Also widely documented on the web.

Figure 7.5: A UML class diagram for a subset of \mathcal{F} expression class hierarchy, showing the simplifier functionality made modular by use of the Visitor design pattern.

Well, it's not quite as simple as Figure 7.5 might lead you to believe: there is some infrastructure that we need to add to make this work properly. Figure 7.6 illustrates that we need to add an `accept(Visitor)` method to each concrete class.³ Similarly, we add a `visit()` for each concrete class to the `Simplifier` class. These `accept(Visitor)` methods can be reused by any Visitor, whether it is the simplifier or the technology mapper that we will develop in this lab.

With the use of the Visitor design pattern we can add new operations (e.g., simplification, technology mapping, etc.) to our complex data structure (AST) in a modular manner, without having to directly modify the AST classes.

The Visitor design pattern involves two main methods: *visit* and *accept*. The *visit* methods are written on the Visitor class and the *accept* methods are written on the AST classes. The abstract superclass for the Visitors that we will write is listed in Figure 7.8.

Notice that our visit methods return an object of type `Expr`. This feature allows our Visitors to rewrite the AST as they visit it, which is necessary for transformations such as the simplifier we wrote previously. The Visitor we write in this lab does not rewrite the AST, and so all of its visit methods will simply return the AST node that they are visiting (i.e., return `e`).

³ Notice that in our \mathcal{F} Expr class hierarchy the only classes that can be directly instantiated (i.e., are concrete) are the leaves of the hierarchy. This is sometimes known as the *abstract superclass rule* (i.e., all super-classes should be abstract), and is a good design guideline for you to follow.

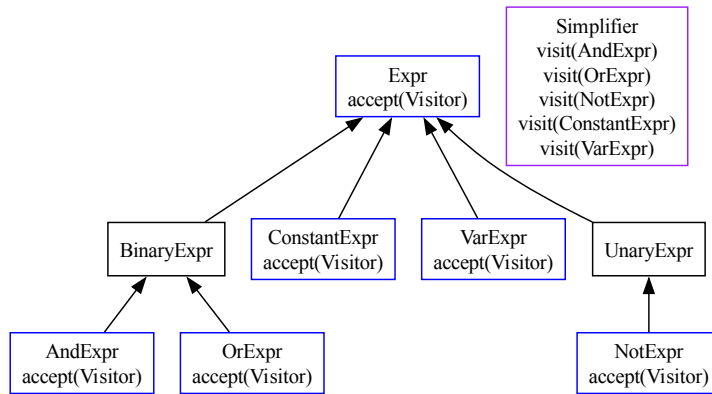


Figure 7.6: A UML class diagram for a subset of \mathcal{F} expression class hierarchy, showing the simplifier functionality made modular by use of the Visitor design pattern.

AST Class	Operation		
	simplify	pretty-printing	conversion to gates
AndExpr	Simplifier	PPrinter	TechMapper
OrExpr	Simplifier	PPrinter	TechMapper
VarExpr	Simplifier	PPrinter	TechMapper
ConstantExpr	Simplifier	PPrinter	TechMapper
NotExpr	Simplifier	PPrinter	TechMapper
NaryAndExpr	Simplifier	PPrinter	TechMapper
NaryOrExpr	Simplifier	PPrinter	TechMapper

Figure 7.7: Matrix view of the Visitor Design Pattern. Rows for AST leaf classes. Columns for operations to be performed on the AST. Cell values indicate where the code for that class/op pair is located. Notice that the pattern here is that the values are repeated down each column, rather than across each row as we saw above in Figure 7.4 for the Interpreter Design Pattern.

In the actual code for our labs we followed the Interpreter Design Pattern for the simplify and pretty-printing operations, so the class names Simplifier and PPrinter are fictitious. We are developing the TechnologyMapper visitor in this lab.

```

public abstract Expr visitConstant(ConstantExpr e);
public abstract Expr visitVar(VarExpr e);
public abstract Expr visitNot(NotExpr e);
public abstract Expr visitAnd(AndExpr e);
public abstract Expr visitOr(OrExpr e);
public abstract Expr visitNaryAnd(NaryAndExpr e);
public abstract Expr visitNaryOr(NaryOrExpr e);
  
```

Figure 7.8: Abstract super class for Visitors for \mathcal{F} expression ASTs. There is a visit method for each concrete \mathcal{F} AST class.

To the FExpr class we add the signature for the accept method, which is then implemented in each concrete AST class as a call to a visit method, as shown in Figure 7.9.

```
public abstract Expr accept(final ExprVisitor exprVisitor);

public Expr accept(final ExprVisitor v) { return v.visitAnd(this); }
```

Together the visit and accept methods implement what is known as *double dispatch*: i.e., select which method to execute based on the polymorphic type of two objects. Languages like Java, C++, C#, and Objective C are all *single dispatch* languages: the target of a polymorphic call is determined just by the type of the receiver object. CLOS and Dylan are examples of *multiple-dispatch* languages, where the target of a polymorphic call is determined by the runtime types of all of arguments.

One of the main points of variation in how the Visitor pattern is implemented is where the traversal code goes: in the visit methods? an iterator? somewhere else? All of these options are used in practice. We have decided to put the traversal code in a set of *traverse* methods in the Visitor classes (Figures 7.10 and 7.11).

We can write Visitors that perform a number of useful tasks for this lab. For example, Figure 7.12 lists a Visitor that builds a set of all the nodes in an \mathcal{F} expression AST.

Figure 7.9: Signature and implementation of accept method. The signature belongs in the Expr class and the implementation is put in each concrete subclass of Expr.

The point of a Visitor is to *traverse*, or ‘walk over’, the nodes in an AST.

7.3.1 Is Visitor always better than Interpreter?[†]

From the discussion above you might think that it is always better to use the Visitor pattern than the Interpreter pattern. This is not correct. Which pattern to choose depends on what changes you anticipate happening to the software in the future.

Consider this: we have a complex structure (*e.g.*, an AST), and a set of complex operations that we would like to apply to that structure (*e.g.*, simplification, technology mapping, *etc.*). In ECE250 you studied some relatively simple structures, such as lists, and some relatively simple operations, such as searching or sorting. Those structures were simple enough that they could be defined in one or two classes, and those operations were simple enough that they could be defined in one or two methods. Now we have a complex structure that is defined across a dozen or more classes, and operations that might similarly be defined across a dozen or more methods.

The choice of how we organize this code (Visitor or Interpreter) has consequences, and those consequences are in how difficult the code will be to change in the future.

If we expect the grammar of our input language (*e.g.*, \mathcal{W} or \mathcal{F}) to be stable, and if we expect to add more transformations (operations) to the AST in future labs, then it is better to use the Visitor pattern, because the Visitor pattern allows us to add new transformations in a modular fashion.

If, on the other hand, we expect the grammar of our input language to change and the set of transformations we might want to perform is small and known *a priori*, then it is better to use Interpreter. Interpreter lets us change the AST class hierarchy in a modular fashion, which is what we would need to do if the grammar of the input language were to change significantly.

What if we anticipate both kinds of change will happen in the future? Then we have a problem, because there is no widely accepted programming language that lets us structure our code to facilitate both kinds of change. This is known as the *Expression Problem* in programming language design. The Expression Problem was named by Wadler in the late nineties⁴, although the idea goes back at least as far as Reynolds work in the mid seventies⁵. More recent research papers have proposed some solutions in Java⁶ or Scala⁷, but none of these proposals is widely accepted and all have drawbacks.

Almost all software evolves or dies. Donald Knuth's T_EX is a notable counter-example: it does not grow new features, it does not have significant bugs, it was conceived complete. Consequently, the version number for T_EX does not increment, it converges: as the (minor) bugs are removed the version number approximates π by another decimal place. This reflects convergence to an original vision, with the acknowledgement that perfection will never truly be achieved within the finite time available to us here on earth.

These references below are for your broader educational enrichment. You will not be tested on them specifically. You are expected to understand the trade-offs between the Visitor and Interpreter design patterns.

⁴ P. Wadler. The expression problem, Nov. 1998. Email to Java Generics list

⁵ J. C. Reynolds. User-defined types and procedural data as complementary approaches to data abstraction. In S. A. Schuman, editor, *New Directions in Algorithmic Languages: IFIP Working Group 2.1 on Algol*. INRIA, 1975

⁶ M. Torgersen. The Expression Problem Revisited: Four new solutions using generics. In M. Odersky, editor, *Proc. 18th ECOOP*, volume 3344 of LNCS, Oslo, Norway, June 2004. Springer-Verlag

⁷ M. Zenger and M. Odersky. Independently extensible solutions to the expression problem. In *Proc. 12th Workshop on Foundations of Object-Oriented Languages*, 2005

Figure 7.10: The traverse methods of ExprVisitor

```

public final Expr traverseExpr(final Expr e) {
    if (e instanceof NaryExpr) {
        return traverseNaryExpr( (NaryExpr) e );
    } else if (e instanceof BinaryExpr) {
        return traverseBinaryExpr( (BinaryExpr) e );
    } else if (e instanceof UnaryExpr) {
        return traverseUnaryExpr( (UnaryExpr) e );
    } else {
        return e.accept(this);
    }
}

public abstract Expr traverseNaryExpr(final NaryExpr e);
public abstract Expr traverseBinaryExpr(final BinaryExpr e);
public abstract Expr traverseUnaryExpr(final UnaryExpr e);

/**
 * Visit/rewrite all of the exprs in this FProgram.
 * @param p input FProgram
 * @return a new FProgram with changes applied
 */
public FProgram traverseFProgram(final FProgram p) {
    FProgram result = new FProgram();
    for (final AssignmentStatement astmt : p.formulas) {
        result = result.append(traverseAssignmentStatement(astmt));
    }
    return result;
}

/**
 * Visit/rewrite the expr in this AssignmentStatement
 * @param astmt the AssignmentStatement to be visited/rewritten
 * @return a new AssignmentStatement with changes applied
 */
public AssignmentStatement traverseAssignmentStatement(final AssignmentStatement astmt) {
    final Expr e = traverseExpr(astmt.expr);
    if (e == astmt.expr) {
        // no change
        return astmt;
    } else {
        // rewrite occurred
        return astmt.varyExpr(e);
    }
}
}

```



```

/**
 * This visitor rewrites the AST from the bottom up.
 * Optimized to only create new parent nodes if children have changed.
 */
public abstract class PostOrderExprVisitor extends ExprVisitor {

    @Override
    public final Expr traverseUnaryExpr(UnaryExpr u) {
        // child first
        final Expr child = traverseExpr(u.expr);
        // only rewrite if something has changed
        if (child != u.expr) {
            u = u.newUnaryExpr(child);
        }
        // now parent
        return u.accept(this);
    }

    @Override
    public final Expr traverseBinaryExpr(BinaryExpr b) {
        // children first
        final Expr left = traverseExpr(b.left);
        final Expr right = traverseExpr(b.right);
        // only rewrite if something has changed
        if (left != b.left || right != b.right) {
            b = b.newBinaryExpr(left, right);
        }
        // now parent
        return b.accept(this);
    }

    @Override
    public final Expr traverseNaryExpr(NaryExpr e) {
        // children first
        ImmutableList<Expr> children = ImmutableList.of();
        boolean change = false;
        for (final Expr c1 : e.children) {
            final Expr c2 = traverseExpr(c1);
            children = children.append(c2);
            if (c2 != c1) { change = true; }
        }
        // only rewrite if something changed
        if (change) {
            e = e.newNaryExpr(children);
        }
        // now parent
        return e.accept(this);
    }
}

```

Figure 7.11: Implementation of Post-OrderExprVisitor. Any Visitors that extend this class will visit the nodes of an \mathcal{F} expression AST in post-order (i.e., parents after children).

Figure 7.12: Implementation of ExtractAllExprs

```

/**
 * Returns a set of all Expr objects in a given FProgram or AssignmentStatement.
 * The result is returned in an IdentityHashSet, which defines object identity
 * by memory address. A regular HashSet defines object identity by the equals()
 * method. Consider two VarExpr objects, X1 and X2, both naming variable X. If
 * we tried to add both of these to a regular HashSet the second add would fail
 * because the regular HashSet would say that it already held a VarExpr for X.
 * The IdentityHashSet, on the other hand, will hold both X1 and X2.
 */
public final class ExtractAllExprs extends PostOrderExprVisitor {

    private final IdentityHashSet<Expr> exprs = new IdentityHashSet<Expr>();

    private ExtractAllExprs(final Expr e) { traverseExpr(e); }

    public static IdentityHashSet<Expr> allExprs(final AssignmentStatement f) {
        final ExtractAllExprs cae = new ExtractAllExprs(f.expr);
        return cae.exprs;
    }

    public static IdentityHashSet<Expr> allExprs(final FProgram p) {
        final IdentityHashSet<Expr> allExprs = new IdentityHashSet<Expr>();
        for (final AssignmentStatement f : p.formulas) {
            allExprs.add(f.outputVar);
            allExprs.addAll(ExtractAllExprs.allExprs(f));
        }
        return allExprs;
    }

    @Override public Expr visitConstant(ConstantExpr e) { exprs.add(e); return e; }
    @Override public Expr visitVar(VarExpr e) { exprs.add(e); return e; }
    @Override public Expr visitNot(NotExpr e) { exprs.add(e); return e; }
    @Override public Expr visitAnd(AndExpr e) { exprs.add(e); return e; }
    @Override public Expr visitOr(OrExpr e) { exprs.add(e); return e; }
    @Override public Expr visitXOr(XOrExpr e) { exprs.add(e); return e; }
    @Override public Expr visitNAnd(NAndExpr e) { exprs.add(e); return e; }
    @Override public Expr visitNOr(NOrExpr e) { exprs.add(e); return e; }
    @Override public Expr visitXNor(XNorExpr e) { exprs.add(e); return e; }
    @Override public Expr visitEqual(EqualExpr e) { exprs.add(e); return e; }
    @Override public Expr visitNaryAnd(NaryAndExpr e) { exprs.add(e); return e; }
    @Override public Expr visitNaryOr(NaryOrExpr e) { exprs.add(e); return e; }
}

```

7.4 Hash Structures, Iteration Order, and Object Identity[†]

When you iterate over the elements in a List you get them in the order that they were added to the List. When you iterate over the elements in a TreeSet you get them sorted lexicographically. When you iterate over the elements in a HashSet or HashMap, what order do you get them in? Unspecified, unknown, and non-deterministic: the order could change the next time you iterate, and will likely change the next time the program executes.

TreeSet, List, HashSet, and HashMap are all part of the standard JDK Collections classes.

List list = new ArrayList();	SortedSet tset = new TreeSet();	Set hset = new HashSet();
list.add(3);	tset.add(3);	hset.add(3);
list.add(1);	tset.add(1);	hset.add(1);
list.add(2);	tset.add(2);	hset.add(2);
System.out.println(list);	System.out.println(tset);	System.out.println(hset);

Figure 7.13: Iteration order for different data structures

Why might the iteration order change with hash structures? Because the slot into which an element gets stored in a hash structure is a function of that element's hash value and *the size of the hash table*. As more elements are added to a hash structure then it will resize itself and rehash all of its existing elements and they'll go into new slots in the new (larger) table. If the same data value always produces the same hash value and the table never grows then it is possible to get deterministic iteration order from a hash structure — although that order will still be nonsense, it will be deterministically repeatable nonsense. But these assumptions often do not hold. For example, if a class does not implement the equals() and hashCode() methods then its memory address is used as its hashCode(). The next execution of the program is highly likely to put that same data value at a different memory address.

Iterating over the elements in a hash structure is one of the most common ways of unintentionally introducing non-determinism into a Java program. Non-determinism makes testing and debugging difficult because each execution of the program behaves differently. So unless there is some benefit to the non-determinism it should be avoided.

The JDK Collections classes provide two hash structures with deterministic iteration ordering: LinkedHashMap and LinkedHashMap. These structures also maintain an internal linked list that records the order in which elements were added to the hash structure. You should usually choose LinkedHashMap, LinkedHashMap, TreeMap, or TreeSet over HashMap and HashSet. The linked structures give elements in their insertion order (as a list would), whereas the tree

What benefit could there be to non-determinism? Not much directly. But non-determinism is often a consequence of parallel and distributed systems. In these circumstances we sometimes choose to tolerate some non-determinism for the sake of performance — but we still try to control or eliminate some of the non-determinism using mechanisms like locks or database engines.

structures give you elements in alphabetical order (assuming that there is some alphabetical ordering for the elements).

You could safely use `HashMap` and `HashSet` without introducing non-determinism into your program *if you never iterate over their elements*. It's hard to keep that promise though. Once you've got a data structure you might want to print out its values, or pass it in to some third party code, *etc.* So it's better to just use a structure with a deterministic iteration ordering.

The skeleton code for this lab makes use of two other hash structures: `IdentityHashMap` and `IdentityHashSet`. What distinguishes these structures from the common structures?

The definition of a set is that it does not contain duplicate elements. How is 'duplicate' determined? We work with four different definitions of 'duplicate' in these labs:

x == y *x* and *y* have the same physical memory address.

x.equals(y) any computation that uses *x* or *y* will have no observable differences.

x.isomorphic(y) *x* and *y* might have some minor structural differences, but are essentially the same.

x.equivalent(y) *x* and *y* are semantically equivalent and might not have any structural similarities.

```
Set hset = new HashSet();
hset.add(new VarExpr("Y"));
hset.add(new VarExpr("X"));
hset.add(new VarExpr("X"));
System.out.println(hset);
```

```
Set lhset = new LinkedHashSet();
lhset.add(new VarExpr("Y"));
lhset.add(new VarExpr("X"));
lhset.add(new VarExpr("X"));
System.out.println(lhset);
```

```
Set iset = new IdentityHashSet();
iset.add(new VarExpr("Y"));
iset.add(new VarExpr("X"));
iset.add(new VarExpr("X"));
System.out.println(iset);
```

The common structures (`HashSet`, `TreeSet`, *etc.*) use the `equals()` method to determine duplicates, whereas the `IdentityHashSet` and `IdentityHashMap` use the memory address (`==`) to determine duplicates. In this lab we want to ensure that our substitution table contains an entry for every single `FExpr` object (physical memory address), so we use `IdentityHashSet` and `IdentityHashMap`. Notice that `ExtractAllExprs` also returns an `IdentityHashSet`.

The skeleton code is careful about its use of `IdentityHashSet/Map`, `LinkedHashSet/Map`, and `TreeSet/Map`. You should think about the concept of duplicate and the iteration ordering of each data structure used in this lab, including your temporary variables. The `GraphvizToF` converter requires that the edges are printed according

`IdentityHashMap` is from the JDK. `IdentityHashSet` is from the open-source project named `Kodkod`. There are a number of open source projects that implement an `IdentityHashSet` due to its omission in the JDK. See Java bug report #4479578.

Figure 7.14: Object identity for different data structures

to a post-order traversal of the AST, so that it can reconstruct the AST bottom-up.

7.5 Non-determinism: Angelic, Demonic, and Arbitrary[†]

We previously saw non-determinism when studying finite automata. In that case we focused on *angelic* non-determinism: the automata would ‘magically’ (non-deterministically) choose the right path that will lead to success (acceptance).

In this lab we’re seeing *arbitrary* non-determinism: the machine choose some path arbitrarily, and that path may lead to a good result or a bad result.

As engineers, we want to avoid non-determinism. It makes testing unnecessarily difficult. We transform non-deterministic machines into deterministic ones. We are careful to avoid unnecessary sources of non-determinism in our programs, such as iterating over hash structures or unregulated concurrency.

In theory they also consider *demonic* non-determinism, where the machine always chooses a path that leads to the worst possible outcome.

7.6 Translate one formula at a time

Flesh out the skeleton synthesizer in the provided TechnologyMapper class. The TechnologyMapper extends one of the Visitor classes, and in its visit methods it will create edges from the child nodes in the AST to their parents.

The TechnologyMapper class contains a field called substitutions that maps FExprs to FExprs. For now populate this data structure by mapping every FExpr AST node to itself. Populating this data structure in a more sophisticated manner is the next task.

7.7 Common Subexpression Elimination

An \mathcal{F} program may contain common subexpressions. A smart synthesizer will synthesize shared hardware for these common subexpressions. Consider the following simple \mathcal{F} program where the subexpression A or B appears twice:

```
X <= D and (A or B);
Y <= E and (A or B);
```

Figure 7.15 shows the circuit produced by a synthesizer that does not perform common subexpression elimination: there are two OR gates that both compute A or B. A more sophisticated synthesizer that does perform common subexpression elimination would synthesize a circuit like the one in Figure 7.16 in which there is only a single OR gate. It is also possible to have common subexpressions within a single formula, for example: $Z <= (A \text{ or } B) \text{ and } \neg(A \text{ or } B)$;

We do not have to worry about variables changing values in \mathcal{F} because all expressions in \mathcal{F} are *referentially transparent*: i.e., they always have the same value no matter where they occur in the program. Referential transparency is a common property of *functional* programming languages such as Scheme, Lisp, Haskell, ML, and \mathcal{F} . Consider the following program fragment written in an *imperative* language:

```
a = 1;
b = 2;
x = a + b; // = 3
a = 3;
y = a + b; // = 5
```

The subexpression $a + b$ occurs twice in this fragment but cannot be eliminated because the value of a has changed. A *dataflow analysis* is required to determine if the values of the variables have changed.

7.8 Designing a Common Subexpression Eliminator for \mathcal{F}

The key question that the eliminator needs to answer is, for each sub-expression, whether that sub-expression should be used, or whether it should be eliminated in favour of another sub-expression.

What measure of ‘sameness’ should be used for this decision? What concept of identity should we use? There are four options: physical memory address (`==`), *equals*, *isomorphic*, and *equivalent*.

Memory address is (hopefully obviously) not a useful choice: if our input ASTs are already sharing physical objects for common subexpressions then the elimination has already been done. *Equivalence* is, in some sense, the ideal measure. But computing *equivalence* for \mathcal{F} formulas is an NP-complete problem, and we want to stick with polynomial complexity. That leaves *equals* and *isomorphic*. The *equals* method is perhaps easier to work with, since it is already used by the standard Java collections classes. *Isomorphic* will give us better results, at the cost of slightly increased design complexity — but no significant change in computational complexity. Possible designs include:

- a. Just use *equals* via the standard Java collections classes. Easy to design and implement. Results not as good as they could be if *isomorphic* were used.
- b. Convert all sub-expressions to a canonical form first, then use *equals* as above. We used this kind of approach in LAB4: NaryExprs have a canonical form because they sort their children lexicographically. We will never have an NaryOrExpr like B or A, because it will be sorted to A or B. Once we have all sub-expressions in a canonical form, then we can use *equals* as above.
- c. Implement a customized data structure based on *isomorphic*. If *hashCode* is consistent with *isomorphic* then it can be used as a first step in the search for substitutable sub-expression objects.
- d. Write some wrapper code that compares each sub-expression to every other one using *isomorphic*, and builds a substitution table based on physical memory address. This is not necessarily the most elegant design, but it is one that is described in some detail below.

These alternatives came out in discussion with the class of the summer of 2017.

There is another design decision about whether the ASTs get rebuilt with the substitutions, or whether the substitution table is just used in the printing process.

7.8.1 An n^2 design using isomorphic and memory address

This is the solution in the staff code, which is hinted at in the skeleton. You are welcome to delete those parts of the skeleton go with another design. Some alternatives are described above.

- Build up a table of FExpr substitutions. Suppose our \mathcal{F} program has three occurrences of A or B: $(A \text{ or } B)_1$, $(A \text{ or } B)_2$, and $(A \text{ or } B)_3$. These three expressions are in the same *equivalence class* according to our `isomorphic()` method. We select one representative from this group, say $(A \text{ or } B)_2$, and use that to represent all of the isomorphic expressions. So our substitution table will contain the tuples $\langle (A \text{ or } B)_1, (A \text{ or } B)_2 \rangle$, $\langle (A \text{ or } B)_2, (A \text{ or } B)_2 \rangle$, $\langle (A \text{ or } B)_3, (A \text{ or } B)_2 \rangle$. In other words, we'll use $(A \text{ or } B)_2$ in place of $(A \text{ or } B)_1$ and $(A \text{ or } B)_3$.

This substitution table can be built in the following way. First, compare every FExpr in the program with every other FExpr with the `isomorphic()` method to discover the equivalence classes. For each equivalence class pick a representative (doesn't matter which one), and then make an entry in the substitution table mapping each FExpr in the equivalence class to the representative.

- Visit the AST and produce edges from children FExprs to parent FExprs using the substitution table.

Why do we use the `isomorphic()` method here instead of the `equivalent()` method? Because the `isomorphic()` method runs in polynomial time whereas the `equivalent()` method runs in exponential time. Our common subexpression elimination only makes engineering sense if it costs polynomial time. If we were to spend exponential time then we could do more sophisticated circuit minimization, such as computing the prime implicants with the Quine-McCluskey algorithm.

See the provided utility methods in the skeleton `TechnologyMapper` class.




Figure 7.15: Synthesized circuit *without* common subexpression elimination for \mathcal{F} program $X \Leftarrow D$ and $(A \text{ or } B)$; $Y \Leftarrow E$ and $(A \text{ or } B)$. There are two gates synthesized that both compute A or B.




Figure 7.16: Synthesized circuit *with* common subexpression elimination for \mathcal{F} program $X \Leftarrow D$ and $(A \text{ or } B)$; $Y \Leftarrow E$ and $(A \text{ or } B)$. There is only one gate that computes the value A or B, and the output of this gate is fed to two places.

7.9 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

The first testing equation for this lab is as follows. Let f be the AST produced by your \mathcal{F} parser.

```
f.equivalent(GraphVizToF(TechnologyMapper(simplify(f))))
```

In other words, take the output of your \mathcal{F} parser, simplify it, convert it to dot (this is the part that you are doing, the technology mapper), then convert that output back to an \mathcal{F} program (we provide this code for you), and compare this result to the original AST produced by the parser.

The second testing equation for this lab compares your output with the staff output (dot file). Let s name the staff dot file, and f names the AST produced by your \mathcal{F} parser.

```
simplify(GraphVizToF(s)).equivalent(GraphVizToF(TechnologyMapper(simplify(f))))
```

The two equations above just check that you produced an equivalent circuit, but they do not measure the number of gates used in that circuit. A third component of the evaluation is comparing the number of gates in your circuit versus the staff circuit.

TestTechnologyMapper	Shared	Secret
correctness	50	20
gate count	20	10

The baseline gate count marks are 16/20 and 8/10. If you produce fewer gates than the staff solution you bump up to 20/20 and 10/10. If you produce more gates than the staff solution then for each extra gate we subtract one point from the baseline. Some strategies that might produce fewer gates but probably take too much time to implement include:

- Using DeMorgan's Laws or other algebraic identities.
- Implementing the Quine-McCluskey or ESPRESSO algorithms to compute the minimal form of the circuit (a much more advanced version of our simplifier).
- Translating our FExprs to Binary Decision Diagrams (BDDs).⁸ BDDs are used in practice for functional technology mapping and circuit equivalence checking.

MORAL: Reducing the gate count is an NP-complete problem. Our simplifier and common subexpression eliminator are both polynomial. No combination of polynomial algorithms is ever going to be a perfect solution to an NP-complete problem. The ESPRESSO algorithm is an example of a very good polynomial approximation for this NP-complete problem.

The Quine-McCluskey algorithm is not used in practice because it is exponential/NP-complete, which means it is too expensive for most real circuits. We can afford the price for the small circuits considered in this course. The ESPRESSO algorithm is efficient enough to be used in practice but is not guaranteed to find the global minimum.
http://en.wikipedia.org/wiki/Espresso_heuristic_logic_minimizer
http://en.wikipedia.org/wiki/Quine-McCluskey_algorithm

⁸ R. E. Bryant. Graph-based algorithms for boolean function manipulation. *IEEE Transactions on Computers*, C-35(8):677–691, Aug. 1986

Lab 8

Simulation: $\mathcal{F} \rightarrow \text{Java}$

Consider the following \mathcal{F} program: $X \leftarrow A \text{ or } B$. A simulator for this \mathcal{F} program would read an input \mathcal{W} file describing the values of A and B at each time step, and would write an output \mathcal{W} file describing the values of X at each time step. Such a simulator is listed in Figure 8.1. Your task is to write a program that, given an input \mathcal{F} program, will produce a Java program like the one in Figure 8.1. In other words, you are going to write a simulator generator. The generated simulators will perform the following steps:

- Read the input \mathcal{W} program.
- Allocate storage for the output \mathcal{W} program.
- Loop over the time steps and compute the outputs.
- Write the output \mathcal{W} program.

The generated simulator will have a method to compute the value of each output pin. In our example \mathcal{F} program listed above the output pin is X , and so the generated simulator in Figure 8.1 has a method named X . The body of an output pin method is generated by performing a *pre-order* traversal of the corresponding \mathcal{F} expression AST. \mathcal{F} expressions are written with operators *in-order*: that is, the operators appear between the variables. For example, in the \mathcal{F} program we have the in-order expression $A + B$, while in the Java translation we have the pre-order expression $\text{or}(A, B)$.

8.1 Name Collision/Capture[†]

When we generate code in another language we need to be careful that the variable names we generate based on our input program are legal in the target language, and that they do not collide/capture an already existing name. All legal \mathcal{W} and \mathcal{F} variable names are also legal Java variable names, so we don't need to worry about the first point here. On the second point, notice that the generated Java variable names in Figure 8.1 are prefixed with 'in_' or 'out_', and that none of the boilerplate names have such prefixes.

Files:
ece351.f.simgen.SimulatorGenerator
Libraries:
ece351.f.analysis.DetermineInputVars
Tests:
ece351.f.simgen.TestSimulatorGenerator

An alternative to generating a simulator would be to write an \mathcal{F} *interpreter*. What we are doing here is writing an \mathcal{F} *compiler*. We call it a 'simulator generator' because the term 'compile' is not clear in the context of \mathcal{F} : in the last lab we 'compiled' \mathcal{F} to a circuit; here we 'compile' it to a simulator.

An *interpreter* evaluates a program with an input, but does not transform the program to another language. A compiler, by contrast, is almost the opposite: it translates a program to another language, but does not evaluate that program on any input.

It is usually the case that it takes less effort to write an interpreter, but that the target program's runtime is faster if it is compiled first.

A name capture problem occurred in some of the utility code. The `FProgram.equivalent()` method checks that the equivalence of two \mathcal{F} programs by translating them to SAT. For convenience, that process first translates the \mathcal{F} program to Alloy, and then the Alloy is translated to SAT. An older version of this translation did not take care to avoid name capture, so if the input \mathcal{F} program had variables corresponding to Alloy keywords (e.g., 'this', 'int') then the Alloy to SAT translation would fail.

```

import java.util.*;
import ece351.w.ast.*;
import ece351.w.parboiled.*;
import static ece351.util.Boolean351.*;
import ece351.util.CommandLine;
import java.io.File;
import java.io.FileWriter;
import java.io.StringWriter;
import java.io.PrintWriter;
import java.io.IOException;
import ece351.util.Debug;

public final class Simulator_ex11 {
    public static void main(final String[] args) {
        final String s = File.separator;
        // read the input F program
        // write the output
        // read input WProgram
        final CommandLine cmd = new CommandLine(args);
        final String input = cmd.readInputSpec();
        final WProgram wprogram = WParboiledParser.parse(input);
        // construct storage for output
        final Map<String,StringBuilder> output = new LinkedHashMap<String,StringBuilder>();
        output.put("x", new StringBuilder());
        // loop over each time step
        final int timeCount = wprogram.timeCount();
        for (int time = 0; time < timeCount; time++) {
            // values of input variables at this time step
            final boolean in_a = wprogram.valueAtTime("a", time);
            final boolean in_b = wprogram.valueAtTime("b", time);
            // values of output variables at this time step
            final String out_x = x(in_a, in_b) ? "1" : "0";
            // store outputs
            output.get("x").append(out_x);
        }
        try {
            final File f = cmd.getOutputFile();
            f.getParentFile().mkdirs();
            final PrintWriter pw = new PrintWriter(new FileWriter(f));
            // write the input
            System.out.println(wprogram.toString());
            pw.println(wprogram.toString());
            // write the output
            System.out.println(f.getAbsolutePath());
            for (final Map.Entry<String,StringBuilder> e : output.entrySet()) {
                System.out.println(e.getKey() + ":" + e.getValue().toString() + ",");
                pw.write(e.getKey() + ":" + e.getValue().toString() + "\n");
            }
            pw.close();
        } catch (final IOException e) {
            Debug.barf(e.getMessage());
        }
    }
    // methods to compute values for output pins
    public static boolean x(final boolean a, final boolean b) { return or(a, b) ; }
}

```

Figure 8.1: Simulator for \mathcal{F} program $x \leq a$ or b ;

Figure 8.2 lists a visitor that determines all of the input variables used by an \mathcal{F} expression AST. This code is listed here for your reference. It is also in your repository. You will need to call this code when writing your simulator generator.

```

public final class DetermineInputVars extends PostOrderExprVisitor {
    private final Set<String> inputVars = new LinkedHashSet<String>();
    private DetermineInputVars(final AssignmentStatement f) { traverseExpr(f.expr); }
    /** Input variables of an AssignmentStatement, in order of occurrence. */
    public static Set<String> inputVars(final AssignmentStatement f) {
        final DetermineInputVars div = new DetermineInputVars(f);
        return Collections.unmodifiableSet(div.inputVars);
    }
    /** Input variables of an FProgram, sorted lexicographically. */
    public static SortedSet<String> inputVars(final FProgram p) {
        final SortedSet<String> vars = new TreeSet<String>();
        for (final AssignmentStatement f : p.formulas) {
            vars.addAll(DetermineInputVars.inputVars(f));
        }
        return vars;
    }
    @Override public Expr visitConstant(final ConstantExpr e) { return e; }
    @Override public Expr visitVar(final VarExpr e) { inputVars.add(e.identifier); return e; }
    @Override public Expr visitNot(final NotExpr e) { return e; }
    @Override public Expr visitAnd(final AndExpr e) { return e; }
    @Override public Expr visitOr(final OrExpr e) { return e; }
    @Override public Expr visitNaryAnd(final NaryAndExpr e) { return e; }
    @Override public Expr visitNaryOr(final NaryOrExpr e) { return e; }
    @Override public Expr visitXOr(final XOrExpr e) { return e; }
    @Override public Expr visitNAnd(final NAndExpr e) { return e; }
    @Override public Expr visitNOr(final NOExpr e) { return e; }
    @Override public Expr visitXNOr(final XNOrExpr e) { return e; }
    @Override public Expr visitEqual(final EqualExpr e) { return e; }
}

```

Figure 8.2: Implementation of DetermineInputVars

8.2 Testing a code generator[†]

In LAB8 here you are writing a code generator: that is, the code that you write produces code as its output, as shown in Figure 8.3.




Figure 8.3: Dataflow of LAB8

How can the course staff evaluate your LAB8 code?

- a. *Reading your hand-written source code.*
 - Hard to mechanize. Doesn't scale.
 - Easy to make mistakes in evaluation.
- b. *Mechanically comparing your hand-written source code to staff code.*
 - *Syntactic comparison:* many legitimate differences.
 - *Semantic comparison:* impossible (in theory) — equivalent to Halting Problem (because Java is Turing Complete).
- c. *Comparing your generated code to staff generated code.*
 - Same problems as above.
- d. *Comparing the \mathcal{W} output produced by your generated code with the \mathcal{W} output from the staff generated code.*
 - \mathcal{W} is regular, so we can compare easily.
 - Engineer's Induction problem: how do we know if we did enough tests? Test Suite Adequacy.

This is one of the main ways that compilers are tested in industry: by running the code that they generate, and examining its output.

8.3 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below. Let 'f' name the \mathcal{F} program we want to simulate, let 'w' name an input \mathcal{W} program for it, and let 'wStaff' name the output of the staff's simulator for this \mathcal{F}/\mathcal{W} combination. The testing equation is:

$$\text{gen(f).simulate(w).isomorphic(wStaff)}$$

	Current	New
TestSimulatorGenerator	70	30

Lab 9

Simulation: $\mathcal{F} \rightarrow \text{x64 Assembly}$

This lab developed by past ECE351 students Nicholas Champion, Akshay Joshi, and Jackson Prange.

In §8 you generated Java code from an \mathcal{F} program. That Java code would read and write \mathcal{W} files, simulating the \mathcal{F} program for each clock tick. Now we will do a similar thing, but instead of Java code we will generate x64 (*i.e.*, x86_64) assembly.

9.1 Reading the \mathcal{W} Program

In §8, your generated Java code called your \mathcal{W} parser to read in the \mathcal{W} input. So your generated Java code could work with different input \mathcal{W} programs.

Calling between Java and native assembler code is possible, but is annoying and time consuming to set up, so we do not want to do it in this lab.

Instead, we will hardcode the \mathcal{W} program into the x64 assembly code. This means that the generated assembly will not read any input when it executes — and also that the generated assembly code will not have to interact with Java code.

9.2 Register Allocation[†]

The programmer may use as many variables as they want, but the machine has only a finite number of registers. Levels of register allocation:

- a. Every variable gets a register.
- b. Variables that are not live at the same time can share a register.

If there are too many variables live at the same time, then some of them need to be *spilled* to main memory.

9.2.1 Register Allocation in This Lab

We need to store the values for the input variables in registers. We will use registers `R8–R15` for the input variables. Conceptual levels of register allocation, in order of increasing sophistication:

- a. Every input variable gets its own register. Ok if the total number of input variables is less than 8.
- b. Each AssignmentStatement (formula) in the input \mathcal{F} program might use only a subset of the input variables. So we could allocate the registers differently for each AssignmentStatement. Ok so long as no individual AssignmentStatement uses more than 8 variables (now the \mathcal{F} program as a whole can use as many variables as it wants).

This is the strategy used in the skeleton code.

9.3 Assembling & Linking

The test harness uses `gcc` for assembling and linking. It assumes that you have `gcc` installed and available in your path.

You can set the `PATH` environment variable in the Eclipse run configuration for the test harness to specify where `gcc` is, if necessary.

`gcc` is the Gnu Compiler Collection, which is the standard compiler on Gnu/Linux systems, and is one of the most commonly used C/C++ compilers in the world.

9.4 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

Let ' f ' name the \mathcal{F} program we want to simulate, let ' w ' name an input \mathcal{W} program for it, and let ' $wStaff$ ' name the output of the staff's simulator for this \mathcal{F}/\mathcal{W} combination. The testing equation is:

`genX64(f).simulate(w).isomorphic(wStaff)`

	Shared Tests	Secret Tests
<code>TestSimulatorGenerator_x86_64</code>	90	10

9.4.1 Bonus Marks

As with all labs, bonus marks are available for improving the lab. A number of such opportunities were discussed in class with respect to this lab, such as:

- improving register allocation
- improving register usage
- improving tests
- switching to `ASMJIT` for assembling instead of `gcc`
- *etc.*

You might have other creative ideas of how to improve this lab or other labs. Bonus marks are assessed by the instructor.

Figure 9.1: x64 assembly for $x \leq a$;

```

.text

.globl out_x
.type out_x, @function
out_x:
    pushq %rbp          // set up a stack frame; rbp is frame pointer
    movq %rsp, %rbp     // set up a stack frame; rbp is frame pointer; rsp is stack pointer
    subq $24, %rsp      // making space on stack
    pushq %r8           // value of input variable "a" is in r8
    popq %rax           // return value goes in rax
    leave
    ret

.globl main
.type main, @function
main:
    pushq %rbp
    movq %rsp, %rbp
    movq $120, %rdi      // put char 'x' in rdi
    call putchar         // print contents of rdi to console
    movq $58, %rdi       // put char ':' in rdi
    call putchar         // print contents of rdi to console
    movq $32, %rdi       // put char ' ' in rdi
    call putchar         // print contents of rdi to console
    movq $0, %r8         // load 0 into r8 (value of "a" at first time step)
    call out_x           // evaluate output pin "x" at first time step
    movq %rax, %rdi      // move x's computed value to rdi
    add $48, %rdi        // convert int to ascii (48 is ASCII value '0')
    call putchar         // print contents of rdi to console
    movq $59, %rdi       // put char ';' in rdi
    call putchar         // print contents of rdi to console
    movq $10, %rdi       // put newline char in rdi
    call putchar         // print contents of rdi to console

    popq %rax
    leave
    ret

```

Lab 10

VHDL Recognizer & Parser

In this lab we will build a recognizer and a parser using Parboiled for a very small and simple subset of VHDL. The grammar for this subset of VHDL is listed in Figure 10.1. Restrictions on our subset of VHDL include:

- only bit (boolean) variables
- no stdlogic
- no nested ifs
- no aggregate assignment
- no combinational loops
- no arrays
- no generate loops
- no case
- no when (switch)
- inside process: either multiple assignment statements or multiple if statements; inside an if there can be multiple assignment statements
- no wait
- no timing
- no postponed
- no malformed syntax (or, no good error messages)
- identifiers are case sensitive
- one architecture per entity, and that single architecture must occur immediately after the entity

Files:

ece351.v.VRecognizer.java
ece351.v.VParser.java

Tests:

ece351.v.test.TestVRecognizerAccept
ece351.v.test.TestVRecognizerReject
ece351.v.test.TestVParser

In w2013 the median time to complete this lab was about 9 hours. Therefore, we have made three changes to this lab for s2013 to hopefully bring this time down to the five hour budget.

First, we added append methods to the VHDL AST classes (as we also did for \mathcal{F} this term). This means that your parser does not need to place ImmutableList objects on the stack directly, which many students found difficult to manage (and certainly the syntax for the casting is complex). This also means that you can construct your AST in a top-down manner, which many students seem to find easier. Both \mathcal{W} and \mathcal{F} were changed to this style this term as well, so this is what you are already familiar with.

Second, we provided implementations of the object contract methods on the VHDL AST classes for you, instead of requiring you to fill them in.

Third, we removed the desugarer and define before use parts of this lab. So now it's just the parser.

<i>Program</i>	→	(<i>DesignUnit</i>)*
<i>DesignUnit</i>	→	<i>EntityDecl</i> <i>ArchBody</i>
<i>EntityDecl</i>	→	'entity' <i>Id</i> 'is' 'port' '(' <i>IdList</i> ':' 'in' 'bit' ';' <i>IdList</i> ':' 'out' 'bit' ')' ';' 'end' ('entity' <i>Id</i>) ';
<i>IdList</i>	→	<i>Id</i> (',' <i>Id</i>)*
<i>ArchBody</i>	→	'architecture' <i>Id</i> 'of' <i>Id</i> 'is' ('signal' <i>IdList</i> ':' 'bit' ';')? 'begin' (<i>CompInst</i>)* (<i>ProcessStmts</i> <i>SigAssnStmts</i>) 'end' <i>Id</i> ;
<i>SigAssnStmts</i>	→	(<i>SigAssnStmt</i>) ⁺
<i>SigAssnStmt</i>	→	<i>Id</i> '<=' <i>Expr</i> ;
<i>ProcessStmts</i>	→	(<i>ProcessStmt</i>) ⁺
<i>ProcessStmt</i>	→	'process' '(' <i>IdList</i> ')' 'begin' (<i>IfElseStmts</i> <i>SigAssnStmts</i>) 'end' 'process' ;
<i>IfElseStmts</i>	→	(<i>IfElseStmt</i>) ⁺
<i>IfElseStmt</i>	→	'if' <i>Expr</i> 'then' <i>SigAssnStmts</i> 'else' <i>SigAssnStmts</i> 'end' 'if' (<i>Id</i>)? ;
<i>CompInst</i>	→	<i>Id</i> ':' 'entity' 'work.' <i>Id</i> 'port' 'map' '(' <i>IdList</i> ')' ;
<i>Expr</i>	→	<i>XnorTerm</i> ('xnor' <i>XnorTerm</i>)*
<i>XnorTerm</i>	→	<i>XorTerm</i> ('xor' <i>XorTerm</i>)*
<i>XorTerm</i>	→	<i>NorTerm</i> ('nor' <i>NorTerm</i>)*
<i>NorTerm</i>	→	<i>NandTerm</i> ('nand' <i>NandTerm</i>)*
<i>NandTerm</i>	→	<i>OrTerm</i> ('or' <i>OrTerm</i>)*
<i>OrTerm</i>	→	<i>AndTerm</i> ('and' <i>AndTerm</i>)*
<i>AndTerm</i>	→	<i>EqTerm</i> ('=' <i>EqTerm</i>)*
<i>EqTerm</i>	→	'not' <i>EqTerm</i> '(' <i>Expr</i> ')' <i>Var</i> <i>Constant</i>
<i>Constant</i>	→	'0' '1'
<i>Var</i>	→	<i>id</i>
<i>Id</i>	→	<i>Char</i> (<i>Char</i> <i>Digit</i> '_')*
<i>Char</i>	→	[A-Za-z]
<i>Digit</i>	→	[0-9]

Figure 10.1: Grammar for VHDL.

This grammar has gone through a number of evolutions. First, David Gao and Rui Kong (3B) downloaded a complete VHDL grammar implementation for the ANTLR tool. They found that grammar was too complicated to use, so they started writing their own with ANTLR. Then Alex Duanmu and Luke Li (1B) reimplemented their grammar in TXL and simplified it. Aman Muthrej (3A) reimplemented it in Parboiled and simplified it further. Wallace Wu (TA) refactored Aman's code. Michael Thiessen (3A) helped make the expression sub-grammar more consistent with the grammar of \mathcal{F} .

There will no doubt be improvements and simplifications that you think of, and we would be happy to incorporate them into the lab.

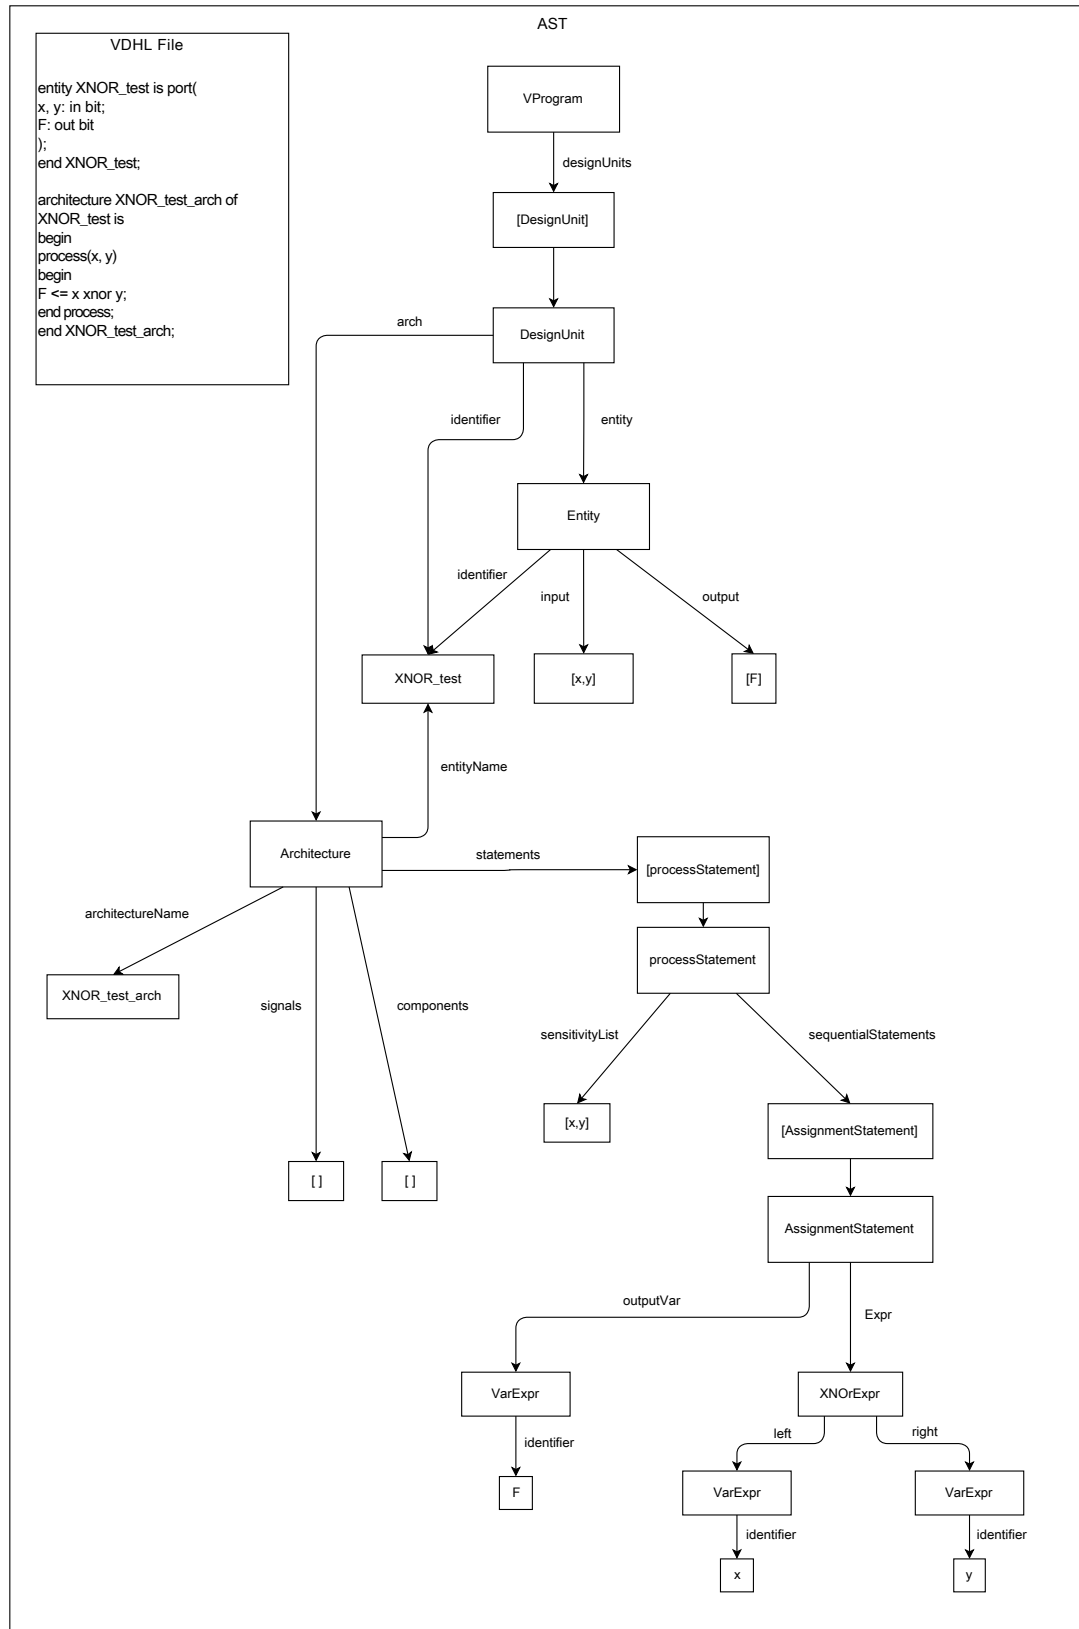


Figure 10.2: Object diagram for example VHDL AST

10.1 Keywords and Whitespaces

To make the code for the recognizer and parser a little bit more legible (and easier to implement), classes `VRecognizer` and `VParser` inherit a common base class called `VBase`,¹ which contains a set of rules that will match keywords found in the language in a case-independent manner (since VHDL is case-insensitive² – e.g., both ‘signal’ and ‘SIGNAL’ represent the same keyword in the language).

Figure 10.3 is an example of one of the rules found in `VBase`. Here, the rule matches the keyword ‘signal’ and as there is no whitespace handling, your code should ensure that at least one whitespace exists between this keyword and the next token/string that is matched.

```
Rule NOR() {
```

¹ `VBase` itself extends `BaseParser351`, which provides extra utility and debugging methods over Parboiled’s `BaseParser` class.

² In some VHDL compilers, there may be some exceptions, where the case sensitivity matters a handful of grammar productions

Figure 10.3: Implementation of rule used to match the keyword ‘signal’

10.2 VHDL Recognizer

Write a recognizer using Parboiled for the VHDL grammar defined in Figure 10.1. Follow the methodology that we have been developing this term:

- Make a method for every non-terminal [§1.3]
- Convert EBNF multiplicities in the grammar into Parboiled library calls [§5.1]
- Convert the lexical specification into Parboiled library calls [§5.1]
- Remember to explicitly check for EOI

You should be able to copy and paste the code from your \mathcal{F} parboiled recognizer with minimal modifications, since \mathcal{F} is a subset of our VHDL grammar.

10.3 VHDL Parser

Write a parser using Parboiled for the VHDL grammar defined in Figure 10.1. You do not need to write a pretty-printer nor the object contract methods: these are provided for you in the new VHDL AST classes, and you already wrote them in the shared AST classes (`ece351.common.ast`). You do not need to edit the AST classes. Just write the parser. Follow the steps we learned earlier in the term:

- Draw out the evolution of the stack for a simple input [Figure 5.3].
- Print out our VHDL grammar.
- Annotate the printed copy of the grammar with where stack manipulations occur according to your diagram.

Sources:

`ece351.v.VRecognizer`

Tests:

`ece351.v.test.TestVRecognizerAccept`

`ece351.v.test.TestVRecognizerReject`

Sources:

`ece351.v.VParser`

Libraries:

`ece351.v.ast.*`

`ece351.common.ast.*`

Tests:

`ece351.v.test.TestVParser`

- Calculate the maximum and expected size of the stack for different points in the grammar.
- Copy your recognizer code into the parser class.
- Add some (but not all) actions (push, pop, *etc.*).
- Add some `checkType()` and `debugmsg()` calls [Figure 5.5].
- Run the `TestVParser` harness to see that it passes.
- Add more actions, `checkType()` calls, *etc.*.
- Repeat until complete.

When you implement the VHDL parser, there may be a few grammar productions shown in Figure 10.1 where you will need to rewrite each of these rules in a different way so that you can instantiate all of the required objects used to construct the AST of the program being parsed.

10.4 Engineering a Grammar[†]

For simple languages like \mathcal{W} and \mathcal{F} we can write a ‘perfect’ recognizer: that is, a recognizer that accepts all grammatical strings and rejects all ungrammatical strings. For more complicated languages, such as VHDL, it is common to design a grammar for a particular task, and that task may be something other than being an ideal recognizer. For example, the grammar might accept some strings that should be rejected.

Industrial programming environments such as Eclipse often include multiple parsers for the same language. Eclipse’s Java development tools include three parsers: one that gives feedback on code while the programmer is editing it (and so the code is probably ungrammatical); one that gives feedback on the code when the programmer saves the file (when the code is expected to be grammatical); and one that is part of the compiler. As you may surmise from this example, the task of giving good feedback to a programmer about ungrammatical code can be quite different from the task of translating grammatical code.

The `TestVRecognizerReject` harness tests for some obvious cases that should be rejected. Do not invest substantial effort into trying to reject more cases.

10.5 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

	Current	New
TestVParboiledRecognizerAccept	20	5
TestVParboiledRecognizerReject	5	0
TestVParboiledParser	45	25

TestVParboiledRecognizer just runs the recognizer to see if it crashes. TestVParboiledParser checks the following equation:

$$\forall v : *.vhd \mid \text{parse}(v).\text{isomorphic}(\text{parse}(\text{prettyprint}(\text{parse}(v))))$$

You do not get any marks for any other parser tests until TestObjectContractF passes everything. You do not get any marks for rejection tests until some acceptance tests pass.

Lab 11

VHDL \rightarrow VHDL: *Desugaring & Elaboration*

In this lab, you will be writing two VHDL to VHDL transformers using the Visitor design pattern. The first, and simpler one, will simply rewrite expressions to replace operators like XOR with their equivalent formulation in terms of AND, OR, and NOT. We call this *desugaring*, where operators like XOR are considered to be *syntactic sugar*: things that might be convenient for the user, but do not enlarge the set of computations that can be described by the language.

The second, and more sophisticated, transformer will expand and inline any component instance declared within an architecture of a design unit if the component is defined within the same VHDL source file. The procedure that this transformer performs is known as *elaboration*. Elaboration is essential for us to process the four bit ripple-carry adder, for example, because it comprises four one-bit full adders.

11.1 VHDL \rightarrow VHDL: *Desugaring*

Often times, languages provide redundant syntactic constructs for the purpose of making code easier to write, read, and maintain for the programmer. These redundant constructs provide ‘syntactic sugar’. In this part of the lab, you will transform VHDL programs that you parse into valid VHDL source code that ‘desugars’ the expressions in the signal assignment statements. In other words, the desugaring process reduces the source code into a smaller set of (i.e., more basic) constructs.

In VHDL, desugaring may be useful in cases where the types of available logic gates are limited by the programmable hardware you might be using to run your VHDL code. For example, if the programmable hardware only comprises of NAND gates, a VHDL synthesizer will be required to rewrite all of the logical expressions in your code using NAND operators.

For this part of the lab, write a ‘desugarer’ that traverses through

Sources:

ece351.v.DeSugarer

Tests:

ece351.v.tests.TestDeSugarer

an AST corresponding to an input VHDL program and rewrites all expressions so that all expressions in the *transformed* AST only consist of AND, OR, and NOT logical operators. This means that expressions containing XOR, NAND, NOR, XNOR, and = operators must be rewritten. For example, the expression $x \oplus y$ (XOR) is equivalent to:

$$x \oplus y \equiv x \cdot !y + !x \cdot y \quad (11.1)$$

Table 11.1 is the truth table for the '=' operator: By observation, we can see that this truth table is equivalent to that of XNOR.

x	y	$x = y$
0	0	1
0	1	0
1	0	0
1	1	1

Table 11.1: Truth table for the '=' operator. Equivalent to XNOR.

11.2 Elaboration

The following sections describe the expected behaviour of the elaborator.

Sources:

ece351.v.Elaborator

Tests:

ece351.v.test.TestElaborator

11.2.1 Inlining Components without Signal List in Architecture

Consider the VHDL source shown in Figure 11.1. Here, we have two design units, OR_gate_2 and four_port_OR_gate_2, where the architecture body corresponding to four_port_OR_gate_2 instantiates two instances of OR_gate_2, namely OR1 and OR2 (lines 20 and 21).

When the elaborator processes this program, it will check the design units sequentially. In this example, OR_gate_2 is checked first. The architecture body corresponding to OR_gate_2 does not instantiate any components, so the elaborator does not do anything to this design unit and moves onto the next design unit. In the architecture body, four_port_structure, we see that there are two components that are instantiated. Since there are components within this architecture body, the elaborator should then proceed to inline the behaviour of the components into four_port_structure and make the appropriate parameter/signal substitutions.

Consider the component OR1. The elaborator will search through the list of design units that make up the program and determine the design unit that the component is instantiating. In this example, OR1 is an instance of the design unit OR_gate_2 (see the string immediately following "work." on line 20). Then the elaborator proceeds to determine how the signals used in the port map relate to the ports defined in the entity declaration of OR_gate_2. OR1 maps the signals

```

1  entity OR_gate_2 is port (
2      x , y: in bit;
3      F: out bit
4  );
5  end OR_gate_2;
6
7  architecture OR_gate_arch of OR_gate_2 is begin
8      F <= x or y;
9  end OR_gate_arch;
10
11 entity four_port_OR_gate_2 is port (
12     a,b,c,d : in bit;
13     result : out bit
14 );
15 end four_port_OR_gate_2;
16
17 architecture four_port_structure of four_port_OR_gate_2 is
18     signal e, f : bit;
19 begin
20     OR1: entity work.OR_gate_2 port map(a,b,e);
21     OR2: entity work.OR_gate_2 port map(c,d,f);
22     result <= e or f;
23 end four_port_structure;

```

Figure 11.1: VHDL program used to illustrate elaboration.

a, b, and e, to the ports x, y, and F of the entity `OR_gate_2`, respectively. Using the private member, `current_map`, in the Elaborator class will help you with the signal substitutions that occurs when a component is inlined. After the mapping is established, the elaborator then proceeds to replace `OR1` by inlining the architecture body corresponding to the entity `OR_gate_2` into `four_port_structure`.

The same procedure is carried out for the component `OR2` and we now will have an equivalent architecture body as shown in Figure 11.2. Lines 5 and 6 in Figure 11.2 corresponds to the inlining of `OR1` and `OR2` (found in lines 20 and 21 from Figure 11.1).

```

1  begin
2      result <= ( e or f );
3      e <= ( a or b );
4      f <= ( c or d );
5  end four_port_structure;
6
7  entity eight_port_OR_gate_2 is port(
8      x0, x1, x2, x3, x4, x5, x6, x7 : in bit;
```

Figure 11.2: Elaborated architecture body, `four_port_structure`.

11.2.2 Inlining Components with Signal List in Architecture

In addition to substituting the input and output pins for a port map, you will also encounter situations where there are signals declared in the architecture body that you are trying to inline to the design unit the elaborator is currently processing.

For example, consider VHDL source in Figure 11.3, which is the extension of the program in Figure 11.1.

The two components in `eight_port_structure` are instances of `four_port_OR_gate_2`; the architecture of `four_port_OR_gate_2` defines signals `e` and `f`. Now, if we elaborate, say, `OR1`, we determine the mapping as before for the input and output pins, but we also need to consider the signals defined within the architecture. If we simply add `e` and `f` to the list of signals of `eight_port_structure`, we will run into problems of multiply defined signals when we elaborate `OR2`; we will obtain a signal list with two `e`'s and two `f`'s. Furthermore, we will change the logical behaviour defined in `eight_port_structure`.

To address this issue, all internal signals that are added as a result of elaboration will be prefixed with '`comp<num>_`', where `<num>` is a unique identifier¹ used to ensure that the elaboration does not change the logical behaviour of the program. The result of elaborating `eight_port_OR_gate_2` is shown Figure 11.4.

¹ This number starts at 1 and increments for each component that is instantiated in the VHDL program. `<num>` is never reset.

```

1  entity eight_port_OR_gate_2 is port (
2      x0, x1, x2, x3, x4, x5, x6, x7 : in bit;
3      y : out bit
4  );
5  end eight_port_OR_gate_2;
6
7  architecture eight_port_structure of eight_port_OR_gate_2 is
8      signal result1, result2 : bit;
9  begin
10     OR1: entity work.four_port_OR_gate_2 port map(x0, x1, x2, x3, result1);
11     OR2: entity work.four_port_OR_gate_2 port map(x4, x5, x6, x7, result2);
12     y <= result1 or result2;
13 end eight_port_structure;

```

Figure 11.3: Extension of the VHDL program shown in Figure 11.1.

```

1  entity eight_port_OR_gate_2 is port(
2      x0, x1, x2, x3, x4, x5, x6, x7 : in bit;
3      y : out bit
4  );
5  end eight_port_OR_gate_2;
6  architecture eight_port_structure of eight_port_OR_gate_2 is
7      signal result1, result2, comp3_e, comp3_f, comp4_e, comp4_f : bit;
8  begin
9      y <= ( result1 or result2 );
10     result1 <= ( comp3_e or comp3_f );
11     comp3_e <= ( x0 or x1 );
12     comp3_f <= ( x2 or x3 );
13     result2 <= ( comp4_e or comp4_f );
14     comp4_e <= ( x4 or x5 );
15     comp4_f <= ( x6 or x7 );
16 end eight_port_structure;

```

Figure 11.4: Elaborated architecture body, eight_port_structure.

11.2.3 Inlining Components with Processes in Architecture

The previous examples demonstrated the behaviour of inlining components with architectures that only contain signal assignment statements. When the elaborator encounters processes in the inlining, a similar procedure is performed. The main difference in the procedure for processes is to make the appropriate signal substitutions in sensitivity lists and if-else statement conditions.

11.2.4 Notes

- The elaborator will only expand components when its corresponding design unit is also defined in the same file.
- The elaborator processes the design units in sequential order. We assume that the VHDL programs we are transforming are written so that you do not encounter cases where the architecture that you are inlining contains components that have not yet been elaborated.
- We will assume that the VHDL programs being elaborated will not result in architecture bodies with a mixture of parallel signal assignment statements and process statements (so that the parser from Lab 10 can parse the transformed programs).

11.3 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below. The testing equation for desugaring is:

`parse(v).desugar().isomorphic(parse(vStaff))`

Similarly, the testing equation for elaboration is:

`parse(v).elaborate().isomorphic(parse(vStaff))`

	Shared	Secret
TestDeSugarer	20	10
TestElaborator	50	20

Choice: either do this lab or do the lab marked as LAB12. Do not do both.

Lab 12

VHDL Process Splitting & Combinational Synthesis

For this lab, you will be writing code to perform two other transformations on VHDL programs. The first transformation is a VHDL to VHDL transformation, which we call *process splitting*. Process splitting involves breaking up if/else statements where multiple signals are being assigned new values in the if and else clauses.

In the second part of this lab, you will be translating VHDL to \mathcal{F} , which we call *combinational synthesis*. The combinational synthesizer will take the VHDL program output from the process splitter and generate a valid \mathcal{F} program from it.

Files:

ece351.v.Splitter.java
ece351.v.Synthesizer.java

Tests:

ece351.v.test.TestSplitter
ece351.v.test.TestSynthesizer

12.1 Process Splitter

The process splitter will perform the necessary transformations to VHDL program ASTs so that exactly one signal is being assigned a value in a process. Consider the VHDL code shown in Figure 12.1.

Here, we have a process in the architecture `behv1` of the entity `Mux` that requires splitting because in the if and else clauses, there are two signals that are being assigned new values: `00` and `01`. The splitter should observe this and proceed to replace the current process with two new processes: one to handle the output signal `00` and the other to handle `01`. Figure 12.2 shows the splitter's output when the code in Figure 12.1 is processed. Note that the sensitivity lists for the two new processes only contain the set of signals that may cause a change to the output signals.

You might also notice that it appears that the condition becomes longer and more complicated from Figure 12.1 to Figure 12.2. That is not related to process splitting: it's the result of desugaring the `'=`' in the condition. Recall that `'=`' does not exist in \mathcal{F} , so it needs to be expressed in terms of the operators that do exist in \mathcal{F} : `and`, `or`, `not`. Table 11.1 shows that `'=`' is equivalent to `xnor`.

This desugaring highlights the importance of syntactic sugar in terms of usability for the programmer, and the importance of desugaring for the compiler engineer. Imagine if the programmer had to write the condition in Figure 12.2: unusable! But conversely, imagine if the compiler engineer had to enhance the entire \mathcal{F} toolchain to deal with so many more operators: too much work; too many opportunities for error.

```
1  entity Mux is port(  
2      I3,I2,I1,I0,S: in bit;  
3      O0,O1: out bit  
4  );  
5  end Mux;  
6  
7  architecture behv1 of Mux is  
8  begin  
9      process(I3,I2,I1,I0,S)  
10         begin  
11             if (S = '0') then  
12                 O0 <= I0;  
13                 O1 <= I2;  
14             else  
15                 O0 <= I1;  
16                 O1 <= I3;  
17             end if;  
18         end process;  
19  
20 end behv1;
```

Figure 12.1: Example VHDL program used to illustrate process splitting.


```

1  entity Mux is port(
2      I3, I2, I1, I0, S : in bit;
3      O0, O1 : out bit
4  );
5  end Mux;
6  architecture behv1 of Mux is
7
8  begin
9      process ( S, I0, I1 )
10         begin
11             if ( ( not ( ( ( S and ( not ( '0' ) ) ) or ( ( not ( S ) ) and '0' ) ) ) ) ) then
12                 O0 <= I0;
13             else
14                 O0 <= I1;
15             end if;
16         end process;
17         process ( S, I2, I3 )
18             begin
19                 if ( ( not ( ( ( S and ( not ( '0' ) ) ) or ( ( not ( S ) ) and '0' ) ) ) ) ) then
20                     O1 <= I2;
21                 else
22                     O1 <= I3;
23                 end if;
24             end process;
25 end behv1;

```

Figure 12.2: The resulting VHDL program of Figure 12.1 after process splitting.

12.2 Splitter Notes

- Assume that there is exactly one assignment statement in the if body and one assignment statement in the else body that write to the same output signal. This implies that you do not need to handle the case where latches are inferred. Making this assumption should reduce the amount of code you need to write for this lab.
- The private variable `usedVarsInExpr` is used to store the variables/signals that are used in a VHDL expression. This is helpful when you are trying to create sensitivity lists for new processes.

12.3 Synthesizer

After parsing the input VHDL source file and performing all of the transformations (*i.e.*, desugaring, elaborating, and process splitting), the synthesizer will traverse the (transformed) AST (of the VHDL program), extract all expressions in the program, and generate an \mathcal{F} program containing these expressions.

In §8 we generated a Java program by constructing a string. We noted that there was no guarantee that the string we generated would be a legal Java program. One way to ensure that we generate a legal program is to build up an AST instead of building up a string. Of course, to do that we need to have AST classes, which we do not have for Java — but do have for \mathcal{F} .

Because the \mathcal{F} grammar and AST classes are a subset of those for VHDL, we can simply reuse many of the expressions from the input VHDL program with only minor modification to rename the variables to avoid name collision.

12.3.1 If/Else Statements

Whenever an if/else statement is encountered, first extract the condition and add it to the \mathcal{F} program. Because there is no assignment that occurs in the condition expression, generate an output variable for the condition when you output the condition to the \mathcal{F} program. The output variable will be of the form: ‘condition<num>’, where <num> = 1, 2, 3, ... <num> is incremented each time you encounter a condition in the VHDL program and is never reset.

After appending the condition of the if/else statement to the \mathcal{F} program, construct an assignment for the output variable like so:¹

```

if ( vexpr ) then
    output <= vexpr1;
else
    output <= vexpr2;
end if;

```

¹ Note that process splitting is useful here because we will only have one signal assignment statement in the if- and else- bodies, where both statements assign the expressions to the same output signal.

The formulas that should be appended to the \mathcal{F} program is:

```

condition<num> <= vexpr;
output <= ( condition<num> and (vexpr1) or (not condition<num>) and (vexpr2) );

```

Let’s refer to the variable named conditionX as an *intermediate variable* because it is not intended as part of the circuit’s final output, but is used in computing the final output pins. Our \mathcal{F} simulator generator and technology mapper do not support these intermediate variables. Removing them is not hard: we just inline their definition wherever they are used. We will provide you with code to do that later.

The synthesizer here is not the only place where these intermediate variables are created. They also occur in elaboration. They are helpful for debugging, so it’s worth keeping them here and removing them later.

12.3.2 Example Synthesizing Signal Assignment Statements

Consider the VHDL program shown in Figure 12.3. When the synthesizer processes this program², an \mathcal{F} program consisting of a single formula (corresponding to) is generated (see Figure 12.4).

² after desugaring, elaborating, and splitting

```

1  entity NOR_gate is port(
2      x,y: in bit;
3      F: out bit
4  );
5  end NOR_gate;
6
7  architecture NOR_gate_arch of NOR_gate is
8  begin
9
10     F <= x nor y;
11
12 end NOR_gate_arch;
```

Figure 12.3: Example used to illustrate synthesizing assignment statements.

```

1  NOR_gateF <= ( not ( NOR_gatex or NOR_gatey ) );
```

Figure 12.4: Synthesized output of the program in Figure 12.3.

12.3.3 Example Synthesizing If/Else Statements

Consider the VHDL program shown in Figure 12.5. After performing all of the VHDL transformations that you have written, the synthesizer will generate the \mathcal{F} program shown in Figure 12.6.

In Figure 12.6, observe that the synthesizer translates the if/else statement in Figure 12.5 into two \mathcal{F} formulae. The first formula that is generated corresponds to the condition that is checked in the statement (i.e., $x = '0'$ and $y = '0'$). The second formula combines the expressions found in the if and else bodies so that the expression in the if-body is assigned to F if the condition is true; otherwise, the expression in the else-body is assigned to F.

```

1  entity NOR_gate is port(
2      x, y: in bit;
3      F: out bit
4  );
5  end NOR_gate;
6
7  architecture NOR_gate_arch of NOR_gate is
8  begin
9      process(x, y)
10     begin
11         if (x='0' and y='0') then
12             F <= '1';
13         else
14             F <= '0';
15         end if;
16     end process;
17 end NOR_gate_arch;

```

Figure 12.5: Example used to illustrate synthesizing assignment statements.

```

1  condition1 <= ( ( not ( ( NOR_gatex and ( not '0' ) ) or ( ( not NOR_gatex )
      and '0' ) ) ) and ( not ( ( NOR_gatey and ( not '0' ) ) or ( ( not
      NOR_gatey ) and '0' ) ) ) );
2  NOR_gateF <= ( condition1 and ( '1' ) ) or ( ( not condition1 ) and ( '0' ) );

```

Figure 12.6: Synthesized output of the program in Figure 12.5.

12.4 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below.

	Current	New
TestSplitter	40	10
TestSynthesizer	40	10

Lab 13

Simulation: $\mathcal{F} \rightarrow \text{Assembly}$

In §8 you generated Java code from an \mathcal{F} program. That Java code would read and write \mathcal{W} files, simulating the \mathcal{F} program for each clock tick. In this lab you will make a variation of the simulator generator from §8. The main loop is still the same, with the call to the \mathcal{W} parser *etc.*. The only part that will be different is the methods you generate to compute the values of the output pins. For example, in §8 you might have generated a method like the one shown in Figure 13.1. The point of this lab is to generate these methods in assembly, and have the main body (still in Java) call the assembly versions.

```
// methods to compute values for output pins
public static boolean x(final boolean a, final boolean b) { return or(a, b) ; }
```

13.1 Which assembler to use?

Options include:

- `MASM32`. Requires a 32 bit `JDK`. This is the option supported by the description below and the skeleton code. The other options have no instructional support.
- `GCC` with inline assembly.
- `JWasm` <http://www.japheth.de/JWasm.html>. `JWasm` is an open-source fork of the Watcom assembler (`Wasm`). It is written in C and is `MASM` compatible. The 'J' here apparently stands for the maintainer (Japheth), and not for Java.

Watcom, as you know, is from the University of Waterloo. UW has a long and famous history with compilers, that you can read a little bit about at http://www.openwatcom.com/index.php/Project_History

This lab by Nicholas Champion.

Choice: You can do either this lab or the original `LAB11`. Don't do both.

There is no official instructional support for this lab. Do not attempt this lab if you are not comfortable learning assembly on your own and configuring your system to work with the assembler that you have chosen.

Figure 13.1: Generate this method in assembly rather than in Java

- *MIPS Assembler and Runtime Simulator (MARS)*. A Java implementation of the RISC machine described in the Hennessey & Patterson computer architecture textbook. <http://courses.missouristate.edu/kenvollmar/mars/> Because this is a simulator written in Java, it might avoid a lot of the systems issues of going to x86 assembly. Also, it might be easier to generate MIPS assembly than x86 (RISC instruction sets are usually simpler than CISC ones).

There might be some bonus marks available if you develop a skeleton and some instructional support for an assembler such as MARS that has no system dependencies outside of a regular JVM.

13.2 MASM32 Configuration

- MASM32 requires Windows (any version should work). MASM32 includes an old version of Microsoft's assembler ml.exe. Newer version are included with releases of Visual Studio releases and support newer instruction.
- Install the 32-bit JDK and configure your Eclipse installation to use this JDK. JNI and Windows operating systems require that 64-bit programs use 64-bit dynamic libraries (*i.e.*, what MASM32 will generate). 64-bit programs cannot use 32-bit libraries.
- A possible multi-platform alternative to experiment with, which is NOT-supported by the skeleton or test harness, is JWASM.
- It can be downloaded from <http://www.masm32.com/>
- It is recommended that you install MASM32 to C:\masm32
- Add (<Install Path>\bin) to your system path environment variable (*i.e.*, C:\masm32\bin)
- To check if it was added to the path, open a new command window and type "ml.exe". It should output "Microsoft (R) Macro Assembler Version 6.14.8444..."

13.3 MASM32 JNI and x86 Assembly

The skeleton and test harness created for solving this lab using MASM32 are based on the Lab 8 and so approximately a third of the redacted skeleton code can be taken from a lab 8 solution. The main difference from lab 8 is the means in which the F-statements are evaluated. In lab 8, methods were created and defined in Java for each statement by walking the tree using pre-order. In lab 12, these methods are to be implemented in x86 assembly. It is not possible to include inline assembly in Java as Java is executed in the JVM. Therefore the x86 is assembled and linked into a shared library (dynamic link library (DLL) on Windows). This library must be loaded by the generated java code using the Java Native Interface (JNI). JNI allows native functions to be added to the class which loads the library. The library load in the skeleton is static while the methods are not static. Hence the main method will instantiate its own class and make calls

You can use the JavaH tool in order to see the function signatures. It is included with the JDK to generate the C header for a java class. You can modify TestSimulatorGeneratorMasm32.java to generate the headers files by uncommenting the lines which write the batch file. You will need to change the path to point to your JDK. To see how java call types map to standard C types (and hence ASM), see jni.h (in the JDK).

to its non-static native methods. JNI also imposes certain naming conventions on the signatures of functions declared in the shared library. For example the native method

```
public native boolean Simulator_opt1_or_true1 (final boolean a);
```

translates to the C shared library export signature

```
JNIEXPORT jboolean JNICALL Java_Simulator_1opt1_1or_1true1_Evalx(JNIEnv *, jobject, jboolean);
```

This C export translates to the x86 assembly signature (for the purposes of the this lab)

```
Java_Simulator_1opt1_1or_1true1_Evalx proc JNIEnv:DWORD, jobject:DWORD, in_a:BYTE
```

It is notable that JNI requires including pointers to the Java environment and the Java class instance. It is also notable that the underscores in the class name and the method name must be replaced with `_1`. Most of the function prototype details are already included in the skeleton.

Figure 13.2 is a simple example of the generate code for a shared library for a \mathcal{F} Program.

```
.386
.model flat,stdcall
option casemap:none

Java_Simulator_1opt1_1or_1true1_Evalx PROTO :DWORD, :DWORD, :BYTE

.code
LibMain proc hInstance:DWORD, reason:DWORD, reserved1:DWORD
    mov eax, 1
    ret
LibMain endp

Java_Simulator_1opt1_1or_1true1_Evalx proc JNIEnv:DWORD, jobject:DWORD, in_a:BYTE
    ; (a, true)or

    mov EAX, 0
    mov EBX, 0
    mov AL, in_a
    mov EBX, 1
    or EAX, EBX
    ret
Java_Simulator_1opt1_1or_1true1_Evalx endp
End LibMain
```

Figure 13.2: Example assembly for an \mathcal{F} Program

Figure 13.3 the linker for MASM32 uses a definitions file to determine which symbols should be exported. Note that the library name does not require the replacing underscores. Additional exports are added on separate lines.

```
LIBRARY Simulator_opt1_or_true1
EXPORTS Java_Simulator_1opt1_1or_1true1_Evalx
```

This lab requires generating the instructions for the F-statement. Most of the other code is generated by the skeleton. To generate the instructions, the AST should be walked in post order as a simple topological sort and stored in the operations list. It is then necessary to assign signals and intermediate signals to registers and write instructions for each AST object. For example VarExpr or ConstantExpr can be converted into 8-bit immediate-mode mov instructions. The other operators can be implemented using various Boolean operator instruction. It is possible to implement all expressions including a NotExpr using a single instruction and no conditionals. The result of the last computation for the statement should be stored in EAX as the return value for the function. x86 has four 32-bit general purpose registers EAX, EBX, ECX, and EDX and supports 16-bit and 8-bit access modes. A helper class is provided which contains the names of that correspond to the other addressing modes (*i.e.*, AL references the low byte of EAX). Since there are only four registers (or fewer if you remove some from the Register.General list for debugging purposes), it is possible that large assignment statements will require the use of memory. A simple FIFO (registerQueue) strategy is used to determine which register to save to memory and reassign before performing an operation. Memory can be allocated using the LOCAL keyword: LOCAL m0 :BYTE. These allocation statements should be immediately following the function signature before any instructions. In order to track which memory is in use the hash map memoryMap maps the name of the memory allocation to a Boolean value of whether it is used. For convenience the IdentityHashMap storageMap maps an Expr to the name of either the memory or register currently storing its output value.

Figure 13.3: Example definitions file for a dynamic link library (DLL)

Edit SimulatorGeneratorMasm32.java
and run TestSimulatorGenerator-
Masm32.java

13.4 *Evaluation*

There is no automated evaluation for this lab due to the variety of possible system dependencies. You will have to meet with the course staff to be evaluated, which will include you answering questions about how your design works. Your generated assembly code should pass all of the tests from §8.

There are potentially bonus marks if you find a cross-platform assembler with minimal system dependencies and develop a bit of a skeleton and instructional support for it. The `MARS MIPS` simulator mentioned above is the best possibility for this that we are aware of.

Lab 14

Simulation: $\mathcal{F} \rightarrow \text{JVM}$

14.1 How *javac* compiles boolean logical operators[†]

The Java operand stack treats booleans as integers because each slot on the operand stack is 32-bits. Zero is false and one is true (as in C). Boolean logic operators are translated into if comparisons against zero. There are no JVM opcodes for boolean logic.

Also note the short-circuiting semantics of boolean logic operators in Java (as in C): operands are only evaluated if necessary. For example, if the first operand of a conjunction (and) is false, then the second operand is not evaluated. Similarly, if the first operand of a disjunction (or) is true, the second operand is not evaluated.

```
public class javac_code {  
    boolean or (boolean x, boolean y) { return x || y; }  
    boolean and(boolean x, boolean y) { return x && y; }  
    boolean not(boolean x)           { return !x; }  
}
```

Compiled from "javac_code.java"

```

public class javac_code {
  public javac_code();
    Code:
      0: aload_0
      1: invokespecial #1          // Method java/lang/Object."<init>":() V
      4: return

  boolean or(boolean, boolean);      // x || y
    Code:
      0: iload_1          // push x
      1: ifne            8   // if (x != 0) goto instruction 8
      4: iload_2          // push y
      5: ifeq           12   // if (y == 0) goto instruction 12
      8: iconst_1         // push true (one)
      9: goto           13   // goto return true
     12: iconst_0         // push false (zero)
     13: ireturn         // return top of stack

  boolean and(boolean, boolean);      // x && y
    Code:
      0: iload_1          // push x
      1: ifeq           12   // if (x == 0) goto instruction 12
      4: iload_2          // push y
      5: ifeq           12   // if (y == 0) goto instruction 12
      8: iconst_1         // push true (one)
      9: goto           13   // goto return true
     12: iconst_0         // push false (zero)
     13: ireturn         // return top of stack

  boolean not(boolean);              // !x
    Code:
      0: iload_1          // push x
      1: ifne            8   // if (x != 0) goto instruction 8
      4: iconst_1         // push true (one)
      5: goto           9    // goto return true
      8: iconst_0         // push false (zero)
      9: ireturn         // return top of stack
}

```

14.2 How we will generate code

We do not need to respect the short-circuiting semantics of boolean logic operators in Java/C — because our input language is not Java/C. Our input language is \mathcal{F} , which does not have short-circuiting semantics. So we can compute disjunction (or) with integer addition and conjunction (and) with integer multiplication.

We can keep all of our temporary values on the JVM operand stack, so we do not need to use any locals/registers beyond what is necessary to hold the method arguments. A simple post-order traversal of the \mathcal{F} AST produces the opcodes we need. In other words, we are converting the \mathcal{F} AST from infix form to postfix (suffix) form (*i.e.*, RPN: Reverse Polish Notation). In LAB8 we converted the \mathcal{F} AST from infix form to prefix form (*i.e.*, Polish Notation). At the end of our generated methods we will convert our computed value to either zero (false) or one (true) before returning it.

```
.class public Simulator_ex00_asm
.super java/lang/Object
.method public <init>()V
    aload_0
    invokespecial java/lang/Object.<init>()V
    return
.end method
.method public static x()Z
    .limit locals 0      ; one for each argument
    .limit stack 1       ; depth of the AST
    iconst_0             ; push false
    ifeq False           ; if top of stack is zero, then jump to False
    True:                ; the true case
    iconst_1             ; push true
    ireturn              ; return top of stack (true)
    False:               ; label we can jump to
    iconst_0             ; push false
    ireturn              ; return top of stack (false)
.end method
```

Files:

ece351.f.SimulatorGeneratorASM
ece351.f.DepthCounter

Tests:

ece351.f.TestSimulatorGeneratorASM

Infix: $a + b$

Prefix: $+ a b$

Postfix: $a b +$

$x \leq 0$;

In these listings there appears to be a space in `Object.<init>()V`. That is a phantom space inserted by the \LaTeX package that formats the source code. There is no space in the underlying source file. You should not have that space in your generated code.

```

.class public Simulator_ex01_asm                                x <= '1';
.super java/lang/Object
.method public <init>()V
    aload_0
    invokespecial java/lang/Object.<init>()V
    return
.end method
.method public static x()Z
    .limit   locals 0      ; one for each argument
    .limit   stack 1      ; depth of the AST
    iconst_1      ; push true
    ifeq False    ; if top of stack is zero, then jump to False
    True:         ; the true case
    iconst_1      ; push true
    ireturn       ; return top of stack (true)
    False:        ; label we can jump to
    iconst_0      ; push false
    ireturn       ; return top of stack (false)
.end method

```

```

.class public Simulator_ex02_asm                                x <= a;
.super java/lang/Object
.method public <init>()V
    aload_0
    invokespecial java/lang/Object.<init>()V
    return
.end method
.method public static x(Z)Z
    .limit   locals 1      ; one for each argument
    .limit   stack 1      ; depth of the AST
    iload_0       ; a
    ifeq False    ; if top of stack is zero, then jump to False
    True:         ; the true case
    iconst_1      ; push true
    ireturn       ; return top of stack (true)
    False:        ; label we can jump to
    iconst_0      ; push false
    ireturn       ; return top of stack (false)
.end method

```

x <= a or b;

```

.class public Simulator_ex05_asm
.super java/lang/Object
.method public <init>()V
    aload_0
    invokespecial java/lang/Object.<init>()V
    return
.end method
.method public static x(ZZ)Z
    .limit locals 2    ; one for each argument
    .limit stack 2     ; depth of the AST
    iload_0             ; a
    iload_1             ; b
    iadd               ; ( a or b )
    ifeq False         ; if top of stack is zero, then jump to False
    True:              ; the true case
    iconst_1           ; push true
    ireturn            ; return top of stack ( true)
    False:             ; label we can jump to
    iconst_0           ; push false
    ireturn            ; return top of stack ( false)
.end method

```

x <= c and (a or b);

```

.class public Simulator_ex06_asm
.super java/lang/Object
.method public <init>()V
    aload_0
    invokespecial java/lang/Object.<init>()V
    return
.end method
.method public static x(ZZZ)Z
    .limit locals 3    ; one for each argument
    .limit stack 3     ; depth of the AST
    iload_0             ; c
    iload_1             ; a
    iload_2             ; b
    iadd               ; ( a or b )
    imul               ; ( c and ( a or b ) )
    ifeq False         ; if top of stack is zero, then jump to False
    True:              ; the true case
    iconst_1           ; push true
    ireturn            ; return top of stack ( true)
    False:             ; label we can jump to
    iconst_0           ; push false
    ireturn            ; return top of stack ( false)
.end method

```

14.3 Negation

Negation (not) needs special consideration in this encoding. There are multiple ways it can be done.

One possibility is to generate conditionals (ifs) to test if the integer on the top of the stack is zero (*i.e.*, false) or non-zero (*i.e.*, true).

There is another possibility that does not introduce any conditionals into the generated code. Conditionals are relatively expensive, compared to integer operations. We can use bitwise xor (exclusive or) will toggle the 1-bit of an integer, which is what logical negation does. For this to work, we need a normalization function that reduces non-zero integers to one, yet still leaves zero as zero. Here is such a function that will work for non-negative integers up to half-max-int:

$$\frac{2x}{1+x}$$

14.4 Evaluation

The last pushed commit before the deadline is evaluated both on the shared test inputs and on a set of secret test inputs according to the weights in the table below. Let ‘f’ name the \mathcal{F} program we want to simulate, let ‘w’ name an input \mathcal{W} program for it, and let ‘wStaff’ name the output of the staff’s simulator for this \mathcal{F}/\mathcal{W} combination. The testing equation is:

`genASM(f).simulate(w).isomorphic(wStaff)`

	Regular Cases	Negation
TestSimulatorGenerator	70	30

14.5 Bonus Lab: \mathcal{F} to x86/x64/etc

See Lab 13. This is completely bonus. In past terms only one or two students have attempted it. It is hard. It is time consuming. It is intellectually rewarding.

The previously mentioned `z_challenge.f` input file turns out to be not that challenging: specifically, it is not triggering overflow in anyone’s implementation. So it did not meet its design objective.

One of the students this term pointed out that the handling of negation previously proposed was not fully correct. So the new ‘challenge’ is to use one of the two correct ways of handling negation now described in this manual (or come up with your own correct way).

Appendix A

Extra VHDL Lab Exercises

A.1 Define-Before-Use Checker

Compilers generally perform some form of semantic analysis on the AST of the program after the input program is parsed. The analysis might include checking that all variables/signals are defined before they are used. In this part of the lab, write a define-before-use checker that traverses through a VHDL AST and determines whether all signals that are used in the signal assignment statements are defined before they are used. In addition, for all signal assignment statements, a signal that is being assigned (left-hand side) to an expression (right-hand side) must not be an input signal. Driving an input signal simultaneously from two sources would cause undefined behaviour at run time.

Within a design unit, a signal may be defined in the entity declaration as an input bit or output bit; it may also be defined in the optional signal declaration list within the body of the corresponding architecture. For example, if we consider the VHDL program shown in Figure A.1, the following signals are defined in this entity-architecture pair: `a0`, `b0`, `a1`, `b1`, `a2`, `b2`, `a3`, `b3`, `Cin`, `sum0`, `sum1`, `sum2`, `sum3`, `Cout`, `V`, `c0`, `c1`, `c2`, `c3`, and `c4`.

The define before use checker should throw an exception if it checks the program shown in Figure A.1. This figure illustrates the two violations your checker should detect:

- a. The assignment statement in line 16 uses a signal called '`c`', which is undefined in this program.
- b. The assignment statement in line 17 tries to assign an input pin/signal ('`Cin`') to an expression.

The checker should throw a `RuntimeException` exception upon the first violation that is encountered.

All of the code that you will write for the checker should be in the class `DefBeforeUseChecker` in the package `ece351.v`. `TestDefBeforeUseCheckerAccept`

and `TestDefBeforeUseCheckerReject` are JUnit tests available for testing the checker. These two classes are found in the package `ece351.v.test`.

A.1.1 *Some Tips*

- For VHDL programs that have multiple design units, apply the violation checks per design unit (i.e., treat them separately).
- Maintain the set of declared signals in such a way that it is easy to identify where the signals are declared in the design unit.

```

1  entity four_bit_ripple_carry_adder is port (
2      a0, b0, a1, b1, a2, b2, a3, b3, Cin : in bit;
3      sum0, sum1, sum2, sum3, Cout, V: out bit
4  );
5  end four_bit_ripple_carry_adder;
6
7
8  architecture fouradder_structure of four_bit_ripple_carry_adder is
9      signal c1, c2, c3, c4: bit;
10 begin
11     FA0 : entity work.full_adder port map(a0,b0,Cin,sum0,c1);
12     FA1 : entity work.full_adder port map(a1,b1,c1,sum1,c2);
13     FA2 : entity work.full_adder port map(a2,b2,c2,sum2,c3);
14     FA3 : entity work.full_adder port map(a3,b3,c3,sum3,c4);
15
16     V <= c xor c4;
17     Cin <= c4;
18 end fouradder_structure;

```

Figure A.1: VHDL program used to illustrate signal declarations and the use of undefined signals in signal assignment statements.

A.2 *Inline \mathcal{F} intermediate variables*

```

X <= A or B;
Y <= X or C;
...
Y <= (A or B) or C;

```

Appendix B

Advanced Programming Concepts

These are some advanced programming ideas that we explored in the labs.

B.1 Immutability

Immutability is not a ‘design pattern’ in the sense that it is not in the design patterns book.¹ It is, however, perhaps the most important general design guideline in this section. Immutability has a number of important advantages:

- Immutable objects can be shared and reused freely. There is no danger of race conditions or confusing aliasing. Those kinds of bugs are subtle and difficult to debug and fix.
- No need for defensive copies. Some object-oriented programming guidelines advocate for making defensive copies of objects in order to avoid aliasing and race condition bugs.
- Representation invariants need to be checked only at object construction time, since there is no other time the object is mutated.

There are also a number of possible disadvantages:

- Object trees must be constructed bottom-up. This discipline can be useful for understanding the program, but sometimes requires a bit of adjustment for programmers who are not familiar with this discipline.
- Changing small parts of large complex object graphs can be difficult. If the object graphs are trees then the visitor design pattern can be used to make the rewriting code fairly simple.
- Sometimes data in the problem domain are mutable, and so are best modelled by mutable objects. A bank account balance is a classic example.

AST classes

¹ E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1995

B.2 Representation Invariants

Tiger: repOk() methods

Representation invariants are general rules that should be true of all objects of a particular class. For example:

- `LinkedList.head != null`
- `LinkedList.size` is the number of nodes in the list
- that the values in a certain field are sorted

These rules are often written in a `repOk()` method, which can be called at the end of every method that mutates the object (for immutable objects these methods are just the constructors). It is a good idea to check these rules right after mutations are made so that the failure (bad behaviour) of the software occurs close to the fault (error in the source code). If these rules are not checked right away, then the failure might not occur until sometime later when the execution is in some other part of the program: in those circumstances it can be difficult to find and fix the fault.

B.3 Functional Programming

Some ideas from functional programming that are useful in the object-oriented programming context:

- immutability
- recursion
- lazy computation
- higher-order functions
- computation as a transformation from inputs to outputs

In w2013 we've talked about immutability, and the code has embodied both immutability and computation as a transformation from inputs to outputs.

Appendix C

Design Patterns

The material in this section is mostly
copied from SourceMaking.com.

Design patterns represent a standard set of ideas and terms used by object-oriented programmers. If you get a job writing object-oriented code these are concepts and terms you will be expected to know.

Design patterns¹ are classified into three groups

{	<i>creational</i>
	<i>structural</i>
	<i>behavioural</i>

¹ E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1995

C.1 Iterator

- Behavioural
- Provide a way to access the elements of an aggregate object sequentially without exposing its underlying representation.
- Enables decoupling of collection classes and algorithms.
- Polymorphic traversal

C.2 Singleton

- Creational
- Ensure a class has only predetermined small number of named instances, and provide a global point of access to them.
- Encapsulated 'lazy' ('just-in-time') initialization.

ConstantExpr

C.3 Composite

- Structural
- Compose objects into tree structures to represent whole-part hierarchies. Composite lets clients treat individual objects and compositions of objects uniformly.
- Recursive composition
- e.g., 'Directories contain entries, each of which could be a file or a directory.'

Expr class hierarchy

C.4 *Template Method*

- Behavioural
- Define the skeleton of an algorithm in an operation, deferring some steps to client subclasses. Template Method lets subclasses redefine certain steps of an algorithm without changing the algorithm's structure.
- Base class declares algorithm 'placeholders', and derived classes implement the placeholders.

BinaryExpr.simplifySelf()

C.5 *Interpreter*

- Behavioural
- Given a language, define a representation for its grammar along with an interpreter that uses the representation to interpret sentences in the language.
- Map a domain to a language, the language to a grammar, and the grammar to a hierarchical object-oriented design.
- In other words, add a method to the AST for the desired operation. Then implement that method for every leaf AST class.
- An alternative to Visitor. Interpreter makes it easier to modify the grammar of the language than Visitor does, but with Visitor it is easier to add new operations.

Expr.simplify()

Tiger: §4.3 + lab manual

C.6 *Visitor*

- Behavioural
- Represent an operation to be performed on the elements of an object structure. Visitor lets you define a new operation without changing the classes of the elements on which it operates.
- Do the right thing based on the type of two objects.
- Double dispatch
- An alternative to Interpreter.

Expr class hierarchy

Appendix D

Lab Instructor Notes (GitLab)

Term names follow the UW convention: *cyy**m*, where *c* is century (0 for 20th, 1 for 21st); *yy* is the last two digits of the year; *m* is the number of the month at the start of the term (either 1, 5, or 9). So 1151 is the term that starts in January of 2015.

D.1 Repository URLs

STAFF: Same repositories used over time.

`https://git.uwaterloo.ca/ece351-notes/ece351-text`
`git@ecgit.uwaterloo.ca:drayside/ece351-code`

Source for Lab Manual + Course Notes
Staff source code

STAFF: Metadata for each offering (create fresh; §E.2):

`git@ecgit.uwaterloo.ca:ece351/term/metadata`

Offering metadata

STUDENTS: Need to be created fresh each term (see §E.2).

`git@ecgit.uwaterloo.ca:ece351/term/ece351-notes`
`git@ecgit.uwaterloo.ca:ece351/term/lib`
`git@ecgit.uwaterloo.ca:ece351/term/skeleton`
`git@ecgit.uwaterloo.ca:ece351/term/student/labs`

Compiled version of lab manual, *etc.*
Shared libraries
Code to be distributed to students
Each student's individual repo

D.2 Setting Up

D.2.1 Creating the Wildcard Repo *ece351/term*

These steps done by the Gitolite administrator:

Gitolite documentation:
<http://gitolite.com/gitolite/conf.html>

- a. `git clone git@ecgit.uwaterloo.ca:gitolite-admin`
- b. `cp conf/courses/ece351-last.conf conf/courses/ece351-term.conf`
- c. edit `conf/courses/ece351-term.conf`
 - define staff group `@ece351-term-staff`
 - define student group `@ece351-term-students`
(student group should include staff IDs)
- d. edit `conf/gitolite.conf`
 - add include for `conf/courses/ece351-term.conf`
- e. `git add conf/gitolite.conf conf/courses/ece351-term.conf`
- f. `git commit -m 'creating new wildcard repo for ECE351 term term'`
- g. `git push`

The subsequent steps can be done by the course lab instructor.

The post-commit hooks might take a long time to run (maybe an hour or more). The permissions changes will take effect almost immediately. What takes so long is updating all of the web pages: it checks which repositories are visible for each user.

D.2.2 Notes/PDF Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/ece351-notes` To create the repo
- `outline.pdf`
- `lab-manual.pdf`
- `course-notes.pdf`
- `practice-questions/`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/ece351-notes + READERS @ece351-term-students`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/ece351-notes + WRITERS @ece351-term-staff`

D.2.3 Lib Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/lib` To create the repo
- `README.txt` as first commit to create master branch
- `jar` files
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/lib + READERS @ece351-term-students`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/lib + WRITERS @ece351-term-staff`

D.2.4 Course Offering Metadata Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/metadata` To create the repo
- `README.txt` as first commit to create master branch
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/metadata + WRITERS @ece351-term-staff`
- download roster from Quest
- extract student userids
- `cd ece351-code` (change to the staff solution repo)
- `cd marking`
- `git submodule add git@ecgit.uwaterloo.ca:ece351/term/metadata metadata-term`

D.2.5 Skeleton Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/skeleton` To create the repo
- `README.txt` as first commit to create master branch
- `git submodule add git@ecgit.uwaterloo.ca:ece351/term/lib`
- `meta/`
- `.project` + `.classpath`
- `.gitignore`
- `build.xml`
- `prelab: src/ece351/util/*` [Note: See §E.6]
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/skeleton + READERS @ece351-term-students`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/skeleton + WRITERS @ece351-term-staff`

Fork student repos (see §E.3 below). Then add:

- `TestPrelab*.java`

Students will have to pull from skeleton during the prelab exercise to get these test files.

- TestImports.java

D.3 Forking Student Repos

See scripts/fork.sh

```
#!/bin/bash

echo "copy this script to the appropriate location and customize it"
exit 1

t=1181

for s in `cat students.txt`
do
    echo "*****"
    echo $s
    git clone git@ecgit.uwaterloo.ca:ece351/$t/$s/labs $s
    cd $s
    git remote add skeleton git@ecgit.uwaterloo.ca:ece351/$t/skeleton
    git pull skeleton master
    git push
    ssh git@ecgit.uwaterloo.ca perms ece351/$t/$s/labs + WRITERS $s
    ssh git@ecgit.uwaterloo.ca perms ece351/$t/$s/labs + WRITERS @ece351-$t-staff
    #ssh git@ecgit.uwaterloo.ca perms ece351/$t/$s/labs -l
    cd ..
done
```

D.4 Marking Script: Build.xml

The students also have a copy of this.

D.5 Web User Interface for Gitolite Server

Written with Django/Python. Source available at:

<https://github.com/eyolfson/site-ecegit>

It depends on two Django apps:

<https://github.com/eyolfson/django-gitolite>

<https://github.com/eyolfson/django-ssh>

D.6 Exporting to Skeleton: *export.sh*

Use the `export.sh` script to prepare the source files for the skeleton. *Never copy source files directly from the staff repo to the skeleton.* The `export.sh` script performs two functions: (1) snips out the parts that the students are supposed to write, and (2) inserts the intellectual property header.

If `export.sh` runs successfully you should see many lines of output that start with `'dos2unix'`. If you get a message that the temporary directory exists and the script is exiting, you need to delete that old temporary directory first and then run `export.sh`.

After running `export.sh` use a visual diff tool such as `meld` to compare the snipped source code with the skeleton. Merge manually. Sometimes there will be changes that are made to the skeleton manually that should not be overwritten by a future export.

Don't forget to also copy the test input files into the skeleton!

Never copy source files directly from the staff repo to the skeleton.

`export.sh` makes a clone of your local copy of the instructor code repo and works from that. So if you have uncommitted changes in your local copy of the repo they will not be included in the export.

D.6.1 Update *header.txt* at the start of term

At the beginning of each term, update the file `src/ece351/header.txt` to have the current term date and number. This file is pre-pended to each source file by the export script.

D.6.2 Export on an as-needed basis

Experience has shown that it's best to do four major releases during the term: prelab, \mathcal{W} , \mathcal{F} , and `VHDL`. Exporting all of the code at the beginning is overwhelming for many students. Also, the gradual export strategy makes it easier for the staff to revise later labs after the term has started.

D.6.3 Exporting to skeleton dev branch

Sometimes it is a good idea to make a dev branch in the skeleton repo and export to that first. This way staff and advanced students can test the code before it is released to the class as a whole. Also, usually the code for lab $k + 1$ is released to dev before the deadline for lab k . After the deadline for k then dev is merged to skeleton master. If code for lab $k + 1$ is released to students before the deadline for lab k it can cause a lot of confusion, especially if that code is mixed with critical patches for lab k .

D.6.4 Files to Release for Each Lab

- o. update header.txt before populating skeleton §E.6.1

README.txt

lib submodule

.gitignore

.project

.classpath

build.xml

meta/hours.txt + meta/collaboration.txt

src/ece351/util/*.java

1. tests/wave/*.wave

src/ece351/w/ast/*

src/ece351/w/rdescent/*

src/ece351/w/regex/*

2. tests/wave/staff.out/*

src/ece351/objectcontract/*

src/ece351/w/svg/*

3. tests/f/* *including ungrammatical/, gates/ and staff.out/*

src/ece351/common/ast/*

src/ece351/common/visitor/* *for compilation dependencies*

src/ece351/f/FParser.java

src/ece351/f/analysis/* *for compilation dependencies*

src/ece351/f/ast/*

src/ece351/f/rdescent/*

src/ece351/f/test/*

4. src/ece351/f/simplifier/*

5. src/ece351/w/parboiled/*

6. src/ece351/f/parboiled/*

7. src/ece351/f/techmapper/*

tests/f/staff.out/graph/*

8. tests/f/staff.out/simulator/*

src/ece351/f/simgen/TestSimulatorGenerator.java

src/ece351/f/simgen/SimulatorGenerator.java

9. tests/v/*

src/ece351/v/VBase.java

src/ece351/v/ast/*

src/ece351/v/test/*

src/ece351/v/test/VRecognizer.java (or class)

src/ece351/v/test/VParser.java (or class)

src/ece351/v/test/TestArchitectureEquivalence

10. src/ece351/v/PostOrderVVisitor.java

src/ece351/v/DeSugarer.java

src/ece351/v/Elaborator.java

After the base files, fork student repos. Then add TestImports.java and TestPrelab.java

Manually remove Todo351Exception from TestWRegexSimpleData.

Excluding tests/f/secret!

```
src/ece351/v/test/TestDeSugarer.java  
src/ece351/v/test/TestElaborator.java
```


Appendix E

Lab Instructor Notes (*gitolite*)

Term names follow the UW convention: *cyy**m*, where *c* is century (0 for 20th, 1 for 21st); *yy* is the last two digits of the year; *m* is the number of the month at the start of the term (either 1, 5, or 9). So 1151 is the term that starts in January of 2015.

E.1 Repository URLs

STAFF: Same repositories used over time.

`https://git.uwaterloo.ca/ece351-notes/ece351-text`
`git@ecgit.uwaterloo.ca:drayside/ece351-code`

Source for Lab Manual + Course Notes
Staff source code

STAFF: Metadata for each offering (create fresh; §E.2):

`git@ecgit.uwaterloo.ca:ece351/term/metadata`

Offering metadata

STUDENTS: Need to be created fresh each term (see §E.2).

`git@ecgit.uwaterloo.ca:ece351/term/ece351-notes`
`git@ecgit.uwaterloo.ca:ece351/term/lib`
`git@ecgit.uwaterloo.ca:ece351/term/skeleton`
`git@ecgit.uwaterloo.ca:ece351/term/student/labs`

Compiled version of lab manual, *etc.*
Shared libraries
Code to be distributed to students
Each student's individual repo

E.2 Setting Up

E.2.1 Creating the Wildcard Repo *ece351/term*

These steps done by the Gitolite administrator:

Gitolite documentation:
<http://gitolite.com/gitolite/conf.html>

- a. `git clone git@ecgit.uwaterloo.ca:gitolite-admin`
- b. `cp conf/courses/ece351-last.conf conf/courses/ece351-term.conf`
- c. edit `conf/courses/ece351-term.conf`
 - define staff group `@ece351-term-staff`
 - define student group `@ece351-term-students`
(student group should include staff IDs)
- d. edit `conf/gitolite.conf`
 - add include for `conf/courses/ece351-term.conf`
- e. `git add conf/gitolite.conf conf/courses/ece351-term.conf`
- f. `git commit -m 'creating new wildcard repo for ECE351 term term'`
- g. `git push`

The subsequent steps can be done by the course lab instructor.

The post-commit hooks might take a long time to run (maybe an hour or more). The permissions changes will take effect almost immediately. What takes so long is updating all of the web pages: it checks which repositories are visible for each user.

E.2.2 Notes/PDF Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/ece351-notes` To create the repo
- `outline.pdf`
- `lab-manual.pdf`
- `course-notes.pdf`
- `practice-questions/`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/ece351-notes + READERS @ece351-term-students`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/ece351-notes + WRITERS @ece351-term-staff`

E.2.3 Lib Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/lib` To create the repo
- `README.txt` as first commit to create master branch
- `jar` files
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/lib + READERS @ece351-term-students`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/lib + WRITERS @ece351-term-staff`

E.2.4 Course Offering Metadata Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/metadata` To create the repo
- `README.txt` as first commit to create master branch
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/metadata + WRITERS @ece351-term-staff`
- download roster from Quest
- extract student userids
- `cd ece351-code` (change to the staff solution repo)
- `cd marking`
- `git submodule add git@ecgit.uwaterloo.ca:ece351/term/metadata metadata-term`

E.2.5 Skeleton Repository

- `git clone git@ecgit.uwaterloo.ca:ece351/term/skeleton` To create the repo
- `README.txt` as first commit to create master branch
- `git submodule add git@ecgit.uwaterloo.ca:ece351/term/lib`
- `meta/`
- `.project` + `.classpath`
- `.gitignore`
- `build.xml`
- `prelab: src/ece351/util/*` [Note: See §E.6]
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/skeleton + READERS @ece351-term-students`
- `ssh git@ecgit.uwaterloo.ca perms ece351/term/skeleton + WRITERS @ece351-term-staff`

Fork student repos (see §E.3 below). Then add:

- `TestPrelab*.java`

Students will have to pull from skeleton during the prelab exercise to get these test files.

- TestImports.java

E.3 *Forking Student Repos*

See scripts/fork.sh

```
#!/bin/bash

echo "copy this script to the appropriate location and customize it"
exit 1

t=1181

for s in `cat students.txt`
do
    echo "*****"
    echo $s
    git clone git@ecgit.uwaterloo.ca:ece351/$t/$s/labs $s
    cd $s
    git remote add skeleton git@ecgit.uwaterloo.ca:ece351/$t/skeleton
    git pull skeleton master
    git push
    ssh git@ecgit.uwaterloo.ca perms ece351/$t/$s/labs + WRITERS $s
    ssh git@ecgit.uwaterloo.ca perms ece351/$t/$s/labs + WRITERS @ece351-$t-staff
    #ssh git@ecgit.uwaterloo.ca perms ece351/$t/$s/labs -l
    cd ..
done
```

E.4 *Marking Script: Build.xml*

The students also have a copy of this.

E.5 *Web User Interface for Gitolite Server*

Written with Django/Python. Source available at:

<https://github.com/eyolfson/site-ecegit>

It depends on two Django apps:

<https://github.com/eyolfson/django-gitolite>

<https://github.com/eyolfson/django-ssh>

E.6 Exporting to Skeleton: *export.sh*

Use the `export.sh` script to prepare the source files for the skeleton. *Never copy source files directly from the staff repo to the skeleton.* The `export.sh` script performs two functions: (1) snips out the parts that the students are supposed to write, and (2) inserts the intellectual property header.

If `export.sh` runs successfully you should see many lines of output that start with `'dos2unix'`. If you get a message that the temporary directory exists and the script is exiting, you need to delete that old temporary directory first and then run `export.sh`.

After running `export.sh` use a visual diff tool such as `meld` to compare the snipped source code with the skeleton. Merge manually. Sometimes there will be changes that are made to the skeleton manually that should not be overwritten by a future export.

Don't forget to also copy the test input files into the skeleton!

Never copy source files directly from the staff repo to the skeleton.

`export.sh` makes a clone of your local copy of the instructor code repo and works from that. So if you have uncommitted changes in your local copy of the repo they will not be included in the export.

E.6.1 Update *header.txt* at the start of term

At the beginning of each term, update the file `src/ece351/header.txt` to have the current term date and number. This file is pre-pended to each source file by the export script.

E.6.2 Export on an as-needed basis

Experience has shown that it's best to do four major releases during the term: prelab, \mathcal{W} , \mathcal{F} , and `VHDL`. Exporting all of the code at the beginning is overwhelming for many students. Also, the gradual export strategy makes it easier for the staff to revise later labs after the term has started.

E.6.3 Exporting to skeleton dev branch

Sometimes it is a good idea to make a dev branch in the skeleton repo and export to that first. This way staff and advanced students can test the code before it is released to the class as a whole. Also, usually the code for lab $k + 1$ is released to dev before the deadline for lab k . After the deadline for k then dev is merged to skeleton master. If code for lab $k + 1$ is released to students before the deadline for lab k it can cause a lot of confusion, especially if that code is mixed with critical patches for lab k .

E.6.4 Files to Release for Each Lab

- o. update header.txt before populating skeleton §E.6.1
 - README.txt
 - lib submodule
 - .gitignore
 - .project
 - .classpath
 - build.xml
 - meta/hours.txt + meta/collaboration.txt
 - src/ece351/util/*.java
- 1. tests/wave/*.wave
 - src/ece351/w/ast/*
 - src/ece351/w/rdescent/*
 - src/ece351/w/regex/*
- 2. tests/wave/staff.out/*
 - src/ece351/objectcontract/*
 - src/ece351/w/svg/*
- 3. tests/f/* *including ungrammatical/, gates/ and staff.out/*
 - src/ece351/common/ast/*
 - src/ece351/common/visitor/* *for compilation dependencies*
 - src/ece351/f/FParser.java
 - src/ece351/f/analysis/* *for compilation dependencies*
 - src/ece351/f/ast/*
 - src/ece351/f/rdescent/*
 - src/ece351/f/test/*
- 4. src/ece351/f/simplifier/*
- 5. src/ece351/w/parboiled/*
- 6. src/ece351/f/parboiled/*
- 7. src/ece351/f/techmapper/*
 - tests/f/staff.out/graph/*
- 8. tests/f/staff.out/simulator/*
 - src/ece351/f/simgen/TestSimulatorGenerator.java
 - src/ece351/f/simgen/SimulatorGenerator.java
- 9. tests/v/*
 - src/ece351/v/VBase.java
 - src/ece351/v/ast/*
 - src/ece351/v/test/*
 - src/ece351/v/test/VRecognizer.java (or class)
 - src/ece351/v/test/VParser.java (or class)
 - src/ece351/v/test/TestArchitectureEquivalence
- 10. src/ece351/v/PostOrderVVisitor.java
 - src/ece351/v/DeSugarer.java
 - src/ece351/v/Elaborator.java

After the base files, fork student repos. Then add TestImports.java and TestPrelab.java

Manually remove Todo351Exception from TestWRegexSimpleData.

Excluding tests/f/secret!

```
src/ece351/v/test/TestDeSugarer.java  
src/ece351/v/test/TestElaborator.java
```


Appendix F

Bibliography

- [1] A. W. Appel and J. Palsberg. *Modern Compiler Implementation in Java*. Cambridge, 2004.
- [2] J. Bloch. *Effective Java*. Addison-Wesley, 2001.
- [3] R. E. Bryant. Graph-based algorithms for boolean function manipulation. *IEEE Transactions on Computers*, C-35(8):677–691, Aug. 1986.
- [4] B. Eckel. *Thinking in Java*. Prentice-Hall, 2002. <http://www.mindview.net/Books/TIJ/>.
- [5] E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1995.
- [6] B. Liskov and J. Guttag. *Program Development in Java: Abstraction, Specification, and Object-Oriented Design*. Addison-Wesley, 2001.
- [7] J. C. Reynolds. User-defined types and procedural data as complementary approaches to data abstraction. In S. A. Schuman, editor, *New Directions in Algorithmic Languages: IFIP Working Group 2.1 on Algol*. INRIA, 1975.
- [8] M. L. Scott. *Programming Language Pragmatics*. Morgan Kaufmann, 3 edition, 2009.
- [9] M. Torgersen. The Expression Problem Revisited: Four new solutions using generics. In M. Odersky, editor, *Proc. 18th ECOOP*, volume 3344 of *LNCS*, Oslo, Norway, June 2004. Springer-Verlag.
- [10] P. Wadler. The expression problem, Nov. 1998. Email to Java Generics list.
- [11] W. Wulf and M. Shaw. Global variables considered harmful. *ACM SIGPLAN Notices*, 8(2):80–86, Feb. 1973.

- [12] M. Zenger and M. Odersky. Independently extensible solutions to the expression problem. In *Proc.12th Workshop on Foundations of Object-Oriented Languages*, 2005.